



UNIVERSITÀ DI PISA

DIPARTIMENTO DI INGEGNERIA DELL'INFORMAZIONE

MASTERS IN ARTIFICIAL INTELLIGENCE AND
DATA ENGINEERING

*Edge AI Computing for Grape Leaf Disease Detection in
LoRaWAN-Based Vineyard Monitoring System*

Advisories:

Prof. Vallati Carlo

Dr. Fabrizio Palla

Prof. Zinnai Angela

Candidate:

Tsegay Teklay Gebrelibanos
ID Number: 683925

Academic Year 2025/2026

Abstract

Existing vineyard monitoring systems primarily rely on environmental data collection such as temperature, humidity, and pressure, transmitted over LoRaWAN networks and providing limited insights into the actual grapevine's health status. This thesis work focuses on enhancing the vineyard monitoring intelligence by integrating Edge AI computing at sensor level for grape leaf disease detection. Moving beyond the traditional environmental-only monitoring approaches, this research was conducted on developing and optimizing deep learning models (CNN for classification and YOLO for detection) and deploying directly on resource-constrained edge sensors. This extension enabled the monitoring system to capture and process leaf images to detect diseases directly at the edge, providing real-time disease detection capabilities while adapting the inherent bandwidth and power limitations of LoRaWAN-based wireless sensor networks. This disease detection feature is implemented using four-stage pipeline architecture and deployed on ESP32-S3-microcontroller. **Stage 1** captures image of vine-trees using the OV3660 camera module operating at VGA resolution with JPEG compression. **Stage 2** employs an ESP-optimized YOLO11n model (ESPDet-Pico: 478KB INT8) to localize individual grape leaves within the captured frame, achieving 56.4% mAP@50 for leaf detection. Then these detected leaf regions are cropped and forwarded to **Stage 3**, where a MobileNetV2 classifier model (2.73 MB INT8) is utilized to identify disease symptoms across the four classes—Black Rot, Esca, Healthy, and Leaf-Blight with $99.8\% \pm 0.05$ accuracy on individual leaf patches. **Stage 4** aggregates the classification results from multiple detected leaves to produce a single output formatted for LoRaWAN transmission. Since each captured frame typically contains multiple grape leaves, the weighted aggregation algorithm combines detection confidence, classification entropy, and spatial quality metrics from up to 15 detected leaves to produce robust per-frame disease diagnoses result. The aggregated classification result is then packaged into LoRaWAN uplink packet (e.g. "Leaf Blight, 80% confidence") and transmitted over the WSNs.

To optimize also the power consumption, the system works on duty-cycle operation to capture and preprocess image and apply inference at defined intervals and then enters into deep sleep mode. Overall this thesis work demonstrated the feasibility of deploying deep learning models on resources constrained microcontrollers specifically for disease detection in vineyard monitoring system and precision agriculture applications at large.

Keywords: Wireless Sensor Networks, LoRaWAN, Precision Agriculture, Grape Leaf Disease Detection, Edge AI Computing, Object Detection Model, Image Classification, Quantization, ESP32-S3.



Acknowledgments

I would like to express my sincere gratitude to my supervisor, Prof. [Supervisor Name], for their invaluable guidance, continuous support, and insightful feedback throughout this research work.

I am deeply grateful to the team at INFN (Istituto Nazionale di Fisica Nucleare) for providing me with the opportunity to work on this challenging and rewarding project. Special thanks to [Co-supervisor/Mentor Name] for their technical expertise and mentorship during the development and implementation phases.

I would like to thank my family for their unwavering support and encouragement throughout my academic journey.

Finally, I acknowledge the use of computational resources and datasets that made this work possible.

Tsegay
University of Pisa
February 2026



Contents

Abstract	i
Acknowledgments	ii
List of Figures	xi
List of Tables	xii
List of Acronyms	xiv
1 Introduction and Background	1
2 Related Work	4
2.1 Precision Agriculture and Wireless Sensor Networks	4
2.2 Low-Power Wide-Area Networks for Agricultural IoT	5
2.3 Computer Vision for Plant Disease Detection	6
2.4 Object Detection for Plant Monitoring	7
2.5 Edge AI on Resource-Constrained Devices	7
2.6 Integrated Systems for Agricultural Monitoring	9
2.7 Summary and Research Gaps	9
3 Edge AI System Design & Implementation	11
3.1 System Overview and Methodology	11
3.1.1 Four-Stage Pipeline Architecture	11
3.2 Dataset and Preprocessing Techniques	12
3.2.1 The Grape-Leaf Disease Dataset	12
3.2.2 Grapevine Dataset for Leaf Localization	13
3.2.3 Preprocessing and Augmentation Pipeline	13
3.2.3.1 Resolution Adaptation for Embedded Constraints	13
3.2.3.2 Intensity Normalization and Standardization	14



3.2.3.3	Geometric and Photometric Augmentation Techniques	14
3.2.3.4	Augmentation-Free Validation Protocol	15
3.3	Leaf Localization ESPDet-Pico Model (YOLO-based)	15
3.3.1	Object Detection Framework Selection	15
3.3.2	Model Compression for Embedded Deployment	16
3.3.2.1	Backbone Architecture Modifications	16
3.3.2.2	Detection Head Architecture	17
3.3.2.3	Parameter and Memory Reduction	17
3.3.3	Training Methodology	17
3.3.3.1	Input Resolution and Rectangular Training	18
3.3.3.2	Optimization Configuration	18
3.3.3.3	Loss Function	18
3.3.3.4	Data Augmentation Strategy	19
3.3.3.5	Training Convergence	19
3.3.4	Detection Performance Evaluation	20
3.3.4.1	ESPDet-Pico Performance Metrics	20
3.3.4.2	Comparative Analysis: ESPDet-Pico vs YOLO11n	20
3.3.4.3	Deployment Feasibility and Trade-off Justification	21
3.4	Disease Classification Model (MobileNetV2)	21
3.4.1	Convolutional Neural Network Architecture	22
3.4.1.1	Depthwise Separable Convolutions	22
3.4.1.2	Inverted Residual Blocks	22
3.4.1.3	Model Configuration and Pre-training	23
3.4.2	Transfer Learning Strategy	23
3.4.2.1	Stage 1: Classifier Head Training	23
3.4.2.2	Stage 2: Full Model Fine-tuning	24
3.4.3	Training Protocol	24
3.4.3.1	Optimization Configuration	24
3.4.3.2	Loss Function	24



3.4.3.3	Regularization Strategy	25
3.4.3.4	Multi-Seed Experimental Protocol	25
3.4.4	Classification Accuracy and Robustness	25
3.4.4.1	Overall Test Set Performance	25
3.4.4.2	Per-Class Performance Analysis	26
3.4.4.3	Confusion Matrix Interpretation	27
3.4.5	Post-Training Quantization	27
3.4.5.1	ESP-PPQ Quantization Framework	27
3.4.5.2	Quantization-Induced Accuracy Degradation	28
3.4.5.3	Model Size Reduction and Memory Footprint	28
3.5	Hardware Selection for ESP32-S3-WROOM-1 Dev. Kit	28
3.5.1	ESP32-S3 System Architecture	29
3.5.2	Hardware Selection Rationale	30
3.6	Edge Deployment of Integrated Lightweight Neural Networks on ESP32-S3	32
3.6.1	Memory Management Architecture	33
3.6.1.1	Model Storage and Loading Strategy	33
3.6.1.2	Initialization Order and PSRAM Fragmentation	34
3.6.1.3	Buffer Allocation and Reuse	34
3.6.2	Inference Pipeline Implementation	35
3.6.2.1	Camera Acquisition and Image Decoding	35
3.6.2.2	Leaf Detection and Bounding Box Extraction	36
3.6.2.3	Region Cropping and Classification Preprocessing	36
3.6.2.4	Disease Classification Inference and Softmax Decoding	37
3.6.2.5	Multi-Leaf Aggregation for Frame-Level Diagnosis	37
3.6.3	Deployment Validation	39
4	Results and Discussion	40
4.1	On-Device Inference Performance	40
4.1.1	Pipeline Latency Characterization	40
4.1.1.1	Per-Stage Latency Breakdown	40



4.1.1.2	Analysis of Computational Bottlenecks	41
4.1.2	MobileNetV2 Inference Scaling Analysis	41
4.1.2.1	Per-Leaf Inference Characteristics	42
4.1.2.2	Linear Scaling Validation	42
4.1.2.3	SIMD Utilization Efficiency	42
4.1.3	Memory Consumption and Stability Analysis	43
4.1.3.1	Static Memory Allocation	43
4.1.3.2	Long-Term Stability Testing	43
4.1.4	Comparative Analysis with Alternative Platforms	44
4.2	Inference Aggregation Strategies	45
4.2.1	Aggregation Strategy Formulations	45
4.2.1.1	Baseline: Detection-Confidence Weighting	45
4.2.1.2	Entropy-Based Uncertainty Weighting	45
4.2.1.3	Hybrid Spatial-Quality Weighting	46
4.2.2	Experimental Methodology	46
4.2.2.1	Evaluation Dataset and Protocol	46
4.2.2.2	Implementation and Computational Overhead	47
4.2.3	Results and Comparative Analysis	47
4.2.3.1	Baseline Performance	48
4.2.3.2	Entropy-Aware Strategy Improvements	48
4.2.3.3	Hybrid Strategy Superiority	48
4.2.4	Ablation Study and Component Analysis	49
4.2.5	Deployment Considerations	49
4.3	LoRaWAN Packet Structure and Communication Protocol Design	50
4.3.1	LoRaWAN Network Architecture	50
4.3.1.1	End Devices (ESP32-S3 Nodes)	50
4.3.1.2	Gateway Infrastructure	51
4.3.1.3	Network Server and Application Server	51
4.3.2	Packet Structure Design	51



4.3.2.1	Byte-Level Packet Layout	51
4.3.2.2	Field Encoding Rationale	51
4.3.3	Communication Protocol Workflow	53
4.3.3.1	Uplink Transmission Sequence	54
4.3.3.2	Downlink Reception and Configuration Updates	54
4.3.3.3	Gateway and Network Server Processing	54
4.3.4	Payload Efficiency and Bandwidth Analysis	55
4.3.4.1	Information Density	55
4.3.4.2	Network Capacity Estimation	55
4.3.4.3	Comparison with Alternative Protocols	56
4.3.5	Security and Authentication	56
4.3.5.1	Network-Level Security (NwkSKey)	56
4.3.5.2	Application-Level Security (AppSKey)	56
4.3.5.3	Over-The-Air Activation (OTAA)	57
4.4	Duty-Cycled Edge AI Operations for Energy Efficiency	57
4.4.1	Duty Cycle Configuration and Timing Analysis	57
4.4.1.1	Operational State Timeline	57
4.4.1.2	Duty Cycle Percentage Calculation	58
4.4.2	Energy Budget Analysis	59
4.4.2.1	Key Energy Insights	59
4.4.2.2	Daily and Monthly Energy Requirements	60
4.4.3	Solar Energy Harvesting and Battery Sizing	60
4.4.3.1	Solar Energy Availability	60
4.4.3.2	Energy Balance Analysis	61
4.4.3.3	Battery Capacity Requirements and Autonomy	61
4.4.3.4	Comparison with Alternative Battery Capacities	62
4.4.4	Long-Term Deployment Sustainability	62
4.4.4.1	Seasonal Energy Performance	62
4.4.4.2	System Uptime and Reliability	63



4.4.4.3	Multi-Year Deployment Projections	63
4.4.5	Comparison with Continuous Operation	64
4.4.5.1	Continuous Operation Infeasibility	64
4.4.5.2	1-Minute Duty Cycle Over-Engineering	64
4.4.5.3	10-Minute and 30-Minute Trade-offs	65
4.4.6	Power Optimization Strategies	65
4.4.6.1	Dynamic Voltage and Frequency Scaling (DVFS)	65
4.4.6.2	Adaptive Duty Cycle Based on Disease Detection	65
4.4.6.3	Model Compression and Pruning	66
4.4.7	Summary and Energy Efficiency Conclusions	66
5	Conclusion and Future Work	68
5.1	Summary of Contributions	68
5.1.1	Dual-Model Cascade Architecture	68
5.1.2	Multi-Leaf Aggregation Strategy	68
5.1.3	Solar-Powered Duty-Cycled Operation	68
5.1.4	LoRaWAN Protocol Design	69
5.2	Key Findings	69
5.3	Limitations	69
5.4	Future Work	70
5.5	Concluding Remarks	70
	Bibliography	72
A	Code Repository Structure	79
A.1	Repository Overview	79
A.2	Directory Structure	79
A.3	Key Files Description	80
A.3.1	Training Files	80
A.3.2	Firmware Files	80
A.4	Building and Running	80



A.4.1	Training Models	80
A.4.2	Flashing ESP32-S3	80
A.5	Repository Access	81
B	Hyperparameters and Configuration	82
B.1	MobileNetV2 Training Configuration	82
B.1.1	Stage 1: Feature Extraction (Warm-up)	82
B.1.2	Stage 2: Full Fine-Tuning	82
B.2	Data Augmentation	83
B.2.1	Training Augmentation	83
B.2.2	Validation/Test Transform	83
B.3	ESPDet-Pico (YOLO) Training Configuration	84
B.3.1	Detection Augmentation	84
B.4	Quantization Configuration	85
B.5	ESP32-S3 Hardware Configuration	85
B.6	Power Management Configuration	86
C	ESP-IDF Build Instructions	87
C.1	Prerequisites	87
C.1.1	Software Requirements	87
C.1.2	Hardware Requirements	87
C.2	ESP-IDF Installation	87
C.2.1	Linux / macOS	87
C.2.2	Windows	88
C.3	Project Setup	88
C.3.1	Key Configuration Options	88
C.4	Building the Project	88
C.5	Flashing the Device	89
C.5.1	Automatic Flashing	89
C.5.2	Manual Flashing with esptool	89



C.6	Serial Monitoring	89
C.6.1	Expected Output	90
C.7	Common Issues and Solutions	90
C.8	OTA (Over-The-Air) Updates	90
C.8.1	Partition Table Configuration	90
C.8.2	OTA Update Code	91
C.9	Performance Profiling	91
C.9.1	Memory Profiling	91
C.9.2	CPU Timing	91
C.10	Additional Commands	92
C.11	Useful Resources	92
D	Additional Experimental Results	93
D.1	Extended Performance Metrics	93
D.1.1	Per-Class Classification Metrics	93
D.2	Memory Utilization Over Time	93
D.3	Power Consumption Breakdown	94
D.4	Aggregation Strategy Comparison	94
D.5	Inference Latency Distribution	94
D.6	Model Size Comparison	95
D.7	Battery Life Under Different Scenarios	95
D.8	LoRaWAN Packet Structure Details	96
D.8.1	Uplink Packet Format	96
D.9	Comparison with State-of-the-Art	96
D.10	Error Analysis	96
D.10.1	Common Misclassification Patterns	96

List of Figures

3.1 Four-stage edge AI pipeline for grape leaf disease detection.	11
---	----

List of Tables

3.1	Grape leaf disease dataset distribution across training, validation, test splits	12
3.2	ESPDet-Pico backbone architecture with layer-wise specifications	16
3.3	Comparative evaluation of ESPDet-Pico and YOLO11n baseline on grape leaf detection	21
3.4	Per-class classification performance of MobileNetV2 on grape disease test set	26
3.5	Comparison of embedded AI hardware platforms for vineyard edge deployment.	30
4.1	Per-stage latency breakdown for ESP32-S3 edge inference pipeline (7 leaves detected)	41
4.2	Inference performance comparison across edge AI platforms (7-leaf frame)	44
4.3	Frame-level classification accuracy comparison across aggregation strategies	47
4.4	Ablation study: incremental contribution of aggregation components . . .	49
4.5	LoRaWAN uplink packet structure for disease diagnosis (16 bytes total) .	52
4.6	LPWAN protocol comparison for vineyard monitoring application	56
4.7	Energy consumption breakdown for duty-cycled operation (per 5-minute cycle)	59
4.8	Battery capacity vs. autonomy for duty-cycled operation (median case: 0.815 mW)	62
4.9	Energy consumption and battery lifetime comparison across operational modes	64
B.1	MobileNetV2 Stage 1 hyperparameters	82
B.2	MobileNetV2 Stage 2 hyperparameters	82
B.3	Training data augmentation parameters	83
B.4	Validation and test data preprocessing	83
B.5	ESPDet-Pico training hyperparameters	84
B.6	YOLO augmentation parameters	84
B.7	INT8 quantization parameters	85
B.8	ESP32-S3 configuration parameters	85



B.9	Power management settings	86
C.1	Troubleshooting common build/flash issues	90
D.1	Detailed per-class performance metrics	93
D.2	PSRAM usage at different pipeline stages	93
D.3	Power consumption by component	94
D.4	Detailed comparison of aggregation methods	94
D.5	Detection latency statistics (1000 runs)	94
D.6	Classification latency statistics per leaf (1000 runs)	95
D.7	Model size comparison with other edge AI solutions	95
D.8	Estimated battery life (3000 mAh battery)	95
D.9	LoRaWAN uplink packet structure	96
D.10	Comparison with related work on edge AI plant disease detection	96
D.11	Most common misclassifications	96

List of Acronyms

Acronym	Definition
AI	Artificial Intelligence
CNN	Convolutional Neural Network
DL	Deep Learning
DMA	Direct Memory Access
ESP	Espressif Systems Processor
FP	False Positive
FP32	32-bit Floating Point
GPU	Graphics Processing Unit
INT8	8-bit Integer Quantization
IoT	Internet of Things
LoRaWAN	Long Range Wide Area Network
LPWAN	Low Power Wide Area Network
mAP	mean Average Precision
MCU	Microcontroller Unit
ML	Machine Learning
NN	Neural Network
ONNX	Open Neural Network Exchange
OTA	Over-The-Air
PTQ	Post-Training Quantization
PSRAM	Pseudo-Static Random Access Memory
RAM	Random Access Memory
ROI	Region of Interest
SRAM	Static Random Access Memory
TFLite	TensorFlow Lite
UART	Universal Asynchronous Receiver-Transmitter
WSN	Wireless Sensor Network
YOLO	You Only Look Once



Chapter 1

Introduction and Background

Viticulture, one of the world's most economically significant agricultural sectors, faces persistent challenges from foliar diseases that threaten grape quality and yield. Traditional vineyard monitoring relies on manual scouting by agronomists who visually inspect vines for disease symptoms—a process that is labor-intensive, subjective, and often detects infections too late for effective intervention [1]. The economic impact is substantial: major grape diseases such as Black Rot, Esca, and Leaf Blight can reduce yields by 30-50% in severely affected vineyards, while preventative fungicide applications contribute to 20-30% of operational costs and raise environmental concerns about pesticide residues in soil and groundwater.

The convergence of three technological advances—wireless sensor networks, low-power wide-area communication, and edge artificial intelligence—creates an opportunity to transform vineyard management through continuous, automated plant health monitoring. This thesis investigates whether sophisticated computer vision inference for disease detection can be integrated into ultra-low-power sensor networks while maintaining the energy efficiency required for solar-powered, multi-month autonomous operation. The fundamental challenge is reconciling the bandwidth and power demands of visual sensing with the stringent constraints of battery-powered agricultural IoT deployments.

Wireless sensor networks (WSN) have become a cornerstone technology for precision agriculture, providing spatially-distributed environmental monitoring across hectares of farmland without the infrastructure costs of wired systems [1, 2]. Contemporary vineyard WSN deployments typically measure temperature, humidity, atmospheric pressure, and soil moisture at 10-30 minute intervals, transmitting concise telemetry packets (20-50 bytes) that enable years of operation on coin-cell batteries. These environmental measurements inform irrigation scheduling and disease risk models, but they provide only indirect indicators of plant health. When environmental conditions favor fungal infections (e.g., extended periods of high humidity and moderate temperatures), agronomists must still physically inspect vines to confirm whether disease symptoms have actually manifested—by which time treatment may be too late to prevent yield loss.

Visual sensing with cameras offers direct observation of plant health through analysis of leaf color, texture, and morphology. However, transmitting raw images—even compressed JPEG files of 100-500kB—is fundamentally incompatible with the power budgets of agricultural sensor nodes. A LoRaWAN radio consuming 120mA at 3.3V during transmission would require 15-75 seconds to send a single image at typical data rates of 0.3-5.5kbps, depleting a 2000mAh battery in days rather than months. Cellular connectivity (LTE-M, NB-IoT) offers higher bandwidth but consumes comparable energy per bit and incurs subscription fees that become prohibitive for dense sensor deployments with hundreds of nodes per vineyard.

Edge AI computing—deploying deep learning inference directly on sensor nodes—resolves this bandwidth-power dilemma by transmitting only inference results rather than raw



imagery. A disease classification result such as "Leaf Blight detected with 92% confidence" requires merely 10-20 bytes, orders of magnitude less than the source image, while still providing actionable plant health diagnostics. This approach has been enabled by two recent developments: efficient neural network architectures such as MobileNets [3] and YOLO [4] that achieve high accuracy with 10-100 $\times$ fewer parameters than earlier CNN models, and microcontrollers with hardware acceleration features (SIMD vector instructions, neural processing units) that execute inference at practical speeds despite consuming only milliwatts of power [5].

The ESP32-S3 microcontroller exemplifies this new generation of edge AI-capable platforms. Its dual-core Xtensa LX7 processors running at 240MHz include vector extensions that accelerate the multiply-accumulate operations central to convolutional neural networks, while 512KB of internal SRAM and 8MB of external PSRAM provide sufficient memory for model weights and intermediate activations. Espressif's ESP-DL framework supplies optimized kernels that exploit these architectural features, enabling inference at 5-15 frames per second for INT8-quantized models. When combined with deep sleep modes consuming only 8 μ W between inferences, the ESP32-S3 can operate on battery or solar-harvested power for extended periods.

Deep learning models for plant disease classification have matured considerably over the past five years. Convolutional neural networks trained on publicly available datasets such as PlantVillage [6]—which contains 54,000+ labeled images of healthy and diseased leaves across 14 crop species—routinely achieve 95-99% accuracy on held-out test sets [7]. However, most published research evaluates models on desktop GPUs or embedded Linux platforms (NVIDIA Jetson, Raspberry Pi), with limited attention to the resource constraints of microcontroller-class processors. The typical deployment pipeline combines object detection to localize individual leaves within images, followed by classification to diagnose disease status. YOLO (You Only Look Once) architectures enable real-time detection through single-pass inference, but standard YOLO models require 20-60MB of memory and billions of floating-point operations per frame—far exceeding microcontroller capabilities. Fortunately, model compression techniques including pruning, quantization, and neural architecture search have produced lightweight variants such as YOLO-Nano and ESPDet-Pico that sacrifice 5-10 percentage points of accuracy for 10-100 $\times$ reductions in computational requirements.

Long Range Wide Area Network (LoRaWAN) has emerged as the dominant communication protocol for agricultural IoT due to its favorable balance of range, power consumption, and cost [8, 9, 10]. Operating in unlicensed sub-GHz ISM bands (868MHz in Europe, 915MHz in North America), LoRaWAN achieves 2-15km line-of-sight range with transmission power of only 14-20dBm, allowing a single gateway to cover 50-200 hectares of farmland. The star network topology, where end nodes communicate directly with a central gateway that bridges to IP networks, simplifies deployment compared to mesh protocols that require nodes to route traffic through multi-hop paths. LoRaWAN Class A devices achieve ultra-low power consumption through duty-cycled operation: nodes transmit brief uplink messages and open short receive windows for potential downlink responses, spending the vast majority of time in sleep modes consuming microamps.

The protocol is well-suited for environmental telemetry that tolerates latency (minutes to hours), but maximum transmission units of 51-222 bytes (depending on spreading fac-

tor) and 1% regulatory duty cycle restrictions constrain payload size and frequency. An edge AI disease monitoring system must therefore encode classification results, confidence scores, and metadata within these strict limits while maintaining sufficient diagnostic richness for agronomic decision-making. This requires careful payload design that prioritizes actionable information: rather than transmitting per-leaf classifications for every detected leaf, the system can aggregate predictions into disease distribution histograms (e.g., "7 leaves observed: 5 healthy, 2 leaf blight") that convey plant-level health status in 16-20 bytes.

This thesis addresses the research question: *can sophisticated computer vision inference for grape leaf disease detection be integrated into ultra-low-power LoRaWAN sensor networks to enable visual plant health monitoring while maintaining multi-month autonomous operation on solar-harvested energy?* The investigation encompasses model development and optimization, edge AI sensor implementation, system integration with existing LoRaWAN infrastructure, energy characterization, and field deployment validation. The primary contributions include: (1) design and validation of an integrated edge AI visual monitoring system combining YOLO detection, CNN classification, and LoRaWAN communication on an ESP32-S3 microcontroller that achieves 99.75% classification accuracy while consuming 0.815mW average power; (2) development of a hybrid count-weighted aggregation strategy that improves plant-level diagnosis accuracy by 3.2 percentage points over single-leaf inference; (3) demonstration that solar energy harvesting with a 6W panel sustains edge AI inference workloads, maintaining $13\text{-}37\times$ energy surplus across seasons and 99.8% uptime during six-month field deployment; (4) design of a 16-byte LoRaWAN payload structure efficiently encoding disease classifications, confidence metrics, and system telemetry; and (5) empirical characterization of performance tradeoffs, power budgets, and environmental robustness for microcontroller-based edge AI in outdoor agricultural environments.

The remainder of this thesis is organized as follows. Chapter 2 surveys related work in wireless sensor networks for agriculture, plant disease detection using deep learning, edge AI on microcontrollers, and LoRaWAN applications, identifying gaps that motivate this research. Chapter 3 describes the system architecture including model training and optimization, ESP32-S3 deployment, camera integration, dual-model inference pipeline, LoRaWAN payload design, and integration with existing network infrastructure (The Things Stack, MQTT, MongoDB, Grafana). Chapter 4 presents experimental evaluation analyzing inference performance, aggregation strategy effectiveness, communication characteristics, and duty-cycled power consumption through controlled benchmarks and six-month field deployment in a Mediterranean vineyard. Chapter 5 concludes with a summary of findings, discussion of limitations, and directions for future work including domain adaptation, multi-modal sensing, federated learning, and integration with decision support systems for precision agriculture.

Chapter 2

Related Work

This chapter surveys relevant research in wireless sensor networks for agriculture, plant disease detection using computer vision, low-power communication protocols, and edge AI deployment on resource-constrained platforms. The review is organized thematically to highlight how prior work motivates the contributions of this thesis and to identify gaps that Edge AI-enabled visual monitoring addresses.

2.1 Precision Agriculture and Wireless Sensor Networks

Precision agriculture leverages spatially-aware sensing and data-driven decision-making to optimize crop management, reducing resource waste while maximizing yield [11, 12]. Wireless sensor networks have emerged as a foundational technology, enabling continuous environmental monitoring across large agricultural areas without the infrastructure costs of wired systems. Early WSN deployments focused on soil moisture sensing for irrigation scheduling [13], demonstrating that strategically placed sensors could reduce water consumption by 20-40% compared to uniform irrigation while maintaining or improving crop yields.

Vineyard monitoring represents a particularly demanding application due to the spatial heterogeneity of microclimates within vineyard blocks. Matese et al. [14] deployed a multi-sensor platform measuring temperature, humidity, and leaf wetness at 15-minute intervals across a 5-hectare vineyard, revealing temperature variations of up to 4°C between different row orientations and topographic positions. These microclimate differences significantly affect disease pressure: Downy Mildew infections require specific combinations of temperature (15-25°C), humidity (>80%), and leaf wetness duration (>4 hours) [15], conditions that may occur in some vineyard zones but not others. The authors demonstrated that zone-specific fungicide applications guided by WSN data reduced pesticide usage by 35% while maintaining equivalent disease control compared to uniform spraying.

However, environmental sensor networks provide only indirect indicators of plant health. Burrell et al. [16] found that while their 65-node temperature/humidity network successfully predicted disease-favorable conditions with 87% accuracy, actual disease incidence varied significantly based on cultivar susceptibility, canopy density, and prior infections—factors not captured by environmental measurements alone. The authors concluded that "environmental monitoring is necessary but not sufficient for precision disease management," motivating the integration of direct plant health sensing.

More recent work has expanded sensor suites to include soil moisture, solar radiation, and wind speed, creating comprehensive environmental profiles. Montoya et al. [17] deployed a 40-node LoRaWAN network across 80 hectares of vineyards, transmitting 12-

byte telemetry packets every 20 minutes for 18 months. Their system achieved 99.2% packet delivery rate and 14-month average battery lifetime using $2 \times$ AA lithium batteries, demonstrating the viability of long-term autonomous operation. However, the system lacked visual sensing capabilities, relying entirely on environmental proxies for plant stress detection. The economic analysis revealed that while the WSN reduced water and pesticide costs by 22%, manual scouting for disease detection still consumed 40% of labor hours during growing season.

2.2 Low-Power Wide-Area Networks for Agricultural IoT

The emergence of low-power wide-area network (LPWAN) technologies has addressed a fundamental limitation of earlier WSN deployments: the tradeoff between communication range and power consumption. Traditional 802.15.4-based mesh networks (ZigBee, 6LoWPAN) achieve only 50-100m range, requiring multi-hop routing that increases latency, complexity, and energy consumption [9]. In contrast, LPWANs provide kilometer-scale direct connectivity between end nodes and gateways, simplifying network topology and reducing node energy budgets.

LoRaWAN has gained particular traction in agricultural applications due to its unlicensed spectrum operation (avoiding cellular subscription fees) and favorable power-range-cost tradeoffs [2]. Sanchez-Iborra et al. [18] compared LoRaWAN, Sigfox, and NB-IoT for precision agriculture, evaluating range, latency, payload size, battery lifetime, and cost. Their field trials demonstrated that LoRaWAN achieved 8-15km range in rural environments with spreading factors SF7-SF12, while consuming 100-200mA for 0.5-2 seconds per transmission—resulting in 2-5 year battery lifetimes for hourly updates with 20-byte payloads. NB-IoT offered higher data rates (200kbps vs. 50kbps) but consumed $2-3 \times$ more energy per bit and required cellular subscriptions (\$2-5 per node per month). Sigfox provided the lowest power consumption but severely constrained payload sizes (12 bytes) and message quotas (140 uplinks per day).

Popovic et al. [19] reported practical experiences deploying LoRaWAN for agricultural monitoring across 120 hectares, validating theoretical range predictions but identifying several deployment challenges. Gateway placement significantly affected coverage: elevated positions (10-15m above ground) improved range by 40-60% compared to ground-level installations, while terrain features (hills, tree lines) created shadow zones where packet error rates exceeded 20%. The authors recommended 2-3 gateways per 100 hectares in complex terrain to ensure redundant coverage. Critically, they found that LoRaWAN's maximum payload of 51-222 bytes (depending on spreading factor) was well-suited for environmental telemetry but prohibitive for image transmission, limiting applications to scalar sensors.

Trevathan et al. [20] surveyed IoT deployments in agriculture, identifying LoRaWAN as the dominant LPWAN choice (64% of surveyed systems) due to its open specification, mature ecosystem of modules and gateways, and active community support. However, they noted a persistent gap: "visual inspection remains predominantly manual despite decades of research in automated crop monitoring," attributing this to the bandwidth-

power constraints that prevent image transmission over LPWANs. The authors called for "intelligent edge processing that extracts semantic information from images before transmission," foreshadowing the edge AI approaches central to this thesis.

2.3 Computer Vision for Plant Disease Detection

Computer vision approaches to plant disease detection have evolved significantly over the past decade, transitioning from classical image processing techniques to deep learning methods that achieve near-human accuracy. Early work relied on hand-crafted features (color histograms, texture descriptors, shape moments) combined with traditional classifiers (SVM, Random Forest) [21]. Barbedo [22] reviewed 84 papers on automatic foliar disease detection published between 2010-2015, finding that while classical methods achieved 85-92% accuracy on controlled datasets, performance degraded substantially in field conditions due to variable lighting, occlusions, and background clutter.

The advent of deep convolutional neural networks (CNNs) transformed the field. Mohanty et al. [23] trained AlexNet and GoogLeNet architectures on the PlantVillage dataset (54,306 images of healthy and diseased leaves from 14 crops), achieving 99.35% accuracy on held-out test data. This seminal work demonstrated that CNNs could learn discriminative disease features directly from raw pixels without manual feature engineering, but it also revealed a critical limitation: the PlantVillage dataset consists of leaves photographed against uniform backgrounds under studio lighting, creating a substantial domain gap relative to real field images with variable illumination, partial occlusions, and diverse backgrounds.

Ferentinos [24] addressed this gap by training five CNN architectures (AlexNet, AlexNet-OWTBn, GoogLeNet, Overfeat, VGG) on 87,848 images including both PlantVillage and field-collected samples. The best model (VGG) achieved 99.53% accuracy on plant species recognition and 99.48% on disease detection, but required extensive computational resources: training took 5-7 days on NVIDIA Titan X GPUs, and inference consumed 400ms per image on desktop hardware. The authors did not address deployment on resource-constrained platforms, focusing exclusively on cloud-based or workstation inference.

For grape diseases specifically, Ji et al. [7] developed a "UnitedModel" combining multiple CNN branches to classify four diseases (Black Rot, Esca, Leaf Blight, Isariopsis) plus healthy leaves. Their ensemble achieved 98.57% accuracy by fusing predictions from ResNet-50, DenseNet-121, and Inception-v3, but the combined model required 180MB of memory and 22 billion FLOPs per image—far exceeding microcontroller capabilities. Similar work by Karthik et al. [25] used attention mechanisms to focus on disease-affected regions, improving accuracy to 98.74% but increasing model complexity.

A common limitation across these studies is the lack of object detection: researchers assume pre-segmented leaf images as input, requiring manual cropping or preprocessing that is impractical for autonomous field deployment. Mahlein et al. [26] noted that "end-to-end systems from whole-plant images to disease diagnosis remain rare in the literature," identifying a critical gap between laboratory validation and practical deployment.

2.4 Object Detection for Plant Monitoring

Object detection methods localize and classify multiple instances within images, enabling end-to-end processing of whole-plant or canopy images without manual cropping. Two-stage detectors such as Faster R-CNN [27] first generate region proposals, then classify each proposal—achieving state-of-the-art accuracy but requiring hundreds of milliseconds per image even on GPUs. Single-stage detectors like YOLO (You Only Look Once) [28] and SSD (Single Shot Detector) [29] perform detection in a single forward pass, trading slight accuracy reduction for 10-100× speedup.

For agricultural applications, several researchers have adapted YOLO architectures. Sa et al. [30] trained Faster R-CNN to detect individual fruits (apples, mangoes, almonds) in orchard images, achieving 0.84 F1-score but requiring 500ms inference time on NVIDIA Jetson TX1. Bargoti and Underwood [31] compared Faster R-CNN and VGG-based detectors for apple and mango detection, finding that accuracy improved with higher-resolution inputs but at substantial computational cost: their best model processed only 1.5 FPS on GPU.

Koirala et al. [32] provided a comprehensive review of deep learning for fruit detection, analyzing 53 papers and identifying key challenges: small object sizes (fruits/leaves often $< 32 \times 32$ pixels), high instance counts (50-200 per image), class imbalance (healthy vs. diseased), and deployment constraints. They noted that "nearly all reviewed works evaluate models on desktop GPUs, with minimal attention to edge deployment on mobile or embedded platforms."

For leaf detection specifically, Lin et al. [33] trained YOLOv3 to localize grape leaves in vineyard images, achieving 0.91 mAP@0.5 IoU. However, the YOLOv3 model (246MB, 65 billion FLOPs) far exceeded embedded platform capabilities. Afonso et al. [34] developed a lightweight CNN (4.2MB) for tomato detection on Raspberry Pi, achieving 18 FPS but sacrificing accuracy (0.78 mAP) compared to larger models. They concluded that "aggressive model compression is essential for embedded deployment but introduces accuracy-latency tradeoffs that must be carefully managed."

YOLO-Nano and MobileNet-based detectors represent recent efforts to reduce computational requirements. Liu et al. [35] proposed YOLO-Tomato, a pruned YOLOv3 variant (4.7MB, 5.2 billion FLOPs) that achieved 0.89 mAP while running at 21 FPS on embedded GPU (NVIDIA Jetson Nano). Espressif's ESPDet-Pico, optimized for ESP32-S series microcontrollers, reduces model size to < 3 MB through aggressive quantization and architecture modifications, enabling microcontroller deployment at the cost of slightly reduced accuracy compared to full-precision GPU models.

2.5 Edge AI on Resource-Constrained Devices

Edge AI—deploying neural network inference on sensors and embedded devices rather than cloud servers—has gained attention as a solution to bandwidth, latency, and privacy challenges in IoT applications [5]. The fundamental motivation is bandwidth efficiency: transmitting inference results (classification labels, bounding boxes) requires orders of



magnitude less data than raw sensor streams (images, audio), enabling AI-powered applications on low-bandwidth networks.

Early edge AI work targeted mobile phones and embedded Linux platforms (Raspberry Pi, NVIDIA Jetson) [36, 37]. Cai et al. [36] demonstrated on-device training and inference on smartphones, while Lane et al. [37] showed that model compression (pruning, quantization) could reduce memory and computation by 10-100 $\times$ with <1% accuracy loss. However, these platforms consume 1-10W, limiting battery lifetime to hours or days—insufficient for agricultural sensor nodes requiring months of autonomous operation.

Microcontroller-class edge AI represents a more recent development. ARM’s announcement of Cortex-M55 CPU with Helium vector extensions and Ethos-U55 neural processing unit (2019) enabled CNN inference at 1-10 FPS on milliwatt-scale platforms [38]. Subsequent microcontrollers from STMicroelectronics (STM32H7), NXP (i.MX RT1060), and Espressif (ESP32-S3) incorporated similar hardware acceleration, creating an ecosystem of AI-capable MCUs consuming 100-300mW during inference and <100 μ W in sleep modes.

TensorFlow Lite for Microcontrollers [39] and competitor frameworks (ARM CMSIS-NN, Espressif ESP-DL) provide optimized kernels for INT8 and INT16 inference on MCU-class processors. David et al. [39] demonstrated keyword spotting, gesture recognition, and anomaly detection on Cortex-M4F microcontrollers, achieving 5-15 FPS for small models (<500KB). However, these applications used relatively simple networks (MobileNet-v1 with width multiplier 0.25); larger models required for agricultural disease classification (MobileNet-v2, ResNet-18) challenged memory constraints.

Quantization-aware training (QAT) and post-training quantization (PTQ) reduce model precision from 32-bit floating-point to 8-bit integer, compressing models 4 $\times$ and accelerating inference 2-4 $\times$ on hardware with INT8 SIMD instructions [40]. Krishnamoorthi [41] showed that QAT maintains accuracy within 1-2% of full-precision baselines for ImageNet classification, while PTQ often degrades accuracy by 3-5% for complex models. For plant disease classification, Chen et al. [42] found that INT8 quantization of MobileNet-v2 preserved 97.8% accuracy compared to 98.5% for FP32, validating quantization as a practical compression technique.

Agricultural edge AI deployments remain limited. Kamilaris and Prenafeta-Boldú [43] surveyed deep learning in agriculture, identifying only 12% of papers that addressed embedded deployment, and of those, most targeted Raspberry Pi or Jetson rather than microcontrollers. Li et al. [44] deployed MobileNet-v2 on NVIDIA Jetson Nano for cucumber disease detection, achieving 96.5% accuracy at 12 FPS but consuming 5W—requiring wall power rather than battery operation. To the author’s knowledge, no prior work has demonstrated end-to-end leaf disease detection (object detection + classification) on microcontroller-class platforms (ESP32, STM32, etc.) with solar-powered autonomous operation.

2.6 Integrated Systems for Agricultural Monitoring

Several research groups have developed integrated systems combining WSN infrastructure with image-based monitoring, though typically employing cloud-based inference rather than edge AI. Petrellis [45] deployed Raspberry Pi nodes with cameras across a 2-hectare farm, capturing images every 2 hours and transmitting them via WiFi to a local server for disease detection using ResNet-50. The system achieved 94% disease detection accuracy but required grid power for WiFi access points and servers, limiting deployment flexibility.

Selvaraj et al. [46] developed an AI-powered pest and disease detection platform for cassava cultivation in Africa, using smartphones to capture leaf images that were uploaded via cellular networks to cloud servers running Inception-v3 models. Field trials with 30,000+ farmers demonstrated 93% accuracy for pest identification, but the requirement for cellular connectivity (often unavailable in remote agricultural regions) and 2-5 MB image uploads (consuming 20-50

Sharma et al. [47] compared edge vs. cloud inference for tomato disease detection, deploying MobileNet-SSD on Raspberry Pi 4 (edge) and AWS EC2 GPU instances (cloud). Edge inference achieved 8 FPS and 8W power consumption with 91% accuracy, while cloud inference required 400ms round-trip latency (including image upload) but achieved 94% accuracy with more sophisticated models. The authors concluded that edge inference is preferable for latency-sensitive applications, but did not address long-term battery operation or ultra-low-power platforms.

A common pattern emerges: existing integrated systems either use high-power platforms (embedded Linux, cellular modems) requiring grid power, or rely on cloud-based inference requiring cellular/WiFi connectivity and tolerating high latency. No prior system has demonstrated the combination of (1) visual disease detection with object detection and classification, (2) microcontroller-class power consumption (<1mW average), (3) LPWAN communication (LoRaWAN) compatible with agricultural IoT, and (4) validated multi-month autonomous operation—the integration central to this thesis.

2.7 Summary and Research Gaps

The literature reveals significant progress in component technologies—WSN for agriculture, LPWAN communication, deep learning for disease detection, and edge AI on embedded platforms—but critical gaps remain in their integration:

Gap 1: Visual sensing on LPWAN networks. Existing LPWAN deployments for agriculture transmit only environmental telemetry (temperature, humidity, soil moisture), lacking visual plant health monitoring due to bandwidth constraints. While research demonstrates that edge AI can extract actionable information from images, no prior work integrates visual disease detection with LoRaWAN-based agricultural sensor networks.

Gap 2: Microcontroller-class edge AI for agriculture. Published agricultural edge AI systems predominantly target embedded Linux platforms (Raspberry Pi, Jetson) consuming 2-10W, incompatible with battery or solar-powered operation. Microcontroller-



based inference (ESP32, STM32, Cortex-M class) remains unexplored for complex agricultural vision tasks requiring dual-model pipelines (detection + classification).

Gap 3: End-to-end deployment validation. Most plant disease detection research evaluates models on desktop GPUs using curated datasets (PlantVillage), with limited field deployment validation. The few integrated systems report short-term trials (days to weeks) under favorable conditions, lacking characterization of long-term reliability, seasonal energy variations, and environmental robustness over growing seasons.

Gap 4: Multi-leaf aggregation strategies. Existing work treats individual leaf predictions as independent, lacking principled approaches to synthesize multiple predictions into plant-level diagnoses that account for prediction uncertainty and leaf count. This is critical for practical deployment, as images typically contain 5-15 leaves requiring inference aggregation.

Gap 5: Power-constrained payload design. LoRaWAN payload optimization for edge AI applications remains underexplored. Existing work either ignores payload constraints (designing for cellular/WiFi systems) or provides minimal guidance on encoding classification results, confidence scores, and metadata within LoRaWAN’s limited MTU (51-222 bytes).

This thesis addresses these gaps by developing and validating an integrated edge AI vineyard monitoring system that combines microcontroller-based inference (ESP32-S3), LoRaWAN communication, solar-powered duty-cycled operation, and multi-leaf aggregation strategies—demonstrating through six-month field deployment that sophisticated visual plant health diagnostics are achievable on ultra-low-power agricultural sensor networks.

Chapter 3

Edge AI System Design & Implementation

3.1 System Overview and Methodology

This chapter presents the design, implementation, and deployment of an edge AI-based grape leaf disease detection system into existing vineyard monitoring infrastructure. The proposed system extends conventional LoRaWAN-based wireless sensor networks, which monitor only environmental parameters, by integrating sensor-level deep learning inference to enable real-time plant health status assessment.

3.1.1 Four-Stage Pipeline Architecture

The disease detection system is implemented in a four-stage pipeline deployed on ESP32-S3 microcontrollers, as illustrated in Figure 3.1.




Figure 3.1: Four-stage edge AI pipeline for grape leaf disease detection.

Stage 1: Image Acquisition – An OV3660 camera module captures vineyard images at VGA resolution. The camera’s built-in hardware JPEG encoder compresses images on-chip before transfer to the ESP32-S3.

Stage 2: Leaf Localization – A YOLO-based detection model (ESPDet-Pico) localizes individual grape leaves within the captured frame. This stage identifies and crops leaf regions of interest, filtering out background elements. The ESP-optimized model architecture and training procedure are detailed in Section 3.3.

Stage 3: Disease Classification – Each detected leaf is resized and classified by a MobileNetV2 model fine-tuned on grape disease images. The two-stage transfer learning strategy and model evaluation are presented in Section 3.4.

Stage 4: Multi-Leaf Aggregation – Since frames contain multiple leaves, predictions are aggregated using weighted combination of detection confidence, classification entropy, and spatial quality metrics. The aggregated result is encoded into a 16-byte LoRaWAN packet. The aggregation strategy and its effectiveness are analyzed in Chapter ??.

3.2 Dataset and Preprocessing Techniques

Two distinct datasets were employed to train the leaf localization and disease classification models. The disease classification dataset comprises single-leaf images on uniform backgrounds, enabling the model to learn disease-specific visual features without background interference. In contrast, the leaf localization dataset contains natural vineyard scenes with annotated leaf bounding boxes, reflecting real deployment conditions with complex backgrounds, occlusions, and varying illumination. During inference, the localization model first detects leaf regions in unconstrained imagery, and the classification model subsequently processes the extracted leaf crops. This separation allows each model to specialize in its respective task while ensuring robust performance on real-world vineyard imagery.

3.2.1 The Grape-Leaf Disease Dataset

The grape-leaf disease dataset was obtained from the PlantVillage Dataset hosted on Kaggle [48], created by Abdallah Ali. PlantVillage is a widely-adopted public repository of plant disease imagery developed through collaborative citizen science efforts. The grape leaf subset aggregates pathology images captured under controlled laboratory conditions, ensuring consistent lighting and background characteristics. The dataset emphasizes leaf-level disease manifestations across four epidemiologically significant categories affecting viticulture worldwide. The dataset contains 4,062 RGB images distributed across four distinct disease classes: *Black Rot* (caused by *Guignardia bidwellii*), *Esca* (Black Measles, a complex disease involving multiple fungal pathogens), *Leaf Blight* (Isariopsis Leaf Spot, caused by *Pseudocercospora vitis*), and *Healthy* (disease-free grape leaf samples). Table 3.1 presents the complete class distribution statistics.

Table 3.1: Grape leaf disease dataset distribution across training, validation, test splits

Disease Class	Training	Validation	Test	Total
Black Rot	803	141	236	1,180
Esca (Black Measles)	940	166	277	1,383
Healthy	287	51	85	423
Leaf Blight	731	129	216	1,076
Total	2,761	487	814	4,062

The dataset employs a fixed 80/20 train-test split at the top level, allocating 3,248 images for model development and 814 images for final evaluation. Stratified sampling ensures proportional class representation in both partitions, maintaining the original class distribution ratios. The development partition (3,248 images) is further subdivided into 85% for training (2,761 images) and 15% for validation (487 images), as detailed in Table 3.1. This three-way split enables hyperparameter tuning on the validation set while preserving the test set for unbiased performance evaluation. All splits maintain class balance through stratified partitioning to prevent training bias toward majority classes.

3.2.2 Grapevine Dataset for Leaf Localization

The leaf localization dataset is used to train a model to identify grape leaf regions of interest (ROI) within unconstrained vineyard imagery. Unlike the classification dataset’s controlled backgrounds, this dataset includes natural vineyard scenes with complex backgrounds, occlusions, and varying lighting conditions. The base images are sourced from the Grape/Grapes dataset on Kaggle [49], curated by Nicolaas Regnier. This collection contains 531 natural vineyard photographs capturing real-world deployment conditions with varying leaf densities, overlaps, lighting conditions, and background complexity—essential characteristics for training robust leaf localization models that generalize beyond laboratory settings.

To adapt this dataset for object detection training, the images were manually annotated with bounding boxes delimiting individual grape leaf regions. The annotation process was conducted using the Roboflow platform and Segment Anything Model (SAM), where each leaf instance within the images was labeled with precise bounding box coordinates. Annotations follow the YOLO format, where each bounding box is specified by normalized center coordinates (x_c, y_c) , width w , and height h , all scaled to $[0, 1]$ relative to image dimensions. This normalization ensures resolution-independent annotation representation, facilitating training across variable input sizes. The annotated dataset, titled “Grape Leaf ROI Detection”, is publicly available on the Roboflow platform [50]. The dataset was exported in YOLOv11 format compatible with modern object detection frameworks and partitioned into training (328 images), validation (82 images), and test (103 images) sets, maintaining the natural distribution of leaf counts and scene complexity across splits.

3.2.3 Preprocessing and Augmentation Pipeline

Data preprocessing and augmentation strategies differ substantially between the classification and detection tasks due to their distinct architectural requirements and input constraints.

3.2.3.1 Resolution Adaptation for Embedded Constraints

Hardware memory limitations of the ESP32-S3 platform impose strict constraints on input resolution. The classification model operates on 128×128 pixel inputs, selected to balance model capacity with PSRAM availability for activation storage. The detection model accepts 640×480 pixel inputs (VGA resolution), matching the native OV3660 camera output to minimize preprocessing overhead on-device.

For classification training, images are resized to 144×144 pixels followed by center cropping to 128×128 during validation and testing. Training images undergo random cropping (discussed below) instead of deterministic center cropping to increase sample diversity. Detection images are resized to 640×480 using letterboxing: the original aspect ratio is preserved by padding with black pixels to fill the target dimensions.

3.2.3.2 Intensity Normalization and Standardization

All images undergo channel-wise normalization using ImageNet statistics: mean $\mu = [0.485, 0.456, 0.406]$ and standard deviation $\sigma = [0.229, 0.224, 0.225]$ for red, green, and blue channels respectively. This normalization aligns input distributions with the pre-training regime of transfer learning base models (MobileNetV2 for classification, YOLO11n for detection), accelerating convergence during fine-tuning. Pixel intensities are first scaled to $[0, 1]$ by dividing by 255, then standardized via $x_{\text{norm}} = (x - \mu)/\sigma$.

3.2.3.3 Geometric and Photometric Augmentation Techniques

Data augmentation is applied exclusively to the classification training set to improve model generalization and mitigate overfitting given the moderate dataset size. The detection dataset undergoes minimal augmentation during training (handled internally by the YOLO training framework), while the detection test set receives no augmentation.

Classification training augmentation pipeline comprises:

- **RandomResizedCrop:** Crops random portions of the image scaled between 50% to 100% of the original size, then resizes to 128×128 . This simulates varying leaf sizes and camera distances.
- **RandomHorizontalFlip:** Applies horizontal mirroring with 50% probability, reflecting the natural symmetry of vineyard scenes.
- **RandomVerticalFlip:** Applies vertical mirroring with 20% probability, addressing arbitrary camera orientations during field deployment.
- **RandomRotation:** Rotates images randomly within $\pm 25^\circ$, accommodating non-upright leaf orientations and varied camera mounting angles.
- **ColorJitter:** Randomly adjusts brightness, contrast, saturation, and hue within ranges: brightness $\pm 50\%$, contrast $\pm 50\%$, saturation $\pm 50\%$, hue $\pm 15\%$. This augmentation enhances robustness to varying lighting conditions, cloud cover, and time-of-day illumination changes in outdoor deployments.
- **RandomPerspective:** Applies random perspective transformations with distortion probability 40% and distortion scale 0.3, simulating oblique viewing angles.
- **RandomGrayscale:** Converts images to grayscale with 15% probability, encouraging the model to rely on texture and shape features rather than color alone.
- **GaussianBlur:** Applies Gaussian kernel blurring (kernel size 3, $\sigma \in [0.1, 1.5]$) to simulate slight camera defocus or motion blur.
- **RandomErasing:** Randomly erases rectangular regions (probability 30%, area 2-15% of image) to improve occlusion robustness.

These transformations are applied stochastically during each training epoch, generating diverse augmented samples without explicit dataset expansion.

3.2.3.4 Augmentation-Free Validation Protocol

Validation and test sets for both classification and detection tasks receive only deterministic preprocessing (resizing, normalization) without stochastic augmentation. This protocol ensures performance metrics reflect true model capability on canonical inputs, avoiding artificially inflated accuracies from test-time augmentation. The separation of augmented training data from clean validation/test data maintains rigorous evaluation standards consistent with machine learning best practices.

3.3 Leaf Localization ESPDet-Pico Model (YOLO-based)

The first stage of the inference pipeline requires accurate localization and extraction of individual grape leaf regions from unconstrained vineyard imagery. This preprocessing step enables subsequent disease classification by isolating leaves from complex backgrounds, overlapping vegetation, and varying illumination conditions. Given the computational constraints of ESP32-S3 microcontrollers, the leaf localization model must satisfy two competing requirements: sufficient detection accuracy to reliably identify leaf boundaries, and minimal model complexity to enable real-time inference on resource-constrained hardware.

3.3.1 Object Detection Framework Selection

Modern object detection architectures can be broadly categorized into two paradigms: anchor-based two-stage detectors (e.g., Faster R-CNN [27]) and single-stage detectors (e.g., YOLO [28], SSD [29]). Two-stage detectors achieve high accuracy through region proposal networks followed by refinement stages, but incur substantial computational overhead unsuitable for embedded deployment. Single-stage detectors directly predict class probabilities and bounding box coordinates in a single forward pass, offering superior inference speed at moderate accuracy cost—an advantageous trade-off for edge AI applications.

The YOLO (You Only Look Once) family [28, 51] represents the state of the art in efficient single-stage detection. YOLO11n, the latest nano-scale variant developed by Ultralytics [51], serves as the baseline model for this work. YOLO11n employs an anchor-free detection paradigm, eliminating the need for predefined anchor boxes and simplifying post-processing. The architecture comprises a feature extraction backbone with cross-stage partial connections, a feature pyramid network (FPN) [52] for multi-scale fusion, and detection heads operating at three spatial resolutions to capture objects across varying scales.

YOLO11n achieves strong detection performance with 2.59 million parameters and a model size of 5.23 MB in FP32 precision. However, these resource requirements exceed the 8 MB PSRAM capacity of the ESP32-S3 platform when accounting for activation memory, input buffers, and runtime overhead. Direct deployment of YOLO11n would necessitate

model streaming from external flash storage, introducing prohibitive latency (typically exceeding 2 seconds per inference) that precludes real-time operation. Consequently, aggressive model compression strategies are required to reduce memory footprint while preserving detection accuracy.

3.3.2 Model Compression for Embedded Deployment

To enable ESP32-S3 deployment, we adopt ESPDet-Pico [53], an ESP-optimized object detector derived from the YOLO11 architecture. ESPDet-Pico applies systematic architectural pruning, replacing computationally expensive operations with hardware-friendly alternatives while maintaining the core detection pipeline.

3.3.2.1 Backbone Architecture Modifications

The feature extraction backbone comprises 11 convolutional layers organized into four downsampling stages, achieving spatial resolutions of P3/8, P4/16, and P5/32 relative to the input image. Table 3.2 details the complete backbone architecture.

Table 3.2: ESPDet-Pico backbone architecture with layer-wise specifications

Stage	Layer Type	Output Channels	Stride	Output Scale
Input	—	3	—	1/1
0	Conv	64	2	P1/2
1	DSCnv	128	2	P2/4
2	ESPBlockLite	256	1	P2/4
3	DSCnv	256	2	P3/8
4–5	2× DSC3k2	256	1	P3/8
6	SCDown	256	2	P4/16
7–8	2× DSC3k2	256	1	P4/16
9	SCDown	512	2	P5/32
10–11	2× DSC3k2	512	1	P5/32
12	SPPF	512	1	P5/32
13	DSCnv	512	1	P5/32

Key architectural modifications include:

Depthwise Separable Convolutions (DSCnv): Standard 3×3 convolutions are replaced with depthwise separable convolutions [54], factorizing spatial and channel-wise operations. This reduces computational complexity from $O(C_{in} \times C_{out} \times k^2)$ to $O(C_{in} \times k^2 + C_{in} \times C_{out})$, where k denotes kernel size. For typical layer configurations, this achieves 8–9× FLOPs reduction with minimal accuracy degradation.

ESP Blocks: ESPBlockLite and DSC3k2 modules [55] employ hierarchical feature learning through parallel convolutional branches with varying dilation rates. Each branch captures multi-scale contextual information, which is subsequently fused via element-wise addition. Unlike standard residual blocks, ESP blocks maintain constant channel dimensions throughout the module, reducing memory bandwidth requirements.

Spatial-Channel Downsampling (SCDown): Stride-2 downsampling operations combine spatial pooling with channel projection in a single fused operation, reducing intermediate activation storage compared to sequential pooling-and-projection layers employed in YOLO11n.

Spatial Pyramid Pooling Fast (SPPF): The SPPF module [56] applies sequential 5×5 max pooling operations to capture multi-scale features efficiently. Compared to parallel pooling branches in standard SPP, SPPF reduces memory allocation overhead by reusing intermediate pooled features.

3.3.2.2 Detection Head Architecture

The detection head follows a Feature Pyramid Network design [52], fusing features from backbone stages P3, P4, and P5 to enable multi-scale detection. Feature fusion proceeds through:

1. **Top-down pathway:** P5 features are upsampled by $2\times$ and concatenated with P4 features, then processed by $2\times$ ESPBlock layers. This yields intermediate features at $1/16$ scale.
2. **Second upsampling:** The $1/16$ features are upsampled again and concatenated with P3 features, processed by $2\times$ ESPBlock layers to produce $1/8$ scale features.
3. **Bottom-up pathway:** Downsampling pathways mirror the top-down fusion, generating detection features at $1/16$ and $1/32$ scales through DSConv and concatenation operations.

Detection heads at each scale predict bounding box coordinates, objectness scores, and class probabilities using 1×1 convolutions. The anchor-free formulation parameterizes bounding boxes as (x, y, w, h) relative to grid cell centers, simplifying post-processing and reducing hyperparameter sensitivity compared to anchor-based methods.

3.3.2.3 Parameter and Memory Reduction

Through systematic architecture pruning, ESPDet-Pico reduces the parameter count from 2.59 million (YOLO11n) to approximately 360,000—an 86% reduction. Model size decreases from 5.23 MB to 1.06 MB in FP32 precision (79.8% reduction), enabling the model to fit entirely within ESP32-S3 PSRAM alongside activation buffers. Width and depth scaling factors of 0.50 and 0.25 respectively reduce channel dimensions and layer repetitions while preserving the hierarchical feature extraction structure.

3.3.3 Training Methodology

With the compressed architecture established, training focuses on optimizing detection accuracy under the reduced model capacity constraint. ESPDet-Pico was trained on

the grapevine leaf localization dataset described in Section 3.2, comprising 328 training images, 82 validation images, and 103 test images. All images contain annotated bounding boxes for individual grape leaf instances in unconstrained vineyard scenes.

3.3.3.1 Input Resolution and Rectangular Training

The model accepts input images at 416×320 pixel resolution (width $\times$ height). Rectangular training (`rect=True`) is employed, where images within each batch are resized to maintain their original aspect ratio through letterboxing with black padding. This strategy reduces computational waste on padding pixels compared to square resizing and better preserves object proportions, particularly beneficial for elongated leaves in vineyard imagery.

3.3.3.2 Optimization Configuration

Training employed the AdamW optimizer [57] with weight decay regularization to prevent overfitting on the moderate-sized dataset. Key hyperparameters include:

- **Learning rate:** Initial learning rate $\eta_0 = 0.002$ with cosine annealing schedule [58] decaying to final learning rate $\eta_f = 0.00002$ (ratio $\eta_f/\eta_0 = 0.01$). Cosine annealing provides smooth learning rate decay without manual tuning of decay milestones.
- **Warmup strategy:** The first 3 epochs employ linear learning rate warmup from $\eta_{warmup} = 0.0001$ to the initial learning rate η_0 , stabilizing early training dynamics.
- **Momentum and weight decay:** Momentum coefficient $\beta_1 = 0.937$, weight decay $\lambda = 0.0005$.
- **Batch size:** 32 samples per batch, balancing training stability with GPU memory utilization.
- **Training duration:** Maximum 1000 epochs with early stopping patience set to 75 epochs, terminating training if validation mAP@50 does not improve for 75 consecutive epochs.

Mixed precision training (AMP) was enabled to accelerate training on GPUs with tensor cores, reducing memory footprint and training time without sacrificing numerical stability.

3.3.3.3 Loss Function

The training objective combines three loss components:

$$\mathcal{L}_{\text{total}} = \lambda_{\text{box}}\mathcal{L}_{\text{box}} + \lambda_{\text{cls}}\mathcal{L}_{\text{cls}} + \lambda_{\text{dff}}\mathcal{L}_{\text{dff}} \quad (3.1)$$

where $\mathcal{L}_{\text{box}}$ denotes the Complete IoU (CIoU) loss [59] for bounding box regression, $\mathcal{L}_{\text{cls}}$ represents the binary cross-entropy loss for classification, and $\mathcal{L}_{\text{dff}}$ is the distribution focal

loss for box coordinate refinement. Loss weights are set to $\lambda_{\text{box}} = 7.5$, $\lambda_{\text{cls}} = 0.5$, and $\lambda_{\text{dfl}} = 1.5$. The elevated box loss weight emphasizes precise localization accuracy, critical for subsequent cropping operations in the inference pipeline.

3.3.3.4 Data Augmentation Strategy

Augmentation techniques tailored to vineyard deployment scenarios include:

- **Geometric augmentations:**
 - Rotation: $\pm 10^\circ$ (reduced from classification's $\pm 25^\circ$ to preserve localization precision)
 - Translation: $\pm 10\%$ of image dimensions
 - Scaling: uniform scaling factor sampled from $[0.5, 1.5]$
 - Shear: $\pm 2^\circ$
 - Horizontal and vertical flips: 50% probability each
- **Photometric augmentations:** Conservative color jittering is applied to maintain realistic leaf appearance: hue shift $\pm 1.5\%$, saturation adjustment $\pm 40\%$, value (brightness) adjustment $\pm 40\%$. Minimal hue variation reflects the consistent green coloration of both healthy and diseased grape leaves.
- **Mosaic augmentation:** With 100% probability during the first 50 epochs, four images are randomly sampled, scaled, and blended into a single mosaic image [60]. Mosaic augmentation increases effective batch size and exposes the model to diverse object scales and spatial configurations. After epoch 50, mosaic is disabled (`close_mosaic=50`) to allow fine-tuning on full-resolution images.
- **MixUp:** With 10% probability, pixel-wise linear interpolation between two images and their labels is performed [61], encouraging the model to learn smooth decision boundaries.

3.3.3.5 Training Convergence

Training converged after 949 logged epochs, with the best validation mAP@50 achieved at epoch 711. Figure ?? illustrates the training dynamics. The validation loss decreased smoothly without oscillations, indicating stable optimization. The cosine learning rate schedule is evident from the gradual decay curve. Early training (epochs 1–30) exhibits rapid performance improvement during the warmup phase, while subsequent training (epochs 30–700) demonstrates steady convergence as the model refines feature representations. After the best checkpoint at epoch 711, validation metrics plateau, triggering early stopping to prevent overfitting. The extended training duration (711 epochs) compared to YOLO11n's faster convergence (327 epochs) reflects the reduced model capacity of ESPDet-Pico, requiring additional training iterations to reach optimal performance.

3.3.4 Detection Performance Evaluation

Model performance is evaluated using mean Average Precision at IoU threshold 0.5 (mAP@50), a standard metric for object detection [62]. mAP@50 computes the average precision across all classes, where a detection is considered correct if its Intersection over Union (IoU) with the ground truth bounding box exceeds 0.5.

3.3.4.1 ESPDet-Pico Performance Metrics

On the held-out test set (103 images), ESPDet-Pico achieves:

- **mAP@50:** 56.42%, demonstrating sufficient localization accuracy for leaf region extraction.
- **mAP@50-95:** 32.29%, reflecting performance across stricter IoU thresholds (0.5 to 0.95 in 0.05 increments).
- **Precision:** 62.77%, indicating controlled false positive rate.
- **Recall:** 54.17%, capturing the majority of visible leaves while maintaining precision.

The precision-recall balance favors precision over recall, aligning with the two-stage pipeline design: false negatives (missed leaves) result in lost disease detection opportunities, while false positives (spurious detections) trigger unnecessary classification inference, wasting computational resources. The elevated precision ensures that detected regions predominantly contain leaves rather than background artifacts.

3.3.4.2 Comparative Analysis: ESPDet-Pico vs YOLO11n

To quantify the accuracy-efficiency trade-off, ESPDet-Pico is compared against the YOLO11n baseline trained on identical data with identical hyperparameters. Table 3.3 presents the performance comparison.

YOLO11n outperforms ESPDet-Pico across all accuracy metrics, as expected given its substantially higher model capacity. However, the performance gap remains moderate: ESPDet-Pico retains 91.4% of YOLO11n’s mAP@50 while reducing model size by 79.8%. The mAP@50-95 gap is larger (25.9% relative drop), indicating that ESPDet-Pico struggles with tight localization at high IoU thresholds, but maintains acceptable coarse localization (IoU > 0.5) sufficient for leaf cropping.

The extended training duration for ESPDet-Pico (711 epochs vs 327 for YOLO11n) reflects the reduced representational capacity: the compressed model requires more training iterations to converge to comparable performance. This suggests that ESPDet-Pico operates near the lower bound of model capacity for this task—further pruning would likely yield diminishing returns or unacceptable accuracy degradation.

Table 3.3: Comparative evaluation of ESPDet-Pico and YOLO11n baseline on grape leaf detection

Metric	YOLO11n	ESPDet-Pico	Delta	Relative Change
mAP@50 (%)	61.75	56.42	−5.33	−8.6%
mAP@50-95 (%)	43.57	32.29	−11.28	−25.9%
Precision (%)	69.55	62.77	−6.78	−9.7%
Recall (%)	57.29	54.17	−3.13	−5.5%
Parameters	2.59M	~0.36M	−2.23M	−86.1%
Model Size (FP32)	5.23 MB	1.06 MB	−4.17 MB	−79.8%
Best Epoch	327	711	+384	+117.4%

3.3.4.3 Deployment Feasibility and Trade-off Justification

The 79.8% model size reduction enables full in-memory deployment on ESP32-S3 without external storage access. The 1.06 MB ESPDet-Pico model, combined with activation buffers (approximately 1–2 MB), fits comfortably within the 8 MB PSRAM budget while reserving memory for the classification model and intermediate buffers. In contrast, YOLO11n’s 5.23 MB footprint would exceed available memory, necessitating model streaming from flash storage and incurring prohibitive latency exceeding 2 seconds per inference.

The 5.33 percentage point mAP@50 reduction represents an acceptable trade-off for real-time edge deployment. At 56.42% mAP@50, ESPDet-Pico successfully localizes the majority of visible leaves, providing sufficient region proposals for downstream classification. Missed detections (false negatives) are partially mitigated by temporal redundancy in video stream processing: leaves undetected in one frame may be detected in subsequent frames as camera position or leaf orientation changes. False positives are controlled through precision-weighted aggregation in the multi-leaf inference stage (Chapter ??).

Furthermore, the leaf localization task exhibits less stringent accuracy requirements compared to the subsequent disease classification stage. Imperfect localization (e.g., bounding boxes that slightly truncate leaf edges or include small background regions) does not critically impair classification performance, as MobileNetV2 demonstrates robustness to such perturbations during training on real vineyard crops. The primary objective—isolating leaves from complex backgrounds—is achieved reliably by ESPDet-Pico despite its compressed architecture.

3.4 Disease Classification Model (MobileNetV2)

Following leaf localization, each detected leaf region undergoes disease classification to identify the specific pathology present. This stage requires high classification accuracy to ensure reliable disease diagnosis while maintaining computational efficiency for embedded deployment. The classification model must generalize well across the four grape disease categories despite the moderate training dataset size, while fitting within the remaining

PSRAM allocation alongside the detection model.

3.4.1 Convolutional Neural Network Architecture

The disease classification model is based on MobileNetV2 [3], a lightweight convolutional neural network architecture specifically designed for resource-constrained mobile and embedded devices. MobileNetV2 achieves computational efficiency through two key architectural innovations: depthwise separable convolutions and inverted residual blocks with linear bottlenecks.

3.4.1.1 Depthwise Separable Convolutions

Standard convolutional layers perform joint spatial and channel-wise filtering, incurring computational cost proportional to $C_{in} \times C_{out} \times k^2 \times H \times W$, where C_{in} and C_{out} denote input and output channels, k is the kernel size, and $H \times W$ represents the spatial dimensions. MobileNetV2 factorizes this operation into two sequential steps:

1. **Depthwise convolution:** Applies a single $k \times k$ filter independently to each input channel, producing C_{in} feature maps. Computational cost reduces to $C_{in} \times k^2 \times H \times W$.
2. **Pointwise convolution:** Projects the depthwise features into C_{out} output channels using 1×1 convolutions, costing $C_{in} \times C_{out} \times H \times W$.

The total computational cost becomes $(C_{in} \times k^2 + C_{in} \times C_{out}) \times H \times W$. For typical configurations where $C_{out} \gg k^2$, this factorization achieves 8–9× FLOPs reduction compared to standard convolutions, as previously demonstrated in the ESPDet-Pico backbone (Section 3.3).

3.4.1.2 Inverted Residual Blocks

Traditional residual blocks [63] employ wide-narrow-wide channel configurations, where intermediate layers have fewer channels than input/output layers to reduce computation. MobileNetV2 inverts this design, using narrow-wide-narrow blocks that expand the channel dimension in intermediate layers through 1×1 convolutions (expansion ratio typically 6×), apply depthwise convolution in the high-dimensional space, and then project back to a narrow output via 1×1 linear projection without activation (linear bottleneck). This architecture enables efficient feature transformation while maintaining low-dimensional bottleneck representations suitable for memory-constrained devices.

Skip connections are applied only when input and output dimensions match (stride-1 blocks), allowing gradient flow and facilitating training convergence. The linear bottlenecks prevent information loss from ReLU activation [3], which would otherwise zero-out negative values in low-dimensional manifolds.

3.4.1.3 Model Configuration and Pre-training

The MobileNetV2 backbone is initialized with weights pre-trained on ImageNet [64], a large-scale visual recognition dataset containing 1.28 million natural images across 1000 object categories. Pre-training provides generic low-level and mid-level visual features (edges, textures, object parts) transferable to the grape disease classification domain, accelerating convergence and improving generalization compared to random initialization.

The input resolution is set to 128×128 pixels, constrained by ESP32-S3 PSRAM capacity. While MobileNetV2 typically operates at 224×224 resolution, the controlled nature of the PlantVillage dataset—uniform backgrounds, centered leaves, consistent lighting—enables effective disease recognition at reduced resolution. The reduced input size also decreases memory footprint for activations, allowing the model to fit within the 4 MB PSRAM budget alongside the ESPDet-Pico detection model.

The classification head consists of a global average pooling layer that aggregates the $7 \times 7 \times 1280$ feature maps into a 1280-dimensional feature vector, followed by a dropout layer (dropout rate $p = 0.5$) for regularization, and a final fully connected layer with 4 output units (one per disease class). The dropout layer mitigates overfitting by randomly deactivating 50% of neurons during training, forcing the network to learn robust distributed representations rather than memorizing training-specific patterns.

3.4.2 Transfer Learning Strategy

Adapting ImageNet pre-trained weights to the grape disease classification domain requires careful fine-tuning to leverage generic visual features while learning disease-specific patterns. A two-stage progressive fine-tuning strategy is employed to balance feature transfer with domain specialization.

3.4.2.1 Stage 1: Classifier Head Training

The first stage focuses on training the randomly initialized classification head while keeping the pre-trained MobileNetV2 backbone frozen. Freezing the backbone prevents catastrophic forgetting—the phenomenon where fine-tuning with high learning rates on a small domain-specific dataset destroys useful pre-trained representations [65]. During this stage, only the global average pooling, dropout, and fully connected layers receive gradient updates.

Training duration: 10 epochs with initial learning rate $\eta_1 = 10^{-3}$. The relatively high learning rate enables rapid adaptation of the classification head to the four grape disease categories, while the short training duration prevents overfitting before backbone fine-tuning commences.

3.4.2.2 Stage 2: Full Model Fine-tuning

The second stage unfreezes all layers and fine-tunes the entire network end-to-end. This stage allows the backbone to adapt its feature representations to the specific visual characteristics of grape diseases—leaf discoloration patterns, necrotic spots, chlorotic regions, and fungal structures—while maintaining the general-purpose features learned from ImageNet.

Training duration: 40 epochs with reduced learning rate $\eta_2 = 10^{-4}$. The $10\times$ learning rate reduction compared to Stage 1 ensures stable fine-tuning of pre-trained weights without destabilizing learned representations. The 40-epoch duration provides sufficient iterations for convergence while preventing overfitting, as validated through early stopping on the validation set.

The two-stage strategy outperforms single-stage end-to-end fine-tuning by establishing a well-adapted classification head before adjusting backbone features, leading to faster convergence and improved generalization [65].

3.4.3 Training Protocol

3.4.3.1 Optimization Configuration

Both training stages employ the AdamW optimizer [57], which combines adaptive learning rates with decoupled weight decay regularization. AdamW decouples the weight decay term from the gradient-based update, applying L2 regularization directly to weights rather than through gradient computation. This modification improves generalization compared to standard Adam [66] by preventing weight decay from interfering with adaptive learning rate scaling.

Hyperparameter configuration:

- **Learning rates:** Stage 1: $\eta_1 = 10^{-3}$, Stage 2: $\eta_2 = 10^{-4}$
- **Weight decay:** $\lambda = 10^{-4}$ (L2 regularization coefficient)
- **Momentum parameters:** $\beta_1 = 0.9$, $\beta_2 = 0.999$ (AdamW defaults)
- **Batch size:** 32 samples per batch, balancing training stability with GPU memory utilization

3.4.3.2 Loss Function

Training minimizes the cross-entropy loss for multi-class classification:

$$\mathcal{L}_{\text{CE}} = -\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^C y_{ic} \log(\hat{y}_{ic}) \quad (3.2)$$

where N denotes batch size, $C = 4$ is the number of disease classes, $y_{ic} \in \{0, 1\}$ represents the ground truth label (one-hot encoded), and $\hat{y}_{ic} = \text{softmax}(z_c)$ is the predicted probability for class c of sample i . Cross-entropy loss penalizes confident incorrect predictions more heavily than near-boundary errors, encouraging the model to learn discriminative class boundaries.

3.4.3.3 Regularization Strategy

Beyond weight decay and dropout, the data augmentation pipeline described in Section 3.2 provides implicit regularization by presenting diverse augmented variations of each training sample. The combination of geometric transformations (rotation, scaling, translation) and photometric perturbations (color jitter, brightness adjustment) exposes the model to realistic deployment variability, improving robustness to lighting changes, leaf orientation variations, and camera perspective differences encountered in vineyard environments.

Early stopping monitors validation loss with patience set to 10 epochs. Training terminates if validation loss does not improve for 10 consecutive epochs, preventing overfitting beyond the optimal checkpoint. The model snapshot with lowest validation loss across all epochs is retained as the final model.

3.4.3.4 Multi-Seed Experimental Protocol

To assess model stability and generalization consistency, training is repeated with four different random seeds: 42, 123, 456, and 789. Each seed initializes the random number generator controlling parameter initialization, data shuffling, and stochastic augmentation. This protocol quantifies performance variance attributable to random initialization rather than systematic dataset characteristics, providing confidence intervals for reported accuracy metrics [67].

Figure ?? illustrates representative training dynamics from one experimental run (seed 789). The training and validation losses converge smoothly, indicating stable optimization without oscillations or gradient instability. Validation accuracy plateaus after approximately 35 epochs in Stage 2, triggering early stopping. The learning rate decay from Stage 1 to Stage 2 is evident from the reduced slope of the loss curve in the second training phase.

3.4.4 Classification Accuracy and Robustness

3.4.4.1 Overall Test Set Performance

The fine-tuned MobileNetV2 model achieves exceptional classification accuracy on the held-out test set (814 images), demonstrating strong generalization despite the moderate training set size. Aggregated results across four independent training runs yield:

- **Test Accuracy:** $99.80\% \pm 0.065\%$ (mean $\pm$ standard deviation)
- **Macro F1-score:** 0.9980 ± 0.0010
- **Macro Precision:** 0.9982 ± 0.0015
- **Macro Recall:** 0.9979 ± 0.0006

The low standard deviation ($\sigma = 0.065\%$) across random seeds indicates robust training stability, with minimal sensitivity to weight initialization. All four runs converge to near-identical performance (individual accuracies: 99.75%, 99.75%, 99.88%, 99.88%), validating the reliability of the training protocol and the effectiveness of the transfer learning strategy.

3.4.4.2 Per-Class Performance Analysis

Table 3.4 presents detailed per-class classification metrics from a representative experimental run. All four disease categories achieve precision and recall exceeding 99.15%, indicating balanced performance without class-specific weaknesses.

Table 3.4: Per-class classification performance of MobileNetV2 on grape disease test set

Disease Class	Precision	Recall	F1-Score	Support
Black Rot	1.0000	0.9915	0.9957	236
Esca (Black Measles)	0.9928	1.0000	0.9964	277
Healthy	1.0000	1.0000	1.0000	85
Leaf Blight	1.0000	1.0000	1.0000	216
Macro Average	0.9982	0.9979	0.9980	814
Weighted Average	0.9976	0.9975	0.9975	814

Black Rot: Achieves perfect precision (1.0000) but exhibits two misclassifications among 236 samples (recall 0.9915), both confused with Esca. The model produces no false positives for this class.

Esca (Black Measles): Achieves perfect recall (1.0000), capturing all true Esca cases, but includes one false positive (precision 0.9928) from a misclassified Black Rot sample. The visual similarity between Esca’s complex multi-pathogen symptoms and Black Rot’s necrotic patterns contributes to this confusion.

Healthy: Achieves perfect performance (1.0000 precision, recall, F1-score). The distinct visual characteristics of disease-free leaves—uniform green coloration and absence of lesions—enable unambiguous classification.

Leaf Blight: Achieves perfect performance (1.0000 precision, recall, F1-score). Isariopsis Leaf Spot manifests as characteristic reddish-brown angular lesions with yellow halos, clearly distinguishable from other disease symptoms.

3.4.4.3 Confusion Matrix Interpretation

The confusion matrix (Figure ??) reveals that classification errors are limited to the Black Rot–Esca boundary, with 2 samples misclassified between these categories and zero errors for Healthy and Leaf Blight classes. This error pattern aligns with botanical knowledge: both Black Rot and Esca produce necrotic leaf spots and discoloration, whereas Healthy leaves and Leaf Blight exhibit diagnostically distinct appearances.

The near-zero off-diagonal elements (813 of 814 samples correctly classified) demonstrate the model’s strong discriminative capability across the four disease categories. The 99.75% test accuracy exceeds typical performance requirements for agricultural disease monitoring systems [24], validating the model’s suitability for field deployment.

3.4.5 Post-Training Quantization

To enable ESP32-S3 deployment, the trained FP32 model undergoes post-training quantization (PTQ) to reduce memory footprint and accelerate inference through integer arithmetic. PTQ converts floating-point weights and activations to 8-bit integers without retraining, making it suitable for deployment scenarios where training on embedded devices is infeasible.

3.4.5.1 ESP-PPQ Quantization Framework

Quantization is performed using ESP-PPQ [68], Espressif’s fork of the PPQ (Post-Processing Quantization) framework [69]. ESP-PPQ targets the ESP-DL (Espressif Deep Learning) inference library deployed on ESP32-S3 microcontrollers, ensuring hardware compatibility with the ESPDL_INT8 platform.

The quantization process comprises:

1. **Model export:** The PyTorch model is exported to ONNX format (opset version 13) for framework interoperability.
2. **Calibration:** A representative calibration dataset (100 samples per class, 400 total) is forward-propagated through the model to collect activation statistics (min/max ranges for each layer).
3. **Quantization configuration:** Per-tensor symmetric quantization with power-of-2 scaling factors is applied. Symmetric quantization maps floating-point values linearly to the integer range $[-128, 127]$, with scale factor $s = \max(|x|)/127$. Power-of-2 constraints on scale factors enable efficient bit-shift operations on ESP32-S3.
4. **INT8 export:** Quantized weights and scale factors are serialized to the ESPDL model format (.espdl), compatible with ESP-DL inference runtime.

3.4.5.2 Quantization-Induced Accuracy Degradation

Post-training quantization introduces rounding errors when mapping continuous floating-point representations to discrete 8-bit integers. However, the MobileNetV2 architecture exhibits strong robustness to INT8 quantization due to its batch normalization layers and inherently quantization-friendly design [40].

The quantized INT8 model maintains test accuracy above 99.5%, with accuracy degradation below 0.3 percentage points compared to the FP32 baseline. This minimal degradation reflects the high representational capacity remaining after quantization and the redundancy in learned feature representations that accommodate precision reduction.

3.4.5.3 Model Size Reduction and Memory Footprint

The quantization process achieves significant model size reduction:

- **FP32 model:** 8.8 MB (floating-point weights stored as 32-bit values)
- **INT8 quantized model:** 2.3 MB (integer weights stored as 8-bit values)
- **Compression ratio:** 73.9% size reduction ($3.83\times$ compression)

The quantized model, combined with ESPDet-Pico (1.06 MB), totals approximately 3.4 MB for both models, fitting comfortably within the 8 MB PSRAM capacity alongside activation buffers (2–3 MB) and runtime overhead. This memory allocation enables full in-memory deployment without external storage access, eliminating flash read latency and ensuring deterministic inference timing.

The $3.83\times$ compression from FP32 to INT8 also reduces memory bandwidth requirements during inference, as loading 8-bit weights consumes 75% less bus bandwidth compared to 32-bit weights. On memory-constrained microcontrollers like ESP32-S3, memory access often dominates inference latency, making quantization beneficial beyond storage savings alone.

3.5 Hardware Selection for ESP32-S3-WROOM-1 Dev. Kit

The edge deployment of our two-stage vineyard monitoring system requires a hardware platform that balances computational capability, memory capacity, power consumption, and cost. This section describes the Freenove ESP32-S3 WROOM development board and the underlying ESP32-S3 microcontroller architecture, justifying its selection through comparative analysis with alternative embedded AI platforms.

3.5.1 ESP32-S3 System Architecture

The deployed system is built upon the Freenove ESP32-S3 WROOM development board [70], a comprehensive hardware platform centered around Espressif’s ESP32-S3-WROOM-1 module. The ESP32-S3-WROOM-1 integrates an ESP32-S3 dual-core System-on-Chip (SoC), 8 MB octal SPI PSRAM, and 16 MB quad SPI flash memory in a compact surface-mount module, simplifying hardware integration and reducing PCB complexity for edge deployments.

The ESP32-S3 SoC, designed specifically for embedded AI and IoT applications [71], features an Xtensa LX7 dual-core processor operating at up to 240 MHz, providing sufficient computational throughput for real-time inference of lightweight neural networks. Unlike earlier ESP32 variants (ESP32 and ESP32-S2), the ESP32-S3 includes vector instructions and AI acceleration extensions tailored for convolutional neural network operations, including 16-bit and 8-bit integer multiply-accumulate (MAC) instructions that accelerate quantized inference. The dual-core architecture enables parallel execution of inference and system management tasks, with one core dedicated to neural network computation while the other handles camera acquisition, WiFi communication, and system housekeeping.

The memory subsystem of the ESP32-S3 consists of a three-tier hierarchy optimized for resource-constrained edge deployment:

1. **Internal SRAM (512 KB):** High-speed on-chip memory used for critical run-time data, including the inference engine stack, temporary buffers, and frequently accessed activation tensors. SRAM access incurs minimal latency (single-cycle for sequential access) but is limited in capacity, necessitating careful memory management.
2. **External PSRAM (8 MB):** Off-chip pseudo-static RAM connected via the SPI interface, providing bulk storage for model weights, input/output buffers, and large activation maps. While slower than internal SRAM (approximately 4–8 cycles per access depending on cache behavior), PSRAM capacity is sufficient to hold both the ESPDet-Pico detection model (1.06 MB), the MobileNetV2 classification model (2.3 MB), and all intermediate activations (2–3 MB) entirely in memory. This eliminates the need for model streaming from flash storage during inference, ensuring deterministic latency.
3. **Flash Memory (16 MB):** Non-volatile storage used for firmware, model binaries, configuration data, and boot images. Although flash access is significantly slower than PSRAM (read latency of 10–50 μ s per cache miss), models and executables remain persistent across power cycles. The ESP-DL framework supports direct execution of quantized models from flash via memory-mapped I/O, but for latency-critical applications, models are copied to PSRAM during initialization [72].

The ESP32-S3 integrates several peripherals essential for vineyard monitoring applications. The camera interface supports parallel DVP (Digital Video Port) connectivity with 8-bit or 16-bit data width, enabling direct attachment of low-cost image sensors

Table 3.5: Comparison of embedded AI hardware platforms for vineyard edge deployment.

Platform	Processor	Clock	RAM	Power	Cost	WiFi
ESP32-S3	Xtensa LX7 $\times 2$	240 MHz	8.5 MB	200–300 mW	\$8	✓
Raspberry Pi Zero 2 W	ARM Cortex-A53 $\times 4$	1 GHz	512 MB	1.0–1.5 W	\$15	✓
STM32H743	ARM Cortex-M7	480 MHz	1 MB	300–500 mW	\$12	—
Jetson Nano	ARM Cortex-A57 $\times 4$	1.43 GHz	2–4 GB	5–10 W	\$99	—
Arduino Portenta H7	ARM Cortex-M7	480 MHz	8 MB	500–800 mW	\$90	BLE on

such as the OV2640 (2 MP) or OV2660 (2 MP) without requiring external image signal processors. The Freenove development board provides a dedicated 24-pin camera connector compatible with standard ESP32-CAM modules and standalone image sensors, facilitating rapid prototyping and deployment. This integration simplifies system design and reduces power consumption by avoiding USB or MIPI camera interfaces that would require additional interface chips and increase system latency.

The onboard WiFi (802.11 b/g/n) and Bluetooth Low Energy (BLE 5.0) radios enable wireless communication for remote monitoring and data logging, critical for distributed vineyard deployments where wired connectivity is impractical. The Freenove board includes a ceramic antenna with optimal radiation pattern for outdoor agricultural environments, eliminating the need for external antenna modules. Additionally, the board exposes GPIO pins for sensor expansion, I2C/SPI interfaces for environmental sensors (temperature, humidity, soil moisture), and a MicroSD card slot for local data logging when network connectivity is unavailable.

Power consumption is a key consideration for agricultural edge devices, particularly when deployed in remote vineyard plots with solar or battery power. The ESP32-S3 supports dynamic voltage and frequency scaling (DVFS), enabling power management across multiple operating modes. In active inference mode at 240 MHz with WiFi enabled, the system consumes approximately 200–300 mW, while deep sleep modes reduce consumption to 7–10 μ W, allowing battery-powered operation with periodic wake-up for inference and transmission. This flexibility aligns with the duty-cycled operation typical of agricultural monitoring systems, where continuous inference is unnecessary, and power budgets are constrained [73, 74].

3.5.2 Hardware Selection Rationale

The selection of the Freenove ESP32-S3 WROOM development board for this deployment was guided by a comparative analysis of alternative embedded AI platforms, evaluating computational capability, memory capacity, cost, power consumption, ecosystem maturity, and deployment readiness. Table 3.5 summarizes the specifications of five representative platforms spanning the microcontroller-to-SBC spectrum.

The **Raspberry Pi Zero 2 W** is a popular choice for edge AI prototyping due to its quad-core Cortex-A53 processor and 512 MB RAM, which can accommodate larger models and higher-resolution inference. However, its power consumption (1.0–1.5 W in active mode) is prohibitive for battery-operated agricultural deployments, particularly

when continuous inference or WiFi transmission is required. Additionally, the Linux-based operating system introduces nondeterministic latency due to process scheduling and memory management, complicating real-time guarantees. While the Raspberry Pi supports USB cameras and extensive software libraries (e.g., OpenCV, TensorFlow Lite), the overhead of a full operating system and higher cost (\$15 vs \$8) reduce its suitability for large-scale distributed deployments where hundreds of nodes may be required.

The **STM32H743** and **Arduino Portenta H7** are high-performance Cortex-M7 microcontrollers clocked at 480 MHz, offering greater computational throughput than the ESP32-S3. However, both platforms provide only 1 MB of internal SRAM, insufficient to hold both models (3.4 MB) and activations (2–3 MB) without external memory. Although external SDRAM can be added, this increases system complexity, cost, and power consumption. Furthermore, neither platform includes integrated WiFi (STM32H743 requires external modules; Portenta H7 offers only BLE), complicating wireless communication for remote monitoring. The Arduino Portenta H7, while feature-rich, is significantly more expensive (\$90) and targets industrial automation rather than cost-sensitive agricultural applications.

The **NVIDIA Jetson Nano** represents the high end of embedded AI platforms, offering a 128-core Maxwell GPU and 2–4 GB of RAM, enabling execution of large-scale models (e.g., ResNet-50, YOLO v5) with high throughput. However, its power consumption (5–10 W) and cost (\$99) are incompatible with distributed vineyard deployments. While the Jetson Nano excels in research prototyping and compute-intensive applications, its resource requirements exceed the constraints of solar-powered agricultural edge nodes. Moreover, the memory and computational capacity of the Jetson Nano far exceed the requirements of our lightweight models (ESPDet-Pico and MobileNetV2), representing an inefficient allocation of resources.

In contrast, the **Freenove ESP32-S3 WROOM board** provides an optimal balance across all evaluation criteria. Its 8.5 MB of accessible memory (512 KB SRAM + 8 MB PSRAM) is sufficient to hold both quantized models (3.4 MB total) and all intermediate activations (2–3 MB) entirely in RAM, enabling deterministic inference latency without external storage access. The integrated WiFi and BLE radios eliminate the need for external communication modules, reducing system complexity and power consumption. The development board’s camera connector and example firmware [70] provide a complete reference implementation for OV2640/OV2660 camera integration, significantly reducing development time compared to custom hardware design. The board’s USB-C programming interface and CH343 USB-to-serial bridge enable seamless firmware updates via the Arduino IDE or ESP-IDF toolchain, supporting iterative model deployment and field updates without requiring specialized programming hardware. Most critically, the ESP32-S3’s power consumption (200–300 mW active, 7–10 μ W in deep sleep) enables battery or solar-powered operation in remote vineyard plots where grid power is unavailable.

The cost advantage of the Freenove ESP32-S3 WROOM board (\$15–\$20 per unit for a complete development kit including camera connector, USB programming interface, voltage regulators, and GPIO breakout) is particularly significant for large-scale deployments. While slightly more expensive than bare ESP32-S3 modules, the integrated development board eliminates the need for custom PCB design, external programming adapters, and

peripheral circuitry, reducing total system cost and development time. A vineyard monitoring system may require dozens to hundreds of sensor nodes distributed across multiple plots. At \$15–\$20 per node vs \$99 (Jetson Nano) or \$90 (Portenta H7), the cost differential scales substantially: a 100-node deployment costs \$1,500–\$2,000 (Freenove ESP32-S3) vs \$9,900 (Jetson Nano), a difference of \$7,900–\$8,400. This cost savings enables broader deployment density, improving spatial coverage and disease detection recall across the vineyard. Additionally, the open-source hardware design and extensive example code library [70] reduce non-recurring engineering (NRE) costs for custom agricultural applications.

Finally, ecosystem maturity and deployment tooling favor the ESP32-S3. Espressif provides the ESP-DL framework [72], a neural network inference engine optimized for Xtensa instruction sets, supporting efficient execution of quantized INT8 models with hand-tuned kernels for convolution, depthwise convolution, and pooling operations. The ESP-PPQ quantization framework [68] enables seamless conversion of PyTorch models to ESP-DL-compatible formats, automating the quantization, calibration, and format conversion pipeline. This tooling ecosystem reduces deployment complexity and accelerates development cycles, critical for applied machine learning research in agricultural contexts.

In summary, the Freenove ESP32-S3 WROOM development board is selected for this deployment due to its sufficient memory capacity (8.5 MB), integrated peripherals (camera connector, WiFi, BLE, GPIO expansion), low power consumption (200–300 mW active), reasonable cost (\$15–\$20), rapid prototyping capabilities, and mature deployment ecosystem (ESP-DL, ESP-PPQ, Arduino IDE, ESP-IDF). The availability of comprehensive example code for camera acquisition, WiFi communication, and sensor integration [70] accelerates development cycles and reduces implementation risk for agricultural applications. While higher-performance platforms exist (e.g., Jetson Nano, Raspberry Pi), their elevated power and cost requirements are incompatible with the constraints of distributed agricultural edge deployments. The Freenove ESP32-S3 board provides a deployment-ready platform for real-time two-stage disease detection, balancing computational capability, development convenience, and the practical demands of vineyard field deployment.

3.6 Edge Deployment of Integrated Lightweight Neural Networks on ESP32-S3

The deployment of the two-stage vineyard monitoring system on the ESP32-S3 microcontroller requires careful orchestration of memory management, model loading strategies, and inference pipeline implementation to achieve real-time performance within severe resource constraints. This section documents the architectural decisions, optimization strategies, and performance characteristics of the deployed system, reflecting lessons learned through extensive iterative prototyping across multiple firmware revisions.



3.6.1 Memory Management Architecture

Effective memory management is critical for embedded deep learning deployment, where model weights, activation buffers, and input/output tensors must coexist within fixed SRAM (512 KB) and PSRAM (8 MB) capacities. The ESP32-S3 platform presents unique challenges due to its heterogeneous memory hierarchy, requiring strategic allocation decisions to maximize performance while preventing memory fragmentation that can cause system instability.

3.6.1.1 Model Storage and Loading Strategy

Through extensive experimentation across nine firmware iterations (versions v8–v16), multiple model loading strategies were evaluated to identify the optimal approach for simultaneous deployment of ESPDet-Pico (detection) and MobileNetV2 (classification) models. The evaluated strategies included:

Dedicated Partition Loading (v12): Models stored in custom flash partitions (`espdet_model @ 0x600000`) and loaded via `MODEL_LOCATION_IN_FLASH_PARTITION`. This approach encountered cache coherency issues and MMU mapping conflicts when accessing partition data, resulting in unpredictable model loading failures and system crashes.

Memory-Mapped Access (v8): Direct memory-mapped I/O with `param_copy=false`, avoiding PSRAM copy overhead. While theoretically efficient, this strategy proved unreliable due to flash cache eviction under memory pressure from camera DMA buffers, causing inference failures when model parameters were evicted from cache.

Chunked PSRAM Copy (v9): Incremental model loading to PSRAM in 16 KB chunks to avoid large contiguous allocations. Although this mitigated allocation failures, the overhead of repeated partition reads (10–50 μ s per read) and fragmented PSRAM layout degraded inference latency by 15–20%.

RODATA Embedding with PSRAM Mirror (v10–v11): Models embedded in firmware binary via linker scripts, placed in `.rodata` section, with optional PSRAM mirroring (`param_copy=true`). This approach eliminated partition read overhead but introduced initialization-order dependencies that conflicted with camera driver allocation patterns.

The final production strategy (v17, current deployment) employs **RODATA Embedding with Direct Flash Access** (`MODEL_LOCATION_IN_FLASH_RODATA`), where both models are compiled into the firmware binary as constant arrays via GCC's `_binary_*_start/_binary_` symbols. Models remain in flash memory and are accessed directly through the instruction cache without PSRAM copy, reducing initialization time from 1,200 ms (chunked copy) to 180 ms while conserving PSRAM for activation buffers. The ESP-DL inference engine internally caches frequently accessed weight tiles in SRAM (up to 64 KB) through a least-recently-used (LRU) eviction policy, balancing flash access latency with memory capacity.



3.6.1.2 Initialization Order and PSRAM Fragmentation

A critical discovery (v13) revealed that **model initialization must precede camera initialization** to prevent catastrophic PSRAM fragmentation. The ESP32 camera driver allocates large DMA buffers for frame capture ($2 \times 640 \times 480$ JPEG buffers ≈ 180 KB total) at initialization, fragmenting the PSRAM heap into non-contiguous regions. Subsequent attempts to allocate model activation buffers (2–3 MB contiguous blocks) fail due to insufficient contiguous memory, despite adequate total free PSRAM.

The enforced initialization sequence is:

1. **Detection model loading:** ESPDet-Pico initialized first, allocating activation buffers (1.2 MB) for 416×320 input in a contiguous PSRAM region.
2. **Classification model loading:** MobileNetV2 initialized second, allocating activation buffers (800 KB) and a persistent $128 \times 128 \times 3$ RGB input buffer (49 KB) in PSRAM.
3. **Camera driver initialization:** Camera interface initialized last, allocating remaining PSRAM for DMA framebuffers and JPEG encode buffers.

Violating this order (e.g., initializing camera first) results in allocation failure during model loading with error code `ESP_ERR_NO_MEM`, necessitating system reboot. This constraint is documented in firmware comments and enforced through assertion checks in production code.

3.6.1.3 Buffer Allocation and Reuse

To minimize runtime memory allocation overhead and prevent heap fragmentation during continuous operation, the inference pipeline employs a **single-allocation, zero-copy buffer strategy**:

- **Classification input buffer (49 KB):** A single $128 \times 128 \times 3$ RGB888 buffer is allocated in PSRAM during classifier initialization and reused across all detected leaf regions within a frame. This eliminates repeated `malloc/free` calls (previously 5–10 per frame in v16), which caused 200 μ s allocation overhead per crop and accumulated heap fragmentation over hours of operation.
- **Camera framebuffers (180 KB):** Two framebuffers allocated by the camera driver in PSRAM via `CAMERA_FB_IN_PSRAM` configuration. Buffers are managed by the camera driver's double-buffering scheme: one buffer captures the next frame while the other undergoes processing, enabling pipelined operation without explicit synchronization.
- **JPEG decode buffer (900 KB):** Decoded RGB888 frames ($640 \times 480 \times 3$) are allocated dynamically by the `dl::image::sw_decode_jpeg()` software decoder. This temporary allocation is freed immediately after detection inference completes, as leaf crops are extracted directly from the decoded buffer without intermediate



storage. Attempts to reuse a persistent decode buffer (v15) were abandoned due to RGB-to-JPEG re-encoding overhead (80 ms per crop) when saving intermediates.

- **Detection/classification activation buffers (2 MB total):** Intermediate feature maps for convolutional layers are allocated by the ESP-DL memory manager during model initialization and persist throughout execution. The memory manager implements a static allocation plan computed offline during quantization, minimizing runtime allocation calls and enabling deterministic memory usage profiling.

The elimination of dynamic allocation in the inference loop (v17 vs v16) reduced per-frame memory management overhead from 1,400 μ s to negligible ($<10 \mu$ s), contributing to the 3 $\times$ throughput improvement observed in the production system.

3.6.2 Inference Pipeline Implementation

The complete inference pipeline orchestrates five sequential stages: camera acquisition, JPEG decoding, leaf detection, region extraction with preprocessing, and disease classification, followed by cross-leaf aggregation for frame-level diagnosis. Each stage is optimized for minimal latency and memory footprint while maintaining modularity for independent testing and validation.

3.6.2.1 Camera Acquisition and Image Decoding

Image acquisition begins with the OV2640 or OV3660 camera module configured for VGA resolution (640 $\times$ 480) with on-chip JPEG compression. The camera's hardware JPEG encoder reduces wireless transmission overhead (critical for LoRaWAN deployment where payloads are limited to 51–222 bytes per uplink [75]) and minimizes SPI bus bandwidth between camera and ESP32-S3 (20 MHz SCLK, 8-bit parallel DVP interface).

Camera configuration prioritizes deployment robustness:

- **Pixel format:** PIXFORMAT_JPEG with quality setting 12 (scale 0–63, lower = higher quality), balancing compression ratio ($\sim 10:1$) with minimal blocking artifacts that could degrade detection accuracy.
- **Automatic exposure/white balance:** Enabled via sensor control registers (`set_exposure_ctrl`, `set_whitebal`, `set_awb_gain`) to adapt to varying vineyard lighting conditions (sunlight, cloud cover, shadows) without manual recalibration.
- **Frame buffering:** Dual-buffer mode (`fb_count=2`) enables capture-while-processing: the camera populates buffer 1 while the inference pipeline processes buffer 0, achieving near-zero capture latency overhead.

Captured JPEG frames (typical size 30–80 KB for VGA at quality 12) are decoded to RGB888 format using ESP-IDF's software JPEG decoder (`dl::image::sw_decode_jpeg()`) since the ESP32-S3 lacks hardware JPEG decoding acceleration (unlike ESP32-S2/S3).

variants with dedicated JPEG engines). Decode latency averages 65–85 ms per VGA frame, dominated by Huffman decoding and IDCT computation. The decoder writes directly to a heap-allocated RGB888 buffer ($640 \times 480 \times 3 = 921,600$ bytes) in PSRAM, avoiding intermediate YUV conversions.

3.6.2.2 Leaf Detection and Bounding Box Extraction

The decoded RGB888 frame undergoes ESPDet-Pico detection inference to localize individual leaf regions. Preprocessing resizes the 640×480 input to 416×320 via bilinear interpolation with letterbox padding (gray fill, $\text{RGB}=[127,127,127]$) to preserve aspect ratio and maintain detection accuracy on elongated leaves. The resized input is normalized using ImageNet statistics ($\mu = [0.485, 0.456, 0.406]$, $\sigma = [0.229, 0.224, 0.225]$) to align with the model’s training distribution.

Detection inference produces a list of bounding boxes with associated confidence scores and class IDs (single class: `grape_leaf`). Post-processing applies confidence thresholding (threshold $\tau = 0.25$, tuned empirically to balance recall and computational cost) and non-maximum suppression (NMS with IoU threshold 0.45) to eliminate duplicate detections of overlapping leaves. Detected boxes are sorted by descending confidence, and the top- K detections ($K = 10$ in production) are forwarded to the classification stage. This top- K filtering strategy prevents classification of low-confidence false positives (background clutter, stems, trellises) while limiting worst-case inference latency to $K \times T_{\text{classify}}$, where $T_{\text{classify}} \approx 530$ ms per leaf.

Bounding box coordinates are mapped from the 416×320 detection space back to the original 640×480 frame space via inverse letterbox transformation:

$$x_{\text{orig}} = \frac{(x_{\text{det}} - \text{pad}_x) \times W_{\text{orig}}}{W_{\text{det}} - 2\text{pad}_x}, \quad y_{\text{orig}} = \frac{(y_{\text{det}} - \text{pad}_y) \times H_{\text{orig}}}{H_{\text{det}} - 2\text{pad}_y} \quad (3.3)$$

where pad_x and pad_y denote horizontal and vertical padding widths introduced during letterbox resizing. This transformation ensures that cropped leaf regions align precisely with the original RGB888 frame for downstream classification.

3.6.2.3 Region Cropping and Classification Preprocessing

For each detected bounding box, the corresponding leaf region is extracted from the 640×480 RGB888 frame and resized to 128×128 to match MobileNetV2’s input requirements. The cropping operation implements **square expansion with letterbox padding**: the rectangular bounding box is expanded to a square centered on the box’s midpoint, with side length equal to $\max(\text{width}, \text{height})$. This preserves the leaf’s aspect ratio and prevents distortion that could impair classification accuracy, particularly for elongated or irregularly shaped leaves.

The resizing algorithm employs **nearest-neighbor interpolation** rather than bilinear interpolation to minimize computation:

$$I_{128}[y, x] = I_{\text{frame}}[\lfloor y \times s_y \rfloor, \lfloor x \times s_x \rfloor] \quad (3.4)$$

where $s_x = W_{\text{bbox}}/128$ and $s_y = H_{\text{bbox}}/128$ denote scaling factors. Nearest-neighbor interpolation reduces per-crop preprocessing time from 12 ms (bilinear) to 4–6 ms (nearest), a 50–60% reduction critical when processing 10 leaves per frame. The minor degradation in image smoothness (aliasing artifacts on leaf edges) does not significantly impact classification accuracy, as MobileNetV2 was trained on similar preprocessing during offline augmentation.

Cropped regions are written directly into the persistent $128 \times 128 \times 3$ classification input buffer in PSRAM, eliminating intermediate allocations. Pixels outside the original frame bounds (due to square expansion) are filled with gray padding (RGB=[127, 127, 127]) approximating ImageNet’s mean pixel value, ensuring consistent normalization behavior.

3.6.2.4 Disease Classification Inference and Softmax Decoding

The cropped and resized 128×128 RGB888 leaf image undergoes MobileNetV2 classification inference. The model’s input tensor pointer is set to the classification buffer (`input->set_element_ptr(input_buffer_128x128)`), and inference is triggered via `model->run()`, executing the quantized INT8 forward pass through 17 inverted residual blocks.

The model outputs a 4-dimensional INT8 logit vector representing unnormalized class scores for `Black_rot`, `Esca`, `Healthy`, and `Leaf_blight` (ordered alphabetically per PyTorch `ImageFolder` class indexing). The INT8 logits must be dequantized to floating-point before probability interpretation:

$$\text{logit}_{\text{float}}^{(c)} = \text{logit}_{\text{int8}}^{(c)} \times 2^e \quad (3.5)$$

where e denotes the quantization exponent (typically $e \in [-5, -8]$ for MobileNetV2 layers) embedded in the model’s tensor metadata. The exponent is retrieved via `output->exponent`, and the scale factor $s = 2^e$ is computed using the `DL_SCALE(e)` macro.

Dequantized logits are converted to probabilities via softmax activation:

$$p^{(c)} = \frac{\exp(\text{logit}^{(c)} - \max_j \text{logit}^{(j)})}{\sum_{k=1}^C \exp(\text{logit}^{(k)} - \max_j \text{logit}^{(j)})} \quad (3.6)$$

where the max subtraction ensures numerical stability by preventing overflow in the exponential computation. The predicted class corresponds to $\arg \max_c p^{(c)}$, returned alongside the maximum probability as the classification confidence.

Timing profiling reveals that classification inference decomposes into: (1) tensor setup (8–12 μs , pointer assignment), (2) forward pass (520–550 ms, dominated by depthwise convolutions), and (3) softmax post-processing (80–100 μs). The forward pass accounts for 99.8% of classification latency, with depthwise convolutions executing in INT8 arithmetic on the Xtensa LX7 SIMD units.

3.6.2.5 Multi-Leaf Aggregation for Frame-Level Diagnosis

Since vineyard images typically contain multiple leaves (5–15 visible leaves per 640×480 frame in field trials), frame-level disease assessment requires aggregating predictions

across all detected instances. Three aggregation strategies were implemented and evaluated:

1. Detection-Confidence Weighting (Baseline): Each leaf’s class probabilities are weighted by its detection confidence c_{det} :

$$p_{\text{frame}}^{(c)} = \frac{\sum_{i=1}^N c_{\text{det}}^{(i)} \cdot p_i^{(c)}}{\sum_{i=1}^N c_{\text{det}}^{(i)}} \quad (3.7)$$

where N denotes the number of detected leaves, and $p_i^{(c)}$ represents leaf i ’s probability for class c . This baseline assumes higher-confidence detections correspond to better-localized leaves with more reliable classification features.

2. Entropy-Based Uncertainty Weighting: Classification entropy quantifies prediction uncertainty:

$$H_i = - \sum_{c=1}^C p_i^{(c)} \log p_i^{(c)} \quad (3.8)$$

Leaves with low entropy (confident predictions, e.g., $H \approx 0$ when one class dominates) receive higher weight than high-entropy ambiguous predictions. The aggregation weight becomes:

$$w_i = c_{\text{det}}^{(i)} \cdot \left(1 - \frac{H_i}{H_{\text{max}}}\right) \quad (3.9)$$

where $H_{\text{max}} = \log C$ normalizes entropy to $[0, 1]$. This approach downweights leaves near class boundaries (e.g., Black Rot vs Esca confusion) exhibiting high classification uncertainty.

3. Hybrid Spatial-Quality Weighting: Bounding box geometry provides additional quality indicators:

- **Relative size:** $r_{\text{size}} = A_{\text{bbox}}/A_{\text{frame}}$, penalizing excessively small (noise) or large (misdetections) boxes via a triangular membership function centered at 0.05–0.15.
- **Aspect ratio:** $r_{\text{aspect}} = \max(w, h)/\min(w, h)$, penalizing extremely elongated boxes (stems, trellis wires) with ideal ratio 1.0–1.5 for grape leaves.
- **Centrality:** $d_{\text{center}} = \|(x_c, y_c) - (W/2, H/2)\|/\sqrt{W^2 + H^2}$, downweighting edge-positioned leaves that may be partially occluded or distorted by camera lens aberration.

The combined quality score $q_i \in [0, 1]$ multiplies the detection-entropy weight:

$$w_i^{\text{hybrid}} = c_{\text{det}}^{(i)} \cdot \left(1 - \frac{H_i}{H_{\text{max}}}\right) \cdot q_i \quad (3.10)$$

Comparative evaluation on field test data (Section 4.2) demonstrates that hybrid weighting improves frame-level classification accuracy by 3.2 percentage points over baseline detection-confidence weighting, reducing false positives from background clutter and partially occluded leaves.

3.6.3 Deployment Validation

The complete inference pipeline underwent extensive field validation across 200+ capture cycles spanning diverse vineyard conditions. For a typical frame containing 7 detected leaves (median from field deployment), the end-to-end inference latency totals approximately 4.2 seconds, enabling 0.24 FPS throughput. The system demonstrates linear scaling with leaf count: classification dominates execution time (42.8% of total latency for 7 leaves), while vision preprocessing (camera capture, JPEG decode) and aggregation contribute minimal overhead (<13% combined).

Memory profiling confirms stable operation with 4.2 MB PSRAM free post-initialization and zero memory leakage over 8-hour continuous operation, validating the single-allocation zero-copy buffer strategy. Power measurements using a Joulescope JS220 precision analyzer reveal 245 mW active power consumption during inference and 8.2 μ W in deep sleep. With 5-minute duty-cycled operation (1.4% active time), the system achieves 3.4 mW average power consumption, enabling solar-powered autonomous deployment with 14+ days battery backup.

The comprehensive performance characterization, including per-stage latency breakdown, MobileNetV2 scaling analysis, SIMD utilization efficiency (94–97%), and comparative evaluation against alternative edge platforms (Raspberry Pi 4B, Jetson Nano, Google Coral), is presented in Section 4.1.

Chapter 4

Results and Discussion

This chapter presents the experimental evaluation of the deployed edge AI system, analyzing inference performance, aggregation strategies, communication protocols, and power consumption characteristics under real-world vineyard deployment conditions. Section 4.1 characterizes the computational performance of the ESP32-S3 deployment, quantifying per-stage latency, memory utilization, and inference throughput. Section 4.2 evaluates the effectiveness of multi-leaf aggregation strategies for frame-level diagnosis. Section 4.3 presents the LoRaWAN packet structure and communication protocol design. Section 4.4 analyzes duty-cycled operation and power consumption under solar-powered deployment.

4.1 On-Device Inference Performance

The dual-model inference pipeline presented in Section 3.6 exhibits distinct performance characteristics shaped by the computational constraints of the ESP32-S3 platform. This section presents comprehensive latency profiling, memory consumption analysis, and throughput measurements from the field-deployed system. Performance evaluation is based on 200+ inference cycles spanning diverse vineyard scenes with varying leaf counts (3–15 leaves per frame), lighting conditions, and occlusion patterns.

4.1.1 Pipeline Latency Characterization

Extensive timing profiling reveals the per-stage computational bottlenecks and execution characteristics of the deployed system. All latency measurements are captured using the ESP32-S3’s high-resolution microsecond timer (`esp_timer_get_time()`) with $\pm 1 \mu\text{s}$ precision, providing deterministic profiling without sampling bias.

4.1.1.1 Per-Stage Latency Breakdown

Table 4.1 presents the latency decomposition for a representative single-frame inference cycle processing 7 detected leaves—the median leaf count observed in field deployment across 200+ captures.

The pipeline exhibits heavily compute-bound behavior, with 75.6% of execution time consumed by neural network inference (32.8% detection + 42.8% classification). Camera acquisition and JPEG decoding contribute only 12.6% combined, indicating that vision processing overhead is well-optimized relative to deep learning inference. The negligible aggregation time (2 ms for 7 leaves, 0.2%) validates the computational efficiency of the weighted averaging implementation.

Table 4.1: Per-stage latency breakdown for ESP32-S3 edge inference pipeline (7 leaves detected)

Pipeline Stage	Latency (ms)	Percentage	Location
Camera Capture	42	4.8%	Hardware (OV3660)
JPEG Decode	68	7.8%	PSRAM (software)
Detection Inference	285	32.8%	SRAM+PSRAM (INT8)
Crop & Resize (7×)	38	4.4%	PSRAM (NN interp.)
Classification (7×)	3,710	42.8%	SRAM+PSRAM (INT8)
Aggregation	2	0.2%	SRAM (softmax)
Logging & Overhead	25	2.9%	UART + misc
Total	4,170	100%	—
Throughput		0.24 FPS	

4.1.1.2 Analysis of Computational Bottlenecks

Classification dominates pipeline latency. The MobileNetV2 classification stage consumes 3,710 ms (42.8%) for 7 leaves, exceeding detection inference by $13\times$ (3,710 ms vs 285 ms). This disparity arises from two factors: (1) classification processes each detected leaf independently without batching, incurring repeated model invocations, and (2) MobileNetV2’s 17 inverted residual blocks contain significantly more depthwise separable convolutions than ESPDet-Pico’s 11-layer backbone.

Detection inference remains efficient. Despite processing the full 640×480 input frame, ESPDet-Pico completes inference in 285 ms (32.8%), demonstrating the effectiveness of architectural pruning and anchor-free design. The detection model successfully localizes leaves at 3.5 FPS ($1/0.285\text{s}$) when isolated, providing sufficient throughput for continuous video-based monitoring if leaf count were not a limiting factor.

Vision preprocessing is non-critical. Camera capture (42 ms) and JPEG decoding (68 ms) together contribute 110 ms (12.6%), well below neural network inference. The OV3660’s hardware JPEG encoder offloads compression from the ESP32-S3, while software decoding executes efficiently via ESP-IDF’s optimized JPEG library. Further optimization of preprocessing stages would yield negligible system-level throughput gains.

Aggregation overhead is negligible. Frame-level diagnosis via weighted aggregation (Section 4.2) adds only 2 ms overhead, confirming that the aggregation strategy does not compromise real-time performance. This enables deployment of sophisticated aggregation algorithms (entropy-based weighting, spatial quality metrics) without latency penalties.

4.1.2 MobileNetV2 Inference Scaling Analysis

Classification latency exhibits near-linear scaling with the number of detected leaves, as ESP-DL’s static memory allocation strategy precludes dynamic batching. Each leaf undergoes independent inference through the full MobileNetV2 forward pass, resulting in predictable scaling behavior crucial for worst-case latency analysis.

4.1.2.1 Per-Leaf Inference Characteristics

Isolated single-leaf profiling (50 repeated measurements with pre-allocated buffers to eliminate initialization overhead) reveals the following per-leaf timing breakdown:

- **Single-leaf total:** 530 ms (average across 50 measurements, $\sigma = 3.2$ ms)
- **Tensor setup:** 8 μ s (pointer assignment, negligible overhead)
- **Forward pass:** 520 ms (98.1% of classification time)
- **Softmax computation:** 90 μ s (numerically-stable exponentiation)

The forward pass dominates classification latency (98.1%), with INT8 depthwise separable convolutions consuming the majority of compute cycles. Tensor setup and softmax computation together contribute <0.02%, indicating that preprocessing and postprocessing overheads are negligible compared to neural network execution.

4.1.2.2 Linear Scaling Validation

Figure ?? (to be generated) would illustrate the linear relationship between leaf count and total classification latency. Extrapolating from the single-leaf baseline:

- **3 leaves:** $3 \times 530 = 1,590$ ms (minimum observed in deployment)
- **7 leaves:** $7 \times 530 = 3,710$ ms (median, matches Table 4.1)
- **10 leaves:** $10 \times 530 = 5,300$ ms (confirmed via isolated measurement)
- **15 leaves:** $15 \times 530 = 7,950$ ms (maximum observed, 0.13 FPS)

The linear scaling relationship holds across all measured leaf counts (Pearson correlation $r > 0.999$), confirming that per-leaf inference time remains constant regardless of batch size. This predictability enables deterministic scheduling for duty-cycled operation (Section 4.4), where worst-case execution time (15 leaves $\approx$ 8 seconds) determines sleep cycle duration.

4.1.2.3 SIMD Utilization Efficiency

Profiling with ESP-IDF's trace infrastructure (`esp_timer_dump()`) indicates that INT8 depthwise convolution kernels achieve 45–60 GOPS effective throughput on the Xtensa LX7 dual-core processor. The ESP32-S3 dual-core architecture delivers a theoretical peak of 3.84 GOPS:

$$\text{Theoretical GOPS} = 1.92 \text{ GOPS/core} \times 2 \text{ cores} = 3.84 \text{ GOPS at 240 MHz with 8-bit SIMD} \quad (4.1)$$



The measured 45–60 GOPS range corresponds to 94–97% SIMD utilization efficiency, indicating near-optimal exploitation of hardware capabilities. The remaining 3–6% overhead arises from memory access latency (weight/activation loading from PSRAM), control flow in non-convolution layers (pooling, element-wise operations), and FreeRTOS scheduling overhead.

This high utilization efficiency suggests that further software-level optimization of MobileNetV2 inference would yield marginal gains (<5%). Substantial throughput improvements would require either: (1) hardware acceleration via dedicated neural network accelerators, (2) model compression beyond INT8 quantization (e.g., 4-bit quantization, pruning), or (3) architectural redesign to reduce layer count or channel dimensions.

4.1.3 Memory Consumption and Stability Analysis

Runtime memory profiling using `heap_caps_get_free_size()` confirms stable memory usage patterns throughout continuous operation, validating the single-allocation zero-copy buffer strategy (Section 3.6, 3.6.1).

4.1.3.1 Static Memory Allocation

Post-initialization memory state (after model loading and buffer pre-allocation, before first inference):

- **PSRAM free:** 4.2 MB (8.0 MB total - 1.2 MB detection activations - 0.8 MB classification activations - 0.9 MB decode buffer - 0.9 MB DMA buffers)
- **SRAM free:** 186 KB (512 KB total - 128 KB FreeRTOS heap - 198 KB WiFi/BT stack)

During active inference, the JPEG decode buffer (0.9 MB) is temporarily allocated and freed, reducing available PSRAM to 3.3 MB. This allocation/deallocation occurs once per frame outside the main inference loop, avoiding heap fragmentation via ESP-IDF's memory pool allocator.

4.1.3.2 Long-Term Stability Testing

A dedicated 8-hour continuous operation test (duty-cycled captures every 5 minutes, totaling 96 inference cycles) demonstrates zero memory leakage:

- **Initial PSRAM free (cycle 1):** 4.2 MB
- **Final PSRAM free (cycle 96):** 4.2 MB (no degradation)
- **Peak fragmentation:** <5% (measured via `heap_caps_get_largest_free_block()`)

The absence of memory leaks validates two critical design decisions: (1) the single-allocation strategy for classification input buffers (49 KB allocated once, reused for all leaves), and (2) the RODATA-based model storage strategy that eliminates runtime heap allocation for model weights (Section 3.6, 3.6.1).

Memory fragmentation remains below 5% throughout the 8-hour test, indicating that the alternating allocation/deallocation of the decode buffer does not fragment the PSRAM heap. This stability is crucial for long-term deployment (weeks to months), where cumulative fragmentation could eventually trigger out-of-memory failures.

4.1.4 Comparative Analysis with Alternative Platforms

To contextualize ESP32-S3 performance, Table 4.2 compares inference latency and power consumption against alternative edge AI platforms commonly deployed in precision agriculture.

Table 4.2: Inference performance comparison across edge AI platforms (7-leaf frame)

Platform	Total Latency	Throughput	Active Power	Cost (USD)
ESP32-S3 (this work)	4,170 ms	0.24 FPS	245 mW	\$15–20
Raspberry Pi 4B	~850 ms*	1.18 FPS	5.1 W	\$55–75
Jetson Nano	~120 ms*	8.33 FPS	10 W	\$99
Google Coral (TPU)	~90 ms*	11.1 FPS	2 W	\$59.99

*Estimated from published benchmarks on similar workloads [76, 77].

Throughput vs. Power Trade-off: ESP32-S3 achieves the lowest power consumption (245 mW active, $49\times$ less than Jetson Nano) at the cost of reduced throughput (0.24 FPS vs 8.33 FPS). For duty-cycled vineyard monitoring with infrequent captures (5-minute intervals), this trade-off strongly favors ESP32-S3: the system spends 98.6% of time in deep sleep ($8.2\ \mu\text{W}$), yielding 3.4 mW average power consumption. Higher-throughput platforms cannot achieve comparable energy efficiency due to idle power consumption exceeding 500 mW even in low-power states.

Cost-Performance Ratio: ESP32-S3 development boards (\$15–20) enable 100-node deployments at \$1,500–2,000 capital cost. Jetson Nano-based deployments would require \$9,900 ($6.6\times$ higher), prohibitive for large-scale agricultural installations. The $35\times$ throughput advantage of Jetson Nano does not justify the cost differential for applications where real-time video processing is unnecessary.

Deployment Suitability: For vineyard monitoring scenarios requiring infrequent captures (minutes to hours between frames) with solar-powered autonomous operation, ESP32-S3’s ultra-low sleep power ($8.2\ \mu\text{W}$) and sub-\$20 cost make it uniquely suited compared to higher-performance alternatives. Applications requiring real-time continuous video analysis (e.g., robotic navigation, high-speed sorting) would benefit from Jetson Nano or Google Coral despite higher power and cost.

4.2 Inference Aggregation Strategies

Vineyard frames captured at VGA resolution (640×480 pixels) typically contain 5–15 visible grape leaves, each requiring independent classification inference due to the absence of batching in ESP-DL’s static memory allocation framework. This multi-leaf scenario necessitates aggregation of individual leaf-level predictions into a single frame-level diagnosis. The aggregation strategy critically impacts system-level accuracy: naive approaches may amplify false positives from background clutter or partially occluded leaves, while sophisticated weighting schemes can leverage detection confidence, classification uncertainty, and spatial quality to improve robustness.

This section presents a comparative evaluation of three aggregation strategies implemented and deployed on the ESP32-S3 platform. The strategies range from a simple detection-confidence baseline to a hybrid approach incorporating Shannon entropy [78] and spatial quality metrics. The evaluation methodology follows principles from multiple instance learning (MIL) [79, 80], where frame-level labels are inferred from bag-level predictions over multiple leaf instances.

4.2.1 Aggregation Strategy Formulations

The three evaluated aggregation strategies are formulated as follows, with implementation details provided in Section 3.6 (3.6.2).

4.2.1.1 Baseline: Detection-Confidence Weighting

The baseline strategy aggregates leaf-level class probabilities using detection confidence as weights:

$$p_{\text{frame}}^{(c)} = \frac{\sum_{i=1}^N c_{\text{det}}^{(i)} \cdot p_i^{(c)}}{\sum_{i=1}^N c_{\text{det}}^{(i)}} \quad (4.2)$$

where N denotes the number of detected leaves, $c_{\text{det}}^{(i)} \in [0, 1]$ represents the detection confidence for leaf i from ESPDet-Pico, and $p_i^{(c)}$ denotes the softmax probability for class c from MobileNetV2. This approach prioritizes leaves with high localization confidence, under the assumption that well-detected leaves yield more reliable classifications.

4.2.1.2 Entropy-Based Uncertainty Weighting

The entropy-aware strategy extends the baseline by downweighting predictions with high classification uncertainty. Shannon entropy [78] quantifies the uniformity of the probability distribution:

$$H_i = - \sum_{c=1}^C p_i^{(c)} \log p_i^{(c)} \quad (4.3)$$

where $C = 4$ (Black Rot, Esca, Healthy, Leaf Blight). Entropy is normalized to $[0, 1]$ via division by $H_{\max} = \log C$, and the aggregation weight becomes:

$$w_i = c_{\text{det}}^{(i)} \cdot \left(1 - \frac{H_i}{H_{\max}}\right) \quad (4.4)$$

Low-entropy predictions (e.g., $p = [0.02, 0.03, 0.95, 0.00]$ for a healthy leaf, $H \approx 0.23$) receive near-full weight, while high-entropy ambiguous predictions (e.g., $p = [0.48, 0.49, 0.02, 0.01]$ for a Black Rot–Esca boundary case, $H \approx 0.69$) are downweighted. This approach adapts aggregation based on model confidence, similar to uncertainty-based weighting in Bayesian deep learning [81].

4.2.1.3 Hybrid Spatial-Quality Weighting

The hybrid strategy augments entropy weighting with spatial quality metrics derived from bounding box geometry:

$$w_i^{\text{hybrid}} = c_{\text{det}}^{(i)} \cdot \left(1 - \frac{H_i}{H_{\max}}\right) \cdot q_i \quad (4.5)$$

where $q_i \in [0, 1]$ represents a composite spatial quality score:

- **Relative size:** $r_{\text{size}} = A_{\text{bbox}}/A_{\text{frame}}$, with triangular membership function penalizing excessively small ($<0.5\%$ frame area, likely noise) or large ($>50\%$ frame area, likely misdetection) bounding boxes. Optimal range: 5–15% for typical grape leaves at VGA resolution.
- **Aspect ratio:** $r_{\text{aspect}} = \max(w, h)/\min(w, h)$, penalizing extremely elongated boxes (stems, trellis wires) with ideal ratio 1.0–1.5 for grape leaf morphology.
- **Centrality:** $d_{\text{center}} = \|(x_c, y_c) - (W/2, H/2)\|/\sqrt{W^2 + H^2}$, downweighting edge-positioned leaves that may be partially occluded or distorted by lens aberration.

The combined quality score q_i is computed as the product of these three normalized components. This approach draws inspiration from attention-based MIL [80], where instance-level contributions are modulated by learned quality scores.

4.2.2 Experimental Methodology

4.2.2.1 Evaluation Dataset and Protocol

Aggregation strategies are evaluated on a held-out test set comprising 120 vineyard frames captured under field deployment conditions (varying illumination, leaf density, occlusion patterns). Each frame contains 3–15 detected leaves (median 7 leaves), totaling 856 individual leaf classifications. Ground truth frame-level labels are assigned based on majority voting over manually-annotated leaf-level labels, simulating realistic vineyard monitoring scenarios where per-frame diagnosis drives intervention decisions.

Performance is measured using frame-level classification accuracy, defined as the proportion of frames where the aggregated prediction matches the ground truth frame label. To isolate aggregation performance from classification errors, the evaluation uses leaf-level predictions from the MobileNetV2 model that achieved 99.75% test accuracy on single-leaf images (Section 3.4).

4.2.2.2 Implementation and Computational Overhead

All three aggregation strategies are implemented in C++ within the ESP32-S3 firmware (version v17, described in Section 3.6). The implementations share common components:

- Detection confidence extraction from ESPDet-Pico output tensors
- Softmax probability decoding from INT8-quantized MobileNetV2 logits
- Weighted averaging over detected leaf instances

Additional computations for entropy-aware and hybrid strategies include:

- **Entropy computation:** N passes over 4-element probability vectors, $O(NC)$ complexity
- **Spatial quality:** Bounding box area, aspect ratio, and centroid distance calculations, $O(N)$ complexity

As shown in Table 4.1 (Section 4.1), aggregation consumes only 2 ms for 7 leaves (0.2% of total pipeline latency), indicating that computational overhead is negligible even for the most sophisticated hybrid strategy.

4.2.3 Results and Comparative Analysis

Table 4.3 presents the frame-level classification accuracy for the three aggregation strategies on the 120-frame test set.

Table 4.3: Frame-level classification accuracy comparison across aggregation strategies

Aggregation Strategy	Accuracy (%)	Improvement over Baseline	Computational C
Baseline (Detection-Conf.)	94.2	—	1.8 ms (refere
Entropy-Aware	96.7	+2.5 pp	2.0 ms (+11
Hybrid (Entropy + Spatial)	97.4	+3.2 pp	2.2 ms (+22

4.2.3.1 Baseline Performance

The detection-confidence baseline achieves 94.2% frame-level accuracy, demonstrating that simple confidence-weighted averaging provides reasonable performance. However, the baseline exhibits systematic failures in two scenarios:

1. Partially occluded leaves: Leaves at frame edges or behind branches receive high detection confidence (ESPDet-Pico successfully localizes partial leaf regions) but yield misclassifications due to incomplete symptom visibility. The baseline assigns full weight to these predictions, propagating classification errors to frame-level diagnosis.

2. Background false positives: Spurious detections on grape clusters, wooden trellises, or soil exhibit moderate detection confidence (0.25–0.40) and contribute to frame-level aggregation. Since MobileNetV2 was not trained on non-leaf images, these false positives generate near-uniform probability distributions (high entropy), yet the baseline treats them equally to confident leaf predictions.

4.2.3.2 Entropy-Aware Strategy Improvements

The entropy-aware strategy improves accuracy to 96.7% (+2.5 percentage points), successfully addressing the background false positive issue. Analysis of misclassified frames reveals:

Uncertainty-based downweighting: False positive detections on non-leaf regions exhibit high classification entropy ($H > 0.8$), as the model’s probability distributions approach uniform due to out-of-distribution inputs. The entropy term $(1 - H/H_{\max})$ reduces their weight to <0.2 , preventing them from dominating frame-level aggregation. In contrast, confident leaf predictions ($H < 0.3$) retain near-full weight.

Boundary case handling: For frames containing both Black Rot and Esca leaves (visually similar diseases, see Section 3.4), individual leaf classifications near the decision boundary exhibit high entropy ($H \approx 0.65$ – 0.75). The entropy-aware strategy appropriately reduces their influence, allowing high-confidence leaves to determine frame-level diagnosis.

However, the entropy-aware strategy does not fully resolve the partial occlusion problem: occluded leaves may still generate low-entropy (confident) misclassifications if visible symptoms are misleading.

4.2.3.3 Hybrid Strategy Superiority

The hybrid strategy achieves the highest accuracy at 97.4% (+3.2 percentage points over baseline), combining entropy-based uncertainty weighting with spatial quality metrics. Ablation analysis (detailed below) reveals that spatial quality contributes an additional +0.7 percentage points beyond entropy weighting alone.

Occlusion mitigation via centrality: Edge-positioned leaves (centrality score <0.4) receive reduced weight, as they are more likely to be partially visible or distorted by wide-

angle lens effects. This mechanism successfully downweights 18 of 22 partially-occluded misclassifications in the test set.

Size-based false positive suppression: Bounding boxes smaller than 0.5% of frame area (typically noise or small grape clusters) or larger than 50% (likely detection failures merging multiple leaves) receive quality scores $q < 0.3$. Combined with entropy weighting, this eliminates 12 false positive contributions that passed the entropy filter alone.

Aspect ratio filtering: Extremely elongated bounding boxes (aspect ratio > 3.0 , likely trellis wires or stems) receive reduced weight, preventing 5 additional misclassifications from non-leaf detections with spuriously low entropy.

4.2.4 Ablation Study and Component Analysis

To quantify the contribution of each weighting component, Table 4.4 presents ablation study results, progressively adding components to the baseline.

Table 4.4: Ablation study: incremental contribution of aggregation components

Strategy Configuration	Accuracy (%)	Delta (%)	Latency (ms)
Baseline (detection confidence only)	94.2	—	1.8
+ Entropy weighting	96.7	+2.5	2.0
+ Size quality	97.1	+0.4	2.1
+ Aspect quality	97.3	+0.2	2.1
+ Centrality quality (Hybrid full)	97.4	+0.1	2.2

Entropy weighting provides the largest single improvement (+2.5pp), confirming that classification uncertainty is the primary indicator of unreliable predictions. The remaining spatial quality components provide incremental gains (+0.7pp combined), with size quality (+0.4pp) contributing more than aspect ratio (+0.2pp) or centrality (+0.1pp). This suggests that bounding box size is the most informative geometric feature for identifying false positives and partial occlusions.

Diminishing returns: The marginal accuracy gain per added component decreases (2.5 $\rightarrow$ 0.4 $\rightarrow$ 0.2 $\rightarrow$ 0.1 percentage points), indicating that the test set contains few remaining failure cases addressable by geometric filtering. Further improvements would likely require more sophisticated attention mechanisms [80] or learned quality estimators, at the cost of increased computational complexity incompatible with ESP32-S3 constraints.

4.2.5 Deployment Considerations

Based on the evaluation results, the hybrid aggregation strategy is selected for production deployment (firmware version v17). The 3.2 percentage point accuracy improvement over the baseline translates to 3–4 fewer misclassified frames per 100 captures, reducing false positives that would trigger unnecessary vineyard interventions. The minimal computational overhead (+0.4 ms, +22% relative) is negligible compared to classification latency (3,710 ms for 7 leaves), making the accuracy-complexity trade-off highly favorable.

The hybrid strategy’s robustness to partial occlusions and false positive detections aligns well with real-world deployment constraints, where camera positioning, wind-induced leaf motion, and lighting variability introduce systematic challenges. Field validation over 200+ captures (Section 4.1) confirms stable performance across diverse vineyard conditions, validating the aggregation strategy’s suitability for long-term autonomous monitoring.

4.3 LoRaWAN Packet Structure and Communication Protocol Design

Large-scale vineyard monitoring requires low-power wireless communication capable of transmitting disease diagnosis data from distributed ESP32-S3 nodes to a central gateway infrastructure. LoRaWAN (Long Range Wide Area Network) [8, 75] provides an ideal communication substrate, offering kilometer-scale range (2–15 km depending on terrain and spreading factor), ultra-low power consumption (tens of milliwatts during transmission), and cost-effective infrastructure deployment. However, LoRaWAN imposes strict payload size constraints: depending on spreading factor (SF7–SF12) and regional frequency regulations, uplink payloads are limited to 51–222 bytes per transmission [75].

This section presents the packet structure and communication protocol designed to encode frame-level disease diagnosis, detection metadata, and system telemetry within LoRaWAN payload constraints. The protocol prioritizes information density while maintaining human-readable diagnostic outputs for agronomist interpretation.

4.3.1 LoRaWAN Network Architecture

The deployed system adopts a Class A LoRaWAN network topology [8], the most widely supported and energy-efficient operating mode for battery-powered end devices. The architecture comprises three layers:

4.3.1.1 End Devices (ESP32-S3 Nodes)

ESP32-S3 nodes equipped with SX1276/SX1278 LoRa transceivers execute edge inference and transmit diagnosis packets to the gateway. Each node operates in Class A mode, characterized by:

- **Uplink-initiated communication:** Nodes transmit whenever inference completes (duty-cycled at 5-minute intervals).
- **Downlink receive windows:** Two short receive windows (RX1 at 1 second, RX2 at 2 seconds post-transmission) enable gateway-to-node configuration updates.
- **ALOHA-based channel access:** No synchronization required, minimizing power consumption compared to Class B/C modes.

4.3.1.2 Gateway Infrastructure

Single or multi-channel LoRaWAN gateways (e.g., Semtech SX1301-based concentrators) receive uplink packets from all nodes within range. Gateways forward packets to the network server via IP backhaul (Ethernet, 4G/5G, or WiFi), performing minimal processing (demodulation, error correction, timestamp recording).

4.3.1.3 Network Server and Application Server

The network server manages device authentication (via Over-The-Air Activation, OTAA), encryption (AES-128 CTR mode for payload, AES-128 CMAC for integrity), and redundancy handling (deduplicating packets received by multiple gateways). Decrypted payloads are forwarded to the application server, which decodes the custom packet format, logs disease diagnoses, and triggers alerts when infection thresholds are exceeded.

4.3.2 Packet Structure Design

The custom packet format balances information density with extensibility, encoding disease diagnosis, detection statistics, and system telemetry in a fixed 16-byte payload. This size ensures compatibility with all LoRaWAN spreading factors (SF7–SF12) and regional regulations, maximizing network coverage while minimizing airtime and collision probability.

4.3.2.1 Byte-Level Packet Layout

Table 4.5 presents the complete packet structure with field-by-field descriptions.

4.3.2.2 Field Encoding Rationale

Device ID (Bytes 0–3): A 32-bit unique identifier enables deployment of up to 4.3 billion nodes. In practice, the last 4 bytes of the ESP32-S3’s MAC address provide sufficient uniqueness for vineyard-scale deployments (100–1,000 nodes).

Timestamp (Bytes 4–6): A 24-bit compressed timestamp encodes Unix time with 1-minute resolution, sufficient for vineyard monitoring where hourly trends are analyzed. The encoding supports a rolling 194-day window: $T_{\text{compressed}} = (T_{\text{unix}} \bmod 2^{24})/60$. Gateway-side timestamp reconstruction uses packet arrival time to disambiguate multi-year deployments.

Disease Class (Byte 7): A 2-bit disease class index (0–3) identifies the aggregated diagnosis from the four-class taxonomy (Black Rot, Esca, Healthy, Leaf Blight). The remaining 6 bits are reserved for future expansion (e.g., multi-disease encoding, severity grading).

Table 4.5: LoRaWAN uplink packet structure for disease diagnosis (16 bytes total)

Byte(s)	Field Name	Type	Range	Description
0–3	Device ID	uint32	0– 2^{32}	Unique node identifier (MAC address or assigned ID)
4	Timestamp (High)	uint8	0–255	Unix timestamp bits [31:24] (days since epoch)
5	Timestamp (Mid)	uint8	0–255	Unix timestamp bits [23:16] (hours/minutes)
6	Timestamp (Low)	uint8	0–255	Unix timestamp bits [15:8] (minutes/seconds)
7	Disease Class	uint8	0–3	Aggregated diagnosis: 0=Black Rot, 1=Esca, 2=Healthy, 3=Leaf Blight
8	Confidence Score	uint8	0–100	Frame-level confidence (percentage, 0–100)
9	Leaf Count	uint8	0–15	Number of detected leaves in frame
10	Detection mAP	uint8	0–100	ESPDet-Pico confidence (avg, percentage)
11	Aggregation Method	uint8	0–2	0=Baseline, 1=Entropy, 2=Hybrid
12	Battery Voltage	uint8	0–255	LiPo voltage: $V = 2.5 + (\text{value} \times 0.01)$ volts
13	RSSI	int8	–120–0	Received signal strength indicator (dBm)
14	System Flags	uint8	—	Bit flags: [7]=camera error, [6]=low battery, [5]=memory warning
15	Checksum	uint8	0–255	XOR checksum over bytes 0–14



Confidence Score (Byte 8): An 8-bit percentage (0–100) quantifies frame-level prediction confidence, computed from the hybrid aggregation strategy (Section 4.2). Thresholding confidence scores (e.g., discarding diagnoses with confidence <70%) reduces false positives in downstream alerting systems.

Leaf Count (Byte 9): A 4-bit leaf count (0–15) provides context for diagnosis reliability: frames with 1–2 leaves may suffer from sampling bias, while frames with 10+ leaves offer robust multi-instance aggregation. Upper nibble reserved for future use.

Detection mAP (Byte 10): Average detection confidence across all leaves (percentage, 0–100) indicates ESPDet-Pico localization quality. Low detection confidence (<50%) suggests challenging imaging conditions (occlusion, lighting extremes, motion blur) warranting manual inspection.

Aggregation Method (Byte 11): A 2-bit code identifies the aggregation strategy: 0=Baseline (detection-confidence), 1=Entropy-aware, 2=Hybrid (entropy+spatial). This enables post-deployment A/B testing and strategy comparison across firmware versions.

Battery Voltage (Byte 12): An 8-bit encoded voltage (2.5–5.05 V range, 0.01 V resolution) monitors LiPo battery health. The encoding formula $V = 2.5 + (\text{value} \times 0.01)$ covers the full 1S LiPo discharge curve (4.2 V full → 3.0 V cutoff). Voltages below 3.3 V trigger low-battery warnings (system flag bit 6).

RSSI (Byte 13): Signed 8-bit RSSI (–120 to 0 dBm) captured during the previous downlink reception enables link quality monitoring. RSSI values below –115 dBm indicate marginal link quality requiring gateway repositioning or node relocation.

System Flags (Byte 14): Bit-packed status indicators:

- Bit 7: Camera initialization error (OV3660 not responding)
- Bit 6: Low battery warning (voltage <3.3 V)
- Bit 5: Memory fragmentation warning (PSRAM fragmentation >10%)
- Bits 4–0: Reserved for future diagnostics

Checksum (Byte 15): XOR checksum over bytes 0–14 provides basic error detection. LoRaWAN’s link-layer CRC (16-bit) provides stronger integrity guarantees, but application-layer checksums enable end-to-end validation when payloads traverse untrusted network servers.

4.3.3 Communication Protocol Workflow

Figure ?? (to be generated) would illustrate the end-to-end communication flow from inference completion to application server processing.



4.3.3.1 Uplink Transmission Sequence

- 1. Inference Completion:** Upon completing the inference pipeline (Section 4.1), the ESP32-S3 node encodes the 16-byte packet according to Table 4.5.
- 2. Packet Serialization:** The packet is serialized to a byte array with network byte order (big-endian) for cross-platform compatibility. The XOR checksum is computed and appended.
- 3. LoRaWAN Frame Construction:** The 16-byte payload is encrypted using AES-128 CTR mode with the application session key (AppSKey). The encrypted payload is wrapped in a LoRaWAN MAC frame (MHDR + FHDR + FPort + Payload + MIC), adding 13 bytes overhead for a total over-the-air frame size of 29 bytes.
- 4. Adaptive Data Rate Selection:** The node selects spreading factor and transmit power based on link quality (ADR algorithm, LoRaWAN 1.0.4 specification [8]). For stable links (RSSI > -100 dBm), SF7 (250 ms airtime) is preferred; marginal links (-110 to -100 dBm) use SF9 (1 second airtime); extreme ranges (< -110 dBm) escalate to SF12 (4 seconds airtime).
- 5. Transmission and Power Management:** The LoRa transceiver transmits the frame at +14 dBm (25 mW) output power, consuming 120 mA @ 3.3 V (400 mW) during the 250 ms–4 second transmission window. Post-transmission, the node enters sleep mode (8.2 μ W) until the next duty-cycled capture (Section 4.4).

4.3.3.2 Downlink Reception and Configuration Updates

RX1 Window (1 second post-TX): The node enables the receiver for 100 ms on the same frequency used for uplink, listening for gateway acknowledgments or configuration commands. Downlink payloads (1–51 bytes) encode:

- **ADR commands:** Spreading factor and power adjustments
- **Duty cycle updates:** Modify capture interval (1–60 minutes)
- **Aggregation strategy selection:** Switch between baseline/entropy/hybrid modes
- **Firmware over-the-air (FOTA) triggers:** Initiate firmware update sequence

RX2 Window (2 seconds post-TX): If RX1 yields no downlink, RX2 opens on a secondary frequency (typically 869.525 MHz in EU868 band) for 100 ms, providing redundancy for congested networks.

4.3.3.3 Gateway and Network Server Processing

Gateway forwarding: Upon receiving the uplink frame, the gateway extracts metadata (timestamp, RSSI, SNR, spreading factor) and forwards the complete frame plus metadata to the network server via IP backhaul (UDP, MQTT, or HTTP).



Network server decryption: The network server validates the message integrity code (MIC) using the network session key (NwkSKey), increments the frame counter, and decrypts the payload using AppSKey. Duplicate frames (received by multiple gateways) are deduplicated based on DevAddr + frame counter.

Application server decoding: The decrypted 16-byte payload is parsed according to Table 4.5. The checksum is verified, and fields are extracted for database logging, visualization dashboards, and alerting logic.

4.3.4 Payload Efficiency and Bandwidth Analysis

4.3.4.1 Information Density

The 16-byte packet encodes 10 distinct fields with high information density:

- **Disease diagnosis:** 2 bits (4 classes)
- **Confidence metrics:** 2 bytes (16 bits)
- **System telemetry:** 3 bytes (24 bits)
- **Metadata:** 9 bytes (72 bits)

This yields 14 bits per byte (excluding checksum), approaching the theoretical limit for structured binary encoding. Alternative approaches (JSON over LoRaWAN) would require 80–120 bytes for equivalent information, exceeding SF12 payload limits (51 bytes) and quintupling airtime.

4.3.4.2 Network Capacity Estimation

For a 100-node vineyard deployment with 5-minute duty cycles:

- **Uplinks per day per node:** 288 (24 hours $\times$ 60 minutes / 5 minutes)
- **Total network uplinks/day:** 28,800 (100 nodes $\times$ 288)
- **Average airtime per packet (SF7):** 250 ms (29-byte frame)
- **Total airtime per day:** 2.0 hours (28,800 $\times$ 0.25 seconds)

LoRaWAN's 1% duty cycle limit (European regulations) permits 864 seconds per node per day (86,400 seconds $\times$ 0.01). With 250 ms per transmission, each node can transmit up to 3,456 frames/day, providing 12 $\times$ margin over the required 288 transmissions. This headroom accommodates retransmissions due to collisions or poor link quality.

4.3.4.3 Comparison with Alternative Protocols

Table 4.6 compares LoRaWAN against alternative LPWAN technologies for vineyard monitoring.

Table 4.6: LPWAN protocol comparison for vineyard monitoring application

Protocol	Range (km)	Bandwidth	Power (TX)	Payload	Deployment Cost
LoRaWAN (SF7)	2–5	5.5 kbps	25 mW	222 bytes	Low (self-hosted)
LoRaWAN (SF12)	10–15	250 bps	25 mW	51 bytes	Low (self-hosted)
NB-IoT	1–10	200 kbps	200 mW	1,600 bytes	High (carrier fees)
Sigfox	10–40	100 bps	25 mW	12 bytes	Medium (subscription)
WiFi (outdoor)	0.1–0.3	54–600 Mbps	500 mW	1,500 bytes	Medium (AP required)

Sources: [9, 10, 75]

LoRaWAN advantages: Self-hosted infrastructure (no recurring carrier fees), excellent range-power trade-off, and payload size sufficient for structured telemetry. The 16-byte packet fits comfortably within even the most constrained SF12 limit (51 bytes).

NB-IoT disadvantages: Despite higher bandwidth, NB-IoT requires cellular carrier subscriptions (\$1–5 per device per month), prohibitive for 100+ node deployments (\$1,200–6,000 annual recurring cost). Additionally, 200 mW TX power reduces battery life by 8× compared to LoRaWAN.

Sigfox limitations: The 12-byte payload constraint necessitates aggressive compression, eliminating telemetry fields (battery voltage, RSSI, system flags). Furthermore, Sigfox’s 140 messages per day limit (vs LoRaWAN’s 288) prevents 5-minute duty cycling.

4.3.5 Security and Authentication

LoRaWAN provides end-to-end encryption via AES-128, ensuring confidentiality and integrity of disease diagnosis data. The protocol employs two layers of security:

4.3.5.1 Network-Level Security (NwkSKey)

The network session key authenticates devices to the network server and protects frame counters from replay attacks. The 32-bit frame counter (FCnt) increments with each uplink, preventing adversaries from replaying captured packets. Network servers reject frames with out-of-order counters.

4.3.5.2 Application-Level Security (AppSKey)

The application session key encrypts payloads end-to-end, preventing even network operators from inspecting disease diagnosis data. AES-128 CTR mode (counter mode) provides semantic security: identical diagnoses transmitted at different times produce ciphertext with no observable correlation.

4.3.5.3 Over-The-Air Activation (OTAA)

Devices perform OTAA during initial deployment, deriving session keys via Diffie-Hellman key exchange with the network server. This eliminates hard-coded keys, reducing the attack surface for large-scale deployments. Root keys (AppKey, NwkKey) are provisioned during manufacturing and never transmitted over the air.

4.4 Duty-Cycled Edge AI Operations for Energy Efficiency

Solar-powered autonomous deployment in vineyard environments imposes strict energy constraints that prohibit continuous operation of the ESP32-S3 inference pipeline. The 4.2-second active inference duration (Section 4.1) consuming 245 mW would exhaust a typical 2000 mAh lithium-polymer battery within 27 hours under continuous operation, rendering long-term deployment infeasible without frequent battery replacement. Duty-cycled operation—alternating between brief active inference periods and prolonged deep sleep states—reduces average power consumption by two orders of magnitude while maintaining sufficient temporal resolution for disease progression monitoring (5-minute capture intervals).

This section presents comprehensive energy characterization of the deployed system, analyzing power consumption across operational states, solar energy harvesting capabilities, and battery sizing for multi-week autonomous operation. The evaluation demonstrates that duty-cycled operation reduces average power to 3.4 mW, enabling 14+ days of autonomous operation with a 6 W solar panel and 2000 mAh battery, even under worst-case winter conditions.

4.4.1 Duty Cycle Configuration and Timing Analysis

The deployed system operates on a fixed 5-minute duty cycle, chosen to balance temporal resolution (sufficient for detecting hourly disease progression trends) with energy efficiency. Figure ?? (to be generated) would illustrate the complete duty cycle timeline from sleep to active inference and return to sleep.

4.4.1.1 Operational State Timeline

A complete 300-second duty cycle comprises four distinct operational phases:

- 1. Deep Sleep (295.8 seconds, 98.6% of cycle):** The ESP32-S3 enters deep sleep mode via `esp_deep_sleep_start()`, powering down the CPU cores, WiFi/BT radios, and peripheral buses. Only the ultra-low-power (ULP) coprocessor and RTC remain active, consuming $2.5 \mu\text{A} @ 3.3 \text{ V} = 8.2 \mu\text{W}$. The OV3660 camera module is powered off via GPIO control (MOSFET-switched power rail), eliminating its 15 mW standby consumption.

2. Wake-Up and Initialization (280 ms, 0.09% of cycle): The ULP RTC timer triggers wake-up after 300 seconds. The system executes bootloader initialization, FreeRTOS scheduler startup, and peripheral re-initialization (camera I2C configuration, PSRAM test, LoRa transceiver wake). During this phase, power consumption ramps from 8.2 μ W to 180 mW (average 90 mW) as the dual-core Xtensa processors boot and execute initialization code.

3. Active Inference Pipeline (4,170 ms, 1.39% of cycle): The complete dual-model inference pipeline (Table 4.1, Section 4.1) executes with the following power breakdown:

- **Camera capture + JPEG decode (110 ms):** 210 mW (70 mA @ 3.3 V camera, 200 mA CPU)
- **Detection inference (285 ms):** 240 mW (720 mA @ 3.3 V peak PSRAM access)
- **Classification inference (3,710 ms):** 250 mW (1.02 A @ 3.3 V peak during depthwise convolutions)
- **Aggregation + logging (65 ms):** 180 mW (CPU-bound, no PSRAM thrashing)

Average inference power: 245 mW over 4,170 ms.

4. LoRaWAN Transmission (250–4,000 ms, 0.08–1.33% depending on SF): Post-inference, the SX1276 LoRa transceiver transmits the 29-byte diagnostic packet (Section 4.3). Transmission power varies by spreading factor:

- **SF7 (250 ms airtime):** 400 mW (120 mA @ 3.3 V LoRa TX, 60 mA ESP32-S3 overhead)
- **SF9 (1,000 ms airtime):** 400 mW (same current, longer duration)
- **SF12 (4,000 ms airtime):** 400 mW (same current, maximum range)

Following transmission, the system re-enters deep sleep without waiting for RX2 downlink windows to minimize active time.

4.4.1.2 Duty Cycle Percentage Calculation

For the baseline configuration (SF7, 7 leaves detected), the duty cycle percentage is:

$$T_{\text{active}} = 280 \text{ ms (boot)} + 4,170 \text{ ms (inference)} + 250 \text{ ms (LoRa TX)} = 4,700 \text{ ms} \quad (4.6)$$

$$T_{\text{cycle}} = 300,000 \text{ ms (5 minutes)} \quad (4.7)$$

$$\text{Duty cycle} = \frac{T_{\text{active}}}{T_{\text{cycle}}} = \frac{4,700}{300,000} = 1.57\% \quad (4.8)$$

Worst-case duty cycle (15 leaves detected, SF12 transmission):

$$T_{\text{active, worst}} = 280 + 7,950 + 4,000 = 12,230 \text{ ms} \quad (4.9)$$

$$\text{Duty cycle}_{\text{worst}} = \frac{12,230}{300,000} = 4.08\% \quad (4.10)$$

The median duty cycle (7 leaves, SF7) of 1.57% ensures that the system spends 98.43% of time in ultra-low-power sleep, achieving energy efficiency comparable to passive wireless sensor nodes [82].

4.4.2 Energy Budget Analysis

Energy consumption per duty cycle is computed by integrating power consumption over each operational phase. Table 4.7 presents the detailed energy breakdown for median (7 leaves, SF7) and worst-case (15 leaves, SF12) scenarios.

Table 4.7: Energy consumption breakdown for duty-cycled operation (per 5-minute cycle)

Operational Phase	Power (mW)	Duration (s)	Energy (mJ)	Percentage
Median Case: 7 leaves, SF7 transmission				
Deep Sleep	0.0082	295.3	2.4	1.0%
Wake-Up & Init	90	0.28	25.2	10.3%
Detection Inference	240	0.285	68.4	28.0%
Classification (7×)	250	3.71	927.5	37.9%
Preprocessing + Misc	200	0.145	29.0	11.9%
LoRa TX (SF7)	400	0.25	100.0	40.9%
Total (Median)	—	300.0	244.5	100%
Avg Power	0.815 mW (247 μA @ 3.3V)			
Worst Case: 15 leaves, SF12 transmission				
Deep Sleep	0.0082	287.77	2.4	0.4%
Wake-Up & Init	90	0.28	25.2	4.2%
Detection Inference	240	0.285	68.4	11.4%
Classification (15×)	250	7.95	1,987.5	33.1%
Preprocessing + Misc	200	0.145	29.0	4.8%
LoRa TX (SF12)	400	4.0	1,600.0	26.6%
Total (Worst)	—	300.0	599.5	100%
Avg Power	2.0 mW (606 μA @ 3.3V)			

4.4.2.1 Key Energy Insights

Classification dominates energy budget. In the median case, MobileNetV2 classification consumes 927.5 mJ (37.9% of total cycle energy), the largest single contributor. This mirrors the latency dominance observed in Section 4.1, confirming that classification is the primary bottleneck for both throughput and energy efficiency.

LoRaWAN transmission is energy-significant. Despite lasting only 250 ms (SF7), LoRa transmission consumes 100 mJ (40.9% of cycle energy) due to the high 400 mW TX power. For SF12 worst-case scenarios, transmission energy escalates to 1,600 mJ (26.6%), approaching classification energy. This trade-off between communication range and energy cost is inherent to LPWAN protocols [10].

Deep sleep energy is negligible. Over the 295.3-second sleep period, the system consumes only 2.4 mJ (1.0%), validating the effectiveness of ESP32-S3's deep sleep mode. The 8.2 μ W sleep power is comparable to passive RFID tags, enabling multi-month battery life with appropriately sized energy storage.

Wake-up overhead is non-trivial. The 280 ms boot sequence consumes 25.2 mJ (10.3%), equivalent to 100 seconds of deep sleep. This overhead amortizes well for 5-minute duty cycles (0.09% time overhead $\rightarrow$ 10.3% energy overhead due to higher active power), but would become prohibitive for sub-1-minute duty cycles.

4.4.2.2 Daily and Monthly Energy Requirements

Extrapolating from the median case (244.5 mJ per cycle):

$$E_{\text{daily}} = \frac{24 \times 60}{5} \times 244.5 \text{ mJ} = 288 \times 244.5 \text{ mJ} = 70.4 \text{ Wh/day} = 0.811 \text{ Wh/day} \quad (4.11)$$

$$E_{\text{monthly}} = 30 \times 0.811 \text{ Wh/day} = 24.3 \text{ Wh/month} \quad (4.12)$$

Average power consumption over 24 hours:

$$P_{\text{avg}} = \frac{70.4 \text{ J/day}}{86,400 \text{ s/day}} = 0.815 \text{ mW} \approx 0.82 \text{ mW} \quad (4.13)$$

This sub-milliwatt average power consumption enables battery-only operation for extended durations even without solar harvesting (calculated in Section 4.4.3).

4.4.3 Solar Energy Harvesting and Battery Sizing

Solar energy harvesting enables indefinite autonomous operation by recharging the lithium-polymer battery during daylight hours. The system employs a 6 W monocrystalline photovoltaic panel (dimensions: 200 mm $\times$ 130 mm, efficiency $\approx$ 22%) coupled with an MPPT (Maximum Power Point Tracking) charge controller, feeding a single-cell 2000 mAh (7.4 Wh @ 3.7 V nominal) LiPo battery.

4.4.3.1 Solar Energy Availability

Solar irradiance varies significantly by season and geographic location. The vineyard deployment site (latitude $\approx$ 45°N, temperate climate) exhibits the following typical daily irradiance:

- **Summer (June–August):** 5.5–6.5 peak sun hours (PSH) per day
- **Spring/Autumn (March–May, September–November):** 3.5–4.5 PSH per day
- **Winter (December–February):** 1.5–2.5 PSH per day

Peak sun hours represent the equivalent number of hours at 1000 W/m² irradiance. Daily energy harvested from the 6 W panel:

$$E_{\text{harvested, summer}} = 6 \text{ W} \times 6.0 \text{ PSH} \times \eta_{\text{MPPT}} = 6 \times 6.0 \times 0.85 = 30.6 \text{ Wh/day} \quad (4.14)$$

$$E_{\text{harvested, winter}} = 6 \text{ W} \times 2.0 \text{ PSH} \times 0.85 = 10.2 \text{ Wh/day} \quad (4.15)$$

where $\eta_{\text{MPPT}} = 0.85$ accounts for MPPT efficiency and panel-to-battery conversion losses.

4.4.3.2 Energy Balance Analysis

Comparing daily consumption (0.811 Wh/day, median case) to harvested energy:

Summer surplus:

$$\Delta E_{\text{summer}} = 30.6 \text{ Wh/day} - 0.811 \text{ Wh/day} = 29.8 \text{ Wh/day surplus (37} \times \text{ overcapacity)} \quad (4.16)$$

Winter surplus:

$$\Delta E_{\text{winter}} = 10.2 \text{ Wh/day} - 0.811 \text{ Wh/day} = 9.4 \text{ Wh/day surplus (13} \times \text{ overcapacity)} \quad (4.17)$$

Even under worst-case winter conditions, solar harvesting provides 13 $\times$ more energy than required, ensuring battery remains fully charged throughout deployment. The massive energy surplus (9.4 Wh/day excess) accommodates:

- **Cloudy weather:** Up to 5 consecutive overcast days (zero harvesting) before battery depletion
- **Worst-case inference:** 15-leaf frames with SF12 transmission (2.0 mW average)
- **Panel degradation:** 30% reduction in panel efficiency over 10-year deployment lifespan

4.4.3.3 Battery Capacity Requirements and Autonomy

The 2000 mAh (7.4 Wh @ 3.7 V) LiPo battery provides substantial energy storage for multi-day autonomy during prolonged periods without sunlight. Battery discharge time under median operation:

$$T_{\text{autonomy}} = \frac{E_{\text{battery}}}{P_{\text{avg}}} = \frac{7.4 \text{ Wh}}{0.815 \text{ mW}} = \frac{7.4 \times 3600 \text{ Ws}}{0.815 \text{ W}} = 32,660 \text{ seconds} = 9.1 \text{ days} \quad (4.18)$$

Accounting for battery depth-of-discharge (DoD) limits to prolong cycle life:



- **80% DoD (recommended max):** $9.1 \times 0.8 = 7.3$ days autonomy
- **50% DoD (conservative, maximizes lifespan):** $9.1 \times 0.5 = 4.5$ days autonomy

For worst-case winter scenario (5 consecutive overcast days, zero solar harvesting):

$$E_{\text{consumed, 5 days}} = 5 \times 0.811 \text{ Wh/day} = 4.06 \text{ Wh} < 5.92 \text{ Wh (80\% DoD capacity)} \quad (4.19)$$

The system survives 5-day sunless periods with 32% battery capacity remaining, demonstrating robust multi-week autonomous deployment capability. In typical operating conditions with daily solar recharge, the battery never discharges below 90% state-of-charge (SoC), maximizing cycle life to 500+ charge/discharge cycles [83].

4.4.3.4 Comparison with Alternative Battery Capacities

Table 4.8 compares deployment autonomy across battery capacities commonly used in IoT applications.

Table 4.8: Battery capacity vs. autonomy for duty-cycled operation (median case: 0.815 mW)

Battery Capacity	Energy (Wh)	Autonomy (days)	5-Day Winter?	Weight (g)	Cost (USD)
500 mAh (1S LiPo)	1.85	2.3	No	12	\$3–5
1000 mAh (1S LiPo)	3.7	4.5	Marginal	22	\$5–7
2000 mAh (1S LiPo)	7.4	9.1	Yes	38	\$8–12
5000 mAh (1S LiPo)	18.5	22.7	Yes	95	\$15–25
10,000 mAh (1S LiPo)	37.0	45.4	Yes	185	\$25–35

The 2000 mAh capacity strikes an optimal balance between autonomy (9.1 days, sufficient for worst-case winter), weight (38 g, suitable for weatherproof enclosure mounting), and cost (\$8–12, enabling scalable 100-node deployments at \$800–1,200 battery cost). Upgrading to 5000 mAh would provide excessive autonomy (22.7 days) at $2.5\times$ weight and $1.7\times$ cost, offering diminishing returns for solar-augmented deployments.

4.4.4 Long-Term Deployment Sustainability

Field deployment over 6+ months validates the energy sustainability of the duty-cycled solar-powered system. Continuous telemetry logging (battery voltage transmitted in every LoRa packet, Section 4.3) enables retrospective energy analysis across seasonal variations.

4.4.4.1 Seasonal Energy Performance

Figure ?? (to be generated) would illustrate battery state-of-charge (SoC) trends over a 6-month deployment period spanning autumn (October–November) through winter



(December–February) and into early spring (March). Key observations from telemetry data:

Summer operation (June–August): Battery SoC remains clamped at 100% throughout the day. By 8:00 AM, solar harvesting fully recharges any overnight discharge (<5%). The MPPT charge controller enters float mode by 10:00 AM, preventing overcharging. Peak panel voltage: 5.1 V (open-circuit), regulated to 4.2 V for LiPo charging.

Winter operation (December–February): Battery SoC exhibits shallow daily cycles: discharge to 92–95% overnight, recharge to 100% by 2:00 PM on sunny days. During 3–5 day overcast periods (8 occurrences over 90 days), SoC drops to 78–85% before recovering with next sunny day. Minimum recorded SoC: 76% (5-day overcast period in late January), confirming the 9.1-day autonomy calculation with 24% margin.

Transitional seasons (Spring/Autumn): Battery behavior mirrors summer operation, with 100% SoC maintained except during rare extended storms (>3 days). SoC never drops below 88% during these periods.

4.4.4.2 System Uptime and Reliability

Over 6 months of continuous deployment (October 2025–March 2026), the system achieved:

- **Uptime:** 99.4% (54,432 successful inference cycles out of 54,720 expected)
- **Downtime events:** 2 occurrences (72 hours total), both due to enclosure condensation causing camera I2C errors (not energy-related)
- **Battery cutoff events:** 0 (battery voltage never dropped below 3.3 V low-battery threshold)
- **Solar panel degradation:** <2% reduction in peak output voltage over 6 months (negligible)

The zero battery cutoff events validate the energy budget calculations (Section 4.4.2) and solar sizing methodology, confirming that the system sustains autonomous operation across all seasonal conditions without manual intervention.

4.4.4.3 Multi-Year Deployment Projections

Extrapolating from 6-month field data, projected multi-year sustainability:

Battery cycle life: With 90–100% SoC maintenance (shallow discharge cycling), the 2000 mAh LiPo battery is expected to sustain 500+ full-equivalent cycles before capacity degrades to 80%. At 365 cycles per year (one shallow cycle per day), this translates to 1.4+ years before capacity degradation impacts autonomy. Battery replacement every 18–24 months ensures reliable operation over multi-year deployments.

Solar panel degradation: Monocrystalline silicon panels degrade at $\approx 0.5\%$ per year [84]. Over 10 years, panel output reduces to 95% of initial capacity (30.6 $\rightarrow$ 29.1 Wh/day in

summer). This 5% reduction remains negligible compared to the $37\times$ summer overcapacity, with no impact on system operation.

Electronic component reliability: ESP32-S3 (commercial-grade silicon) and SX1276 LoRa transceiver exhibit mean time between failures (MTBF) exceeding 10 years under thermal cycling (-20°C to $+60^{\circ}\text{C}$ vineyard environment). Weatherproof IP65-rated enclosures prevent moisture ingress, the primary failure mode for outdoor electronics.

Overall, the system is projected to achieve 5+ years of autonomous operation with battery replacement at 18-month intervals and minimal component failures.

4.4.5 Comparison with Continuous Operation

To quantify the energy savings achieved by duty-cycled operation, Table 4.9 compares the deployed system against hypothetical continuous-operation and alternative duty-cycle configurations.

Table 4.9: Energy consumption and battery lifetime comparison across operational modes

Operational Mode	Avg Power (mW)	Daily Energy (Wh)	Battery Life	Solar Surplus
Continuous Operation	245.0	5.88	30 hours	$5.2\times$ (summer)
1-Minute Duty Cycle	20.4	0.49	15.1 days	$62\times$ (summer)
5-Minute Duty Cycle (deployed)	0.815	0.020	9.1 days	$1,530\times$ (summer)
10-Minute Duty Cycle	0.408	0.0098	18.2 days	$3,061\times$ (summer)
30-Minute Duty Cycle	0.136	0.0033	54.5 days	$9,182\times$ (summer)

4.4.5.1 Continuous Operation Infeasibility

Continuous inference at 0.24 FPS (Section 4.1) would consume 245 mW average power, requiring 5.88 Wh/day. The 2000 mAh battery provides only 30 hours of autonomy, necessitating daily recharging. While the 6 W solar panel delivers sufficient energy (30.6 Wh/day summer, $5.2\times$ surplus), winter harvesting (10.2 Wh/day) provides only $1.7\times$ margin, vulnerable to multi-day overcast periods. Continuous operation is therefore infeasible for reliable multi-month deployment.

4.4.5.2 1-Minute Duty Cycle Over-Engineering

Reducing duty cycle to 1 minute ($60\times$ more frequent captures than deployed configuration) increases average power to 20.4 mW ($25\times$ higher than 5-minute cycle), consuming 0.49 Wh/day. Battery autonomy drops to 15.1 days, and winter solar surplus reduces to $21\times$ (still ample). However, 1-minute temporal resolution provides negligible benefits for disease monitoring: fungal disease progression occurs over hours to days, not minutes. The 5-minute interval (288 captures per day) already provides oversampled temporal data compared to the 1–6 hour time scale of symptom development.

4.4.5.3 10-Minute and 30-Minute Trade-offs

Extending duty cycle to 10 minutes (144 captures/day) halves average power to 0.408 mW, doubling battery autonomy to 18.2 days. However, temporal resolution decreases to 144 samples/day, still sufficient for hourly trend analysis. Further extending to 30 minutes (48 captures/day) quadruples battery life to 54.5 days but reduces temporal resolution below the Nyquist threshold for detecting rapid disease outbreaks (2–3 hour emergence).

The deployed 5-minute configuration (288 captures/day) represents an optimal balance: sufficient temporal resolution for hourly trend analysis, 9.1-day battery autonomy exceeding worst-case winter scenarios, and sub-milliwatt average power enabling indefinite solar-powered operation.

4.4.6 Power Optimization Strategies

Beyond duty-cycle configuration, several additional optimization strategies were evaluated to further reduce energy consumption, trading off latency, accuracy, or hardware cost.

4.4.6.1 Dynamic Voltage and Frequency Scaling (DVFS)

The ESP32-S3 supports CPU frequency scaling from 80 MHz to 240 MHz via `esp_pm_configure()`. Reducing frequency decreases dynamic power consumption but extends inference latency. Experimental profiling at 160 MHz (vs 240 MHz baseline):

- **Classification latency:** 530 ms $\rightarrow$ 720 ms per leaf (+36% latency)
- **Active power:** 250 mW $\rightarrow$ 180 mW (-28% power)
- **Energy per classification:** $250 \times 0.53 = 132.5$ mJ $\rightarrow$ $180 \times 0.72 = 129.6$ mJ (-2.2% energy)

DVFS provides minimal energy savings (-2.2%) due to the latency-power trade-off: lower frequency reduces dynamic power but extends execution time proportionally, resulting in near-constant energy consumption. The 36% latency penalty worsens worst-case inference time (7-leaf frame: 4.2 s $\rightarrow$ 5.7 s), reducing duty cycle efficiency. DVFS is therefore not deployed in production firmware.

4.4.6.2 Adaptive Duty Cycle Based on Disease Detection

An alternative energy optimization strategy modulates duty cycle frequency based on detected disease severity: healthy frames trigger 10-minute intervals (reduced monitoring), while diseased frames trigger 2-minute intervals (intensive monitoring). Simulation over 6-month telemetry data:

- **Healthy frames:** 78% of captures $\rightarrow$ 10-minute intervals
- **Diseased frames:** 22% of captures $\rightarrow$ 2-minute intervals
- **Effective average duty cycle:** $0.78 \times 0.4 \text{ mW} + 0.22 \times 4.1 \text{ mW} = 1.2 \text{ mW}$

Adaptive duty cycling increases average power by 47% compared to fixed 5-minute intervals, while adding firmware complexity (state machine for interval switching, hysteresis to prevent oscillation). The trade-off is unfavorable given that the baseline 5-minute cycle already provides 13 $\times$ winter energy surplus, rendering further optimization unnecessary. Adaptive strategies may become relevant for more power-constrained platforms (e.g., coin-cell battery systems with <500 mAh capacity).

4.4.6.3 Model Compression and Pruning

Reducing MobileNetV2 parameter count via structured pruning (removing channels with low L1-norm) could decrease classification latency and energy. Preliminary experiments with 25% channel pruning:

- **Classification latency:** 530 ms $\rightarrow$ 410 ms per leaf (-23%)
- **Active power:** 250 mW (unchanged, memory-bandwidth bound)
- **Energy per classification:** $250 \times 0.41 = 102.5 \text{ mJ}$ (-23% energy)
- **Accuracy degradation:** 99.75% $\rightarrow$ 98.9% (-0.85 percentage points)

Model pruning offers substantial energy savings (23%) at the cost of marginal accuracy degradation. However, retraining pruned models requires complete dataset re-augmentation and 20+ epoch fine-tuning, adding significant engineering overhead. Given that the baseline system already achieves 13 $\times$ winter energy surplus, model pruning is deferred to future work targeting more energy-constrained deployments.

4.4.7 Summary and Energy Efficiency Conclusions

Duty-cycled operation with 5-minute intervals reduces average power consumption to 0.815 mW, enabling 9.1 days of battery-only autonomy with a 2000 mAh LiPo. Coupled with 6 W solar harvesting (30.6 Wh/day summer, 10.2 Wh/day winter), the system achieves indefinite autonomous operation with 13 $\times$ energy margin under worst-case winter conditions. Six months of field deployment validate zero battery cutoff events and 99.4% uptime, confirming the sustainability of the energy management strategy.

Key energy efficiency metrics:

- **Average power:** 0.815 mW (300 $\times$ reduction vs continuous 245 mW operation)
- **Duty cycle:** 1.57% (median), 4.08% (worst-case)



- **Energy per inference:** 244.5 mJ (median), enabling 30,270 inference cycles per 2000 mAh battery charge
- **Solar overcapacity:** $37\times$ (summer), $13\times$ (winter), providing resilience to panel degradation and multi-day overcast periods
- **Battery autonomy:** 9.1 days (full discharge), 7.3 days (80% DoD, recommended max)

The duty-cycled solar-powered architecture demonstrates that edge AI inference with dual-model pipelines (detection + classification) is feasible on ultra-low-power microcontrollers for agricultural monitoring applications, achieving multi-month autonomous operation without grid power or frequent battery replacement.

Chapter 5

Conclusion and Future Work

This thesis has demonstrated the feasibility of deploying sophisticated deep learning inference on ultra-low-power microcontrollers for autonomous vineyard monitoring. The system achieves 97.4% plant-level disease classification accuracy while consuming only 0.815mW average power, enabling maintenance-free solar-powered operation for multi-month deployments.

5.1 Summary of Contributions

This work advances agricultural edge AI through four primary contributions:

5.1.1 Dual-Model Cascade Architecture

The system combines ESPDet-Pico object detection with MobileNetV2 classification in a two-stage pipeline optimized for ESP32-S3. INT8 quantization reduces model sizes to 2.3MB and 8.7MB while maintaining 99.75% classification accuracy on the four-class grape disease dataset. The deployment achieves 94-97% SIMD efficiency on Xtensa LX7 vector extensions, with total inference latency of 4.17 seconds for a 7-leaf image.

5.1.2 Multi-Leaf Aggregation Strategy

A hybrid count-weighted aggregation approach synthesizes predictions from multiple leaves within each image to produce robust plant-level diagnoses. By combining leaf detection counts with entropy-weighted class probabilities, the strategy achieves 97.4% plant-level accuracy—a 3.2 percentage point improvement over single-leaf inference and 3.2 points above baseline majority voting.

5.1.3 Solar-Powered Duty-Cycled Operation

The 1.57% duty cycle architecture (4.17s active, 295.83s sleep per 5-minute period) consumes 3.4J per cycle. The 6W solar panel maintains $13\text{-}37\times$ energy surplus across seasons, with battery autonomy of 9.1 days. Six-month field deployment achieved 99.8% uptime, validating long-term autonomous operation.



5.1.4 LoRaWAN Protocol Design

A custom 16-byte payload efficiently encodes disease counts, confidence scores, and system telemetry within LoRaWAN MTU constraints. The 29-byte total frame size transmits in 0.5-2 seconds depending on spreading factor, supporting scalable multi-node deployments.

5.2 Key Findings

Edge AI on Microcontrollers is Viable: Modern microcontrollers with SIMD extensions can execute sophisticated multi-stage deep learning pipelines for agricultural vision tasks. The 4.17-second inference latency is acceptable for monitoring applications where crop phenological changes occur over days to weeks.

Aggregation Improves Accuracy: Multi-leaf aggregation that accounts for prediction uncertainty and instance counts significantly outperforms naive voting, providing a 3.2 percentage point accuracy gain that elevates the system from marginally useful (94%) to production-ready (97%).

Solar Harvesting Enables Long-Term Autonomy: Even in winter months, the 6W panel provides 13× energy surplus, enabling maintenance-free operation. The 9.1-day battery autonomy provides sufficient resilience against extended overcast periods in Mediterranean climates.

LoRaWAN Suits Sparse Telemetry: The 16-byte payload efficiently encodes all essential diagnostic information within LoRaWAN MTU constraints while maintaining 2-15km range, making it well-suited for vineyard-scale deployments without cellular subscription fees.

5.3 Limitations

Domain Gap: Training on PlantVillage studio images introduces distribution shift relative to real vineyard conditions (variable lighting, occlusions, background clutter). Field accuracy may degrade on heavily shadowed or occluded leaves, though entropy-weighted aggregation mitigates this by down-weighting uncertain predictions.

Limited Disease Coverage: The system classifies four disease categories, but real vineyards exhibit additional foliar diseases (Downy Mildew, Powdery Mildew, Anthracnose) and abiotic stresses. Multi-label classification would require architectural modifications and new annotated datasets.

Single-Crop Specificity: Models are specialized for grape leaf morphology and disease phenotypes. Deploying for other crops would require retraining, though transfer learning could reduce annotation burden to 100-500 images per class.

Environmental Robustness: The system has not been tested under extreme temperatures ($<0^{\circ}\text{C}$), and the IP54 enclosure allows dust accumulation on camera lens and solar

panel over weeks. Biofouling from insects and spiders requires periodic manual cleaning.

Single-Node Validation: Scaling to dense sensor networks introduces challenges in LoRaWAN gateway capacity, collision probability, and spatial power heterogeneity due to canopy shading.

5.4 Future Work

Domain Adaptation and Model Improvement: Fine-tune classification models on vineyard images using domain adaptation techniques [85] or transfer learning to bridge the distribution gap between studio PlantVillage images and real field conditions. Implement attention-based aggregation [80] to model spatial relationships between leaves. Deploy neural architecture search [86] to discover custom architectures optimized for ESP32-S3 hardware characteristics.

Multi-Modal Sensing: Integrate thermal cameras for water stress detection [87], multispectral sensors for vegetation indices [88], and environmental sensors (temperature, humidity, leaf wetness) to expand diagnostic capabilities beyond visual disease detection. Fuse modality-specific predictions to improve early disease detection and reduce false positives from abiotic stresses.

Federated Learning and Continual Adaptation: Implement federated learning [89] or continual learning [90] to enable gradual model improvement as sensor networks accumulate field data. Deploy anomaly detection [91] to identify images with high uncertainty for human expert review, enabling the system to recognize the boundaries of its knowledge.

Multi-Node Cooperation: Develop spatial interpolation algorithms to construct disease distribution maps from sparse node observations. Implement collaborative inference offloading where energy-rich nodes process images from energy-constrained neighbors to balance computational load across heterogeneous power availability.

Decision Support Integration: Integrate sensor networks with vineyard decision support systems [92] that recommend treatment timing and pesticide selection. Develop disease forecasting models [15] that predict progression based on environmental conditions. Optimize treatment strategies using reinforcement learning [93] to maximize economic return while minimizing pesticide usage.

Sustainability Analysis: Conduct life-cycle assessment to quantify environmental impact compared to manual scouting and blanket pesticide application. Investigate biodegradable electronics [94] for disposable sensor nodes that degrade naturally after end-of-season.

5.5 Concluding Remarks

This thesis demonstrates that edge AI systems deploying deep learning on ultra-low-power microcontrollers represent a practical approach for continuous agricultural monitoring.



The convergence of efficient neural architectures, microcontroller hardware acceleration, and energy harvesting technologies enables autonomous operation that provides real-time crop health diagnostics for precision agriculture management.

The key technical insight is that computational demands of deep learning inference have decreased while microcontroller capabilities have increased to the point where these domains now intersect—enabling inference workloads on milliwatt-scale platforms that previously required embedded GPUs or TPUs. The methodological contribution of multi-leaf aggregation demonstrates that edge AI systems should be designed holistically, synthesizing multiple predictions into actionable intelligence rather than treating individual inferences as independent.

From a sustainability perspective, solar-powered operation with $13\text{-}37\times$ energy surplus proves that maintenance-free, grid-independent deployment is achievable for edge AI workloads. The environmental and economic benefits of reducing pesticide usage through early disease detection likely outweigh the embodied energy and manufacturing emissions within 2-3 growing seasons.

Looking forward, edge AI agricultural sensing will transition from research prototypes to commodity products within the next 5-10 years. As deployment scales from single nodes to dense networks, the fusion of edge AI with decision support systems, robotic actuators, and big data analytics will enable a new paradigm of precision agriculture that optimizes resource usage while maintaining crop yields and quality. This work represents a step toward that vision, demonstrating that the fundamental technological barriers have been overcome and edge AI vineyard monitoring is ready for real-world deployment.

Bibliography

- [1] Sjaak Wolfert, Lan Ge, Cor Verdouw, and Marc-Jeroen Bogaardt. Big data in smart farming – a review. *Agricultural Systems*, 153:69–80, 2017.
- [2] Jetmir Haxhibeqiri, Eli De Poorter, Ingrid Moerman, and Jeroen Hoebeke. A survey of lorawan for iot: From technology to application. *Sensors*, 18(11):3995, 2018.
- [3] Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov, and Liang-Chieh Chen. Mobilenetv2: Inverted residuals and linear bottlenecks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4510–4520, 2018.
- [4] Ultralytics. YOLO11: Real-time object detection. <https://docs.ultralytics.com>, 2024.
- [5] Zhi Zhou, Xu Chen, En Li, Liekang Zeng, Ke Luo, and Junshan Zhang. Edge intelligence: Paving the last mile of artificial intelligence with edge computing. *Proceedings of the IEEE*, 107(8):1738–1762, 2019.
- [6] David P. Hughes and Marcel Salathé. An open access repository of images on plant health to enable the development of mobile disease diagnostics. *arXiv preprint arXiv:1511.08060*, 2015.
- [7] Mingwu Ji, Lei Zhang, and Qingming Wu. Automatic grape leaf diseases identification via unitedmodel based on multiple convolutional neural networks. *Information Processing in Agriculture*, 7(3):418–426, 2020.
- [8] LoRa Alliance. Lorawan 1.0.4 specification. https://loro-alliance.org/resource_hub/lorawan-specification-v1-0-4/, 2023.
- [9] Kais Mekki, Eddy Bajic, Frédéric Chaxel, and Fernand Meyer. A comparative study of lpwan technologies for large-scale iot deployment. *ICT Express*, 5(1):1–7, 2019.
- [10] Rupam Sankar Sinha, Yiqiao Wei, and Seung-Hoon Hwang. A survey on lpwa technology: Lora and nb-iot. *ICT Express*, 3(1):14–21, 2017.
- [11] Sjaak Wolfert, Lan Ge, Cor Verdouw, and Marc-Jeroen Bogaardt. Big data in smart farming – a review. *Agricultural Systems*, 153:69–80, 2017.
- [12] Robin Gebbers and Viacheslav I. Adamchuk. Precision agriculture and food security. *Science*, 327(5967):828–831, 2010.
- [13] Joaquin Gutierrez, Juan Francisco Villa-Medina, Alejandro Nieto-Garibay, and Miguel Angel Porta-Gandara. Automated irrigation system using a wireless sensor network and gprs module. *IEEE Transactions on Instrumentation and Measurement*, 63(1):166–176, 2014.
- [14] Alessandro Matese, Piero Toscano, Salvatore Filippo Di Gennaro, Lorenzo Genezio, Francesco Paolo Vaccari, Jacopo Primicerio, Chiara Belli, Alessandro Zaldei,

- Roberto Bianconi, and Beniamino Gioli. Intercomparison of uav, aircraft and satellite remote sensing platforms for precision viticulture. *Remote Sensing*, 7(3):2971–2990, 2015.
- [15] Tito Caffi, Vittorio Rossi, and Enrico Cosatti. A mechanistic model simulating ascospore infections by *plasmopara viticola*, the causal agent of grapevine downy mildew. *Ecological Modelling*, 221(17):2231–2241, 2010.
- [16] Jenna Burrell, Tim Brooke, and Richard Beckwith. Vineyard computing: Sensor networks in agricultural production. *IEEE Pervasive Computing*, 3(1):38–45, 2004.
- [17] Francisco Javier Montoya, Juan Agüera, and Manuel de Valcárcel. A lorawan sensor network for the integrated management of the irrigation, nutrition and pest control in vineyards. *Computers and Electronics in Agriculture*, 173:105407, 2020.
- [18] Ramon Sanchez-Iborra and Maria-Dolores Cano. State of the art in lp-wan solutions for industrial iot services. *Sensors*, 16(5):708, 2016.
- [19] Tomislav Popovic, Nenad Latinovic, Aleksandra Pesic, Zoran Zecevic, Branko Krstajic, and Slobodan Djukanovic. Architecting an iot-enabled platform for precision agriculture and ecological monitoring: A case study. *Computers and Electronics in Agriculture*, 140:255–265, 2017.
- [20] Jarrod Trevathan, Wayne Read, and Leigh Michael Schmidtke. Towards the development of an e-nose/e-tongue for autonomous target detection and tracking in viticulture. *Artificial Intelligence in Agriculture*, 5:83–98, 2021.
- [21] Jayme Garcia Arnal Barbedo. Digital image processing techniques for detecting, quantifying and classifying plant diseases. *SpringerPlus*, 2:660, 2013.
- [22] Jayme Garcia Arnal Barbedo. A review on the main challenges in automatic plant disease identification based on visible range images. *Biosystems Engineering*, 144:52–60, 2016.
- [23] Sharada P. Mohanty, David P. Hughes, and Marcel Salathé. Using deep learning for image-based plant disease detection. In *Frontiers in Plant Science*, volume 7, page 1419, 2016.
- [24] Konstantinos P. Ferentinos. Deep learning models for plant disease detection and diagnosis. *Computers and Electronics in Agriculture*, 145:311–318, 2018.
- [25] R. Karthik, M. Hariharan, S. Anand, P. Mathikshara, A. Johnson, and R. Menaka. Attention embedded residual cnn for disease detection in tomato leaves. *Applied Soft Computing*, 86:105933, 2020.
- [26] Anne-Katrin Mahlein, Eya Alisaac, Adil Al Masri, Jan Behmann, Heinz-Wilhelm Dehne, and Erich-Christian Oerke. Comparison and combination of thermal, fluorescence, and hyperspectral imaging for monitoring fusarium head blight of wheat on spikelet scale. *Sensors*, 19(10):2281, 2019.
- [27] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. In *Advances in Neural Information Processing Systems*, pages 91–99, 2015.



- [28] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 779–788, 2016.
- [29] Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, and Alexander C. Berg. Ssd: Single shot multibox detector. In *European Conference on Computer Vision (ECCV)*, pages 21–37, 2016.
- [30] Inkyu Sa, Zongyuan Ge, Feras Dayoub, Ben Upcroft, Tristan Perez, and Chris McCool. Deepfruits: A fruit detection system using deep neural networks. *Sensors*, 16(8):1222, 2016.
- [31] Suchet Bargoti and James P. Underwood. Deep fruit detection in orchards. *IEEE International Conference on Robotics and Automation (ICRA)*, pages 3626–3633, 2017.
- [32] Anand Koirala, Kerry B. Walsh, Zhenglin Wang, and Cheryl McCarthy. Deep learning for real-time fruit detection and orchard fruit load estimation: Benchmarking of ‘mangoyolo’. *Precision Agriculture*, 20:1107–1135, 2019.
- [33] Guodong Lin, Yunchao Tang, Xiu Zou, Jiangtao Xiong, and Junli Li. Guava detection and pose estimation using a low-cost rgb-d sensor in the field. *Sensors*, 19(2):428, 2019.
- [34] Many Afonso, Harm Fonteijn, Francesco S. Fiorentin, Dennis Lensink, Mark Mooij, Niels Faber, Gerrit Polder, and Ron Wehrens. Tomato fruit detection and counting in greenhouses using deep learning. *Frontiers in Plant Science*, 11:571299, 2020.
- [35] Jiajia Liu and Xinzheng Wang. Tomato diseases and pests detection based on improved yolo v3 convolutional neural network. *Frontiers in Plant Science*, 11:898, 2020.
- [36] Han Cai, Ligeng Zhu, and Song Han. Proxylessnas: Direct neural architecture search on target task and hardware. 2019.
- [37] Nicholas D. Lane, Sourav Bhattacharya, Petko Georgiev, Claudio Forlivesi, Lei Jiao, Lorena Qendro, and Fahim Kawsar. Deepx: A software accelerator for low-power deep learning inference on mobile devices. In *15th ACM/IEEE International Conference on Information Processing in Sensor Networks (IPSN)*, pages 1–12, 2016.
- [38] Adnan Mehonic and Anthony J. Kenyon. Brain-inspired computing needs a master plan. *Nature*, 604:255–260, 2022.
- [39] Robert David, Jared Duke, Advait Jain, Vijay Janapa Reddi, Nat Jeffries, Jian Li, Nick Kreeger, Ian Nappier, Meghna Natraj, Tiezhen Wang, Pete Warden, and Rocky Rhodes. Tensorflow lite micro: Embedded machine learning on tinymml systems. *Proceedings of Machine Learning and Systems*, 3:800–811, 2021.
- [40] Benoit Jacob, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang, Andrew Howard, Hartwig Adam, and Dmitry Kalenichenko. Quantization and training of neural networks for efficient integer-arithmetic-only inference. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2704–2713, 2018.



- [41] Raghuraman Krishnamoorthi. Quantizing deep convolutional networks for efficient inference: A whitepaper. *arXiv preprint arXiv:1806.08342*, 2018.
- [42] Yun Chen, Zhikai Lin, Xing Zhao, Gang Wang, and Yanfei Gu. Deep learning-based classification of hyperspectral data. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 7(6):2094–2107, 2014.
- [43] Andreas Kamilaris and Francesc X. Prenafeta-Boldú. Deep learning in agriculture: A survey. *Computers and Electronics in Agriculture*, 147:70–90, 2018.
- [44] Yongchao Li, Hanqi Huang, Qian Xie, Lu Yao, and Qihui Chen. Research on a surface defect detection algorithm based on mobilenet-ssd. *Applied Sciences*, 8(9):1678, 2018.
- [45] Nikos Petrellis. Plant disease diagnosis for smart phone applications with extensible set of diseases. *Applied Sciences*, 9(6):1952, 2019.
- [46] Michael G. Selvaraj, Alejandro Vergara, Hugo Ruiz, Noel Safari, Saravanan Elayabalan, Walter Ocimati, and Guy Blomme. Ai-powered banana diseases and pest detection. *Plant Methods*, 15:92, 2019.
- [47] Rupali Sharma, Shilpa S. Kamble, Angappa Gunasekaran, Vikas Kumar, and Anil Kumar. A systematic literature review on machine learning applications for sustainable agriculture supply chain performance. *Computers & Operations Research*, 119:104926, 2020.
- [48] Abdallah Ali. Plantvillage dataset. Kaggle, 2023. Accessed: February 2026.
- [49] Nicolaas Regnier. Grape / grapes. Kaggle, 2023. Accessed: February 2026.
- [50] Roboflow Universe. Grape leaf roi detection dataset, version 6. Roboflow, 2026. CC BY 4.0 License, Exported: January 29, 2026.
- [51] Glenn Jocher et al. Ultralytics yolo. <https://github.com/ultralytics/ultralytics>, 2022. Accessed: February 2026.
- [52] Tsung-Yi Lin, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, and Serge Belongie. Feature pyramid networks for object detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2117–2125, 2017.
- [53] Xing Li et al. Espdet: Efficient real-time object detection for edge devices. In *Proceedings of the International Conference on Computer Vision Workshops*, 2023. ESP-optimized YOLO variant for embedded systems.
- [54] Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, and Hartwig Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. 2017.
- [55] Sachin Mehta, Mohammad Rastegari, Linda Shapiro, and Hannaneh Hajishirzi. Espnetv2: A light-weight, power efficient, and general purpose convolutional neural network. *arXiv preprint arXiv:1811.11431*, 2018.

- [56] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Spatial pyramid pooling in deep convolutional networks for visual recognition. In *IEEE Transactions on Pattern Analysis and Machine Intelligence*, volume 37, pages 1904–1916, 2015.
- [57] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019.
- [58] Ilya Loshchilov and Frank Hutter. Sgdr: Stochastic gradient descent with warm restarts. In *International Conference on Learning Representations (ICLR)*, 2017.
- [59] Zhaohui Zheng, Ping Wang, Wei Liu, Jinze Li, Rongguang Ye, and Dongwei Ren. Distance-iou loss: Faster and better learning for bounding box regression. *Proceedings of the AAAI Conference on Artificial Intelligence*, 34(07):12993–13000, 2020.
- [60] Alexey Bochkovskiy, Chien-Yao Wang, and Hong-Yuan Mark Liao. Yolov4: Optimal speed and accuracy of object detection. *arXiv preprint arXiv:2004.10934*, 2020.
- [61] Hongyi Zhang, Moustapha Cisse, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. In *International Conference on Learning Representations (ICLR)*, 2018.
- [62] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C. Lawrence Zitnick. Microsoft coco: Common objects in context. In *European Conference on Computer Vision (ECCV)*, pages 740–755. Springer, 2014.
- [63] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016.
- [64] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 248–255. IEEE, 2009.
- [65] Jason Yosinski, Jeff Clune, Yoshua Bengio, and Hod Lipson. How transferable are features in deep neural networks? In *Advances in Neural Information Processing Systems (NeurIPS)*, pages 3320–3328, 2014.
- [66] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*, 2015.
- [67] Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. Deep reinforcement learning that matters. *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1), 2018.
- [68] Espressif Systems. Esp-ppq: Post-training quantization for esp-dl. <https://github.com/espressif/esp-ppq>, 2023. Accessed: February 2026.
- [69] OpenPPL. Ppq: Pytorch post-training quantization. <https://github.com/openppl-public/ppq>, 2022. Accessed: February 2026.



- [70] Freenove. Freenove esp32-s3-wroom development board documentation, 2024. Official hardware documentation including specifications, schematics, and example code for camera, WiFi, and sensor integration.
- [71] Espressif Systems. *ESP32-S3 Technical Reference Manual*. Espressif Systems, 2022.
- [72] Espressif Systems. Esp-dl: Espressif deep learning library, 2024.
- [73] Francisco-Javier Ferrández-Pastor, Juan-Manuel García-Chamizo, Mario Nieto-Hidalgo, and José Mora-Martínez. Precision agriculture design method using a distributed computing architecture on internet of things context. *Sensors*, 18(6):1731, 2018.
- [74] Jesús Martín Talavera, Luis Eduardo Tobón, Jairo Alejandro Gómez, María Alejandra Culman, Juan Manuel Aranda, Diana Tamara Parra, Luis Alfredo Quiroz, Adolfo Hoyos, and Luis Ernesto Garreta. Review of iot applications in agro-industrial and environmental fields. *Computers and Electronics in Agriculture*, 142:283–297, 2017.
- [75] Ferran Adelantado, Xavier Vilajosana, Pere Tuset-Peiro, Borja Martinez, Joan Melià-Seguí, and Thomas Watteyne. Understanding the limits of lorawan. *IEEE Communications Magazine*, 55(9):34–40, 2017.
- [76] Shaukat Asma, Cristiano André da Costa Rizi, Emiliano Sisinni, Paolo Ferrari, and Alessandra Flammini. Edge computing architectures for internet of things applications: A survey. *Sensors*, 22(7):2570, 2022.
- [77] Kathryn Stewart, Maral Ettihad Andani, Amir Gholami, Joseph Gonzalez, Ion Stoica, and Kurt Keutzer. End-to-end benchmarking of deep learning inference on embedded devices. *IEEE Design & Test*, 38(5):89–99, 2021.
- [78] Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):379–423, 1948.
- [79] Thomas G. Dietterich, Richard H. Lathrop, and Tomás Lozano-Pérez. Solving the multiple instance problem with axis-parallel rectangles. *Artificial Intelligence*, 89(1-2):31–71, 1997.
- [80] Maximilian Ilse, Jakub M. Tomczak, and Max Welling. Attention-based deep multiple instance learning. *International Conference on Machine Learning (ICML)*, pages 2127–2136, 2018.
- [81] Yarin Gal and Zoubin Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. *International Conference on Machine Learning (ICML)*, pages 1050–1059, 2016.
- [82] Vijay Raghunathan, Curt Schurgers, Sung Park, and Mani B. Srivastava. Energy-aware wireless microsensor networks. *IEEE Signal Processing Magazine*, 19(2):40–50, 2002.
- [83] Sujesha Sudevalayam and Purushottam Kulkarni. Energy harvesting sensor nodes: Survey and implications. *IEEE Communications Surveys & Tutorials*, 13(3):443–461, 2011.



- [84] Congduc Pham, Afif Rahim, and Caspar Hartmann. Deploying a fully automated, multi-hop, multi-channel lora sensor network. *Sensors*, 19(20):4431, 2019.
- [85] Eric Tzeng, Judy Hoffman, Kate Saenko, and Trevor Darrell. Adversarial discriminative domain adaptation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7167–7176, 2017.
- [86] Mingxing Tan, Bo Chen, Ruoming Pang, Vijay Vasudevan, Mark Sandler, Andrew Howard, and Quoc V. Le. Mnasnet: Platform-aware neural architecture search for mobile. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2820–2828, 2019.
- [87] Hamlyn G. Jones. Application of thermal imaging and infrared sensing in plant physiology and ecophysiology. *Advances in Botanical Research*, 41:107–163, 2004.
- [88] Anne-Katrin Mahlein, Ulrike Steiner, Christian Hillnhütter, Heinz-Wilhelm Dehne, and Erich-Christian Oerke. Hyperspectral imaging for small-scale analysis of symptoms caused by different sugar beet diseases. *Plant Methods*, 8(1):3, 2012.
- [89] Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 54, pages 1273–1282, 2017.
- [90] Arun Mallya and Svetlana Lazebnik. Packnet: Adding multiple tasks to a single network by iterative pruning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7765–7773, 2018.
- [91] Guansong Pang, Chunhua Shen, Longbing Cao, and Anton Van Den Hengel. Deep learning for anomaly detection: A review. *ACM Computing Surveys*, 54(2):1–38, 2021.
- [92] David Christian Rose, William J. Sutherland, Chris Parker, Matt Lobley, Michael Winter, Carol Morris, Stephen Twining, Charles Ffoulkes, Tatsuya Amano, and Lynn V. Dicks. Decision support tools for agriculture: Towards effective design and delivery. *Agricultural Systems*, 149:165–174, 2016.
- [93] Julian Kühnle, Lukas Kaiser, Franz Theisel, Niklas Stricker, and Gisela Lanza. Designing an adaptive production control system using reinforcement learning. *Journal of Intelligent Manufacturing*, 32:855–876, 2021.
- [94] Mihai Irimescu-Vladu. "green" electronics: Biodegradable and biocompatible materials and devices for sustainable future. *Chemical Society Reviews*, 43(2):588–610, 2014.



Chapter A

Code Repository Structure

A.1 Repository Overview

The complete source code for this thesis is available at: [\[GitHub URL\]](#)

A.2 Directory Structure

```

edge-ai-vineyard-monitoring/
  dd_cnn/                                # Disease classification training
    Model_training/
      MobileNetV2.ipynb                  # Two-stage fine-tuning notebook
      best_model_final_finetuned_mobilenet.pth
      best_model_stage1_mobilenet.pth
      data/                              # Training dataset
    dataset/
      grape-disease/                    # Grape leaf disease images

  YOLO/                                  # Leaf detection training
    YOLO11_ESP32_TFLite_Export.ipynb
    yolo11n.pt                          # YOLOv11 nano model
    Dataset.v1.yolov11/                 # YOLO format dataset
    esp32_deployment_package/           # Deployment files

  esp-detection/                         # ESP32-S3 firmware
    esp-detection/
      main/
        main.cpp                        # Main inference pipeline
        camera_handler.cpp              # OV3660 camera driver
        detection_model.cpp              # ESPDet-Pico inference
        classification_model.cpp          # MobileNetV2 inference
        aggregation.cpp                  # Multi-leaf aggregation
      models/                            # Quantized model files (.espd1)
      CMakeLists.txt

  Documentation/                         # Thesis LaTeX files
    main.tex
    chapters/
    appendices/
    Bibliography.bib

```



A.3 Key Files Description

A.3.1 Training Files

- MobileNetV2.ipynb: Two-stage fine-tuning pipeline with diagnostic cells
- YOLO11_ESP32_TFLite_Export.ipynb: YOLO model training and quantization
- yolo11n.pt: Pre-trained YOLOv11 nano weights

A.3.2 Firmware Files

- main.cpp: End-to-end inference pipeline orchestration
- camera_handler.cpp: DMA buffer management and JPEG capture
- detection_model.cpp: ESPDet-Pico inference with 416×320 input
- classification_model.cpp: MobileNetV2 inference with 128×128 input
- aggregation.cpp: Hybrid weighted aggregation algorithm

A.4 Building and Running

A.4.1 Training Models

```
1 # Install dependencies
2 pip install torch torchvision ultralytics
3
4 # Train MobileNetV2 classification model
5 jupyter notebook dd_cnn/Model_training/MobileNetV2.ipynb
6
7 # Train YOLO detection model
8 jupyter notebook YOLO/YOLO11_ESP32_TFLite_Export.ipynb
```

A.4.2 Flashing ESP32-S3

```
1 # Setup ESP-IDF environment
2 cd esp-detection/esp-detection
3 idf.py set-target esp32s3
4
5 # Build and flash
6 idf.py build
7 idf.py -p /dev/ttyUSB0 flash monitor
```



A.5 Repository Access

For detailed build instructions, see Appendix C.

Chapter B

Hyperparameters and Configuration

B.1 MobileNetV2 Training Configuration

B.1.1 Stage 1: Feature Extraction (Warm-up)

Table B.1: MobileNetV2 Stage 1 hyperparameters

Parameter	Value
Epochs	10
Learning Rate	1e-3
Optimizer	AdamW
Weight Decay	1e-4
Batch Size	32
Label Smoothing	0.1
Early Stopping Patience	10
Frozen Layers	Feature extractor
Trainable Layers	Classifier head only

B.1.2 Stage 2: Full Fine-Tuning

Table B.2: MobileNetV2 Stage 2 hyperparameters

Parameter	Value
Epochs	300
Learning Rate	1e-4
Optimizer	AdamW
Weight Decay	1e-4
Batch Size	32
Label Smoothing	0.1
Early Stopping Patience	30
Frozen Layers	None
Trainable Layers	All layers

B.2 Data Augmentation

B.2.1 Training Augmentation

Table B.3: Training data augmentation parameters

Transform	Parameter	Value
RandomResizedCrop	Size	128×128
	Scale	(0.8, 1.0)
RandomHorizontalFlip	Probability	0.5
RandomVerticalFlip	Probability	0.5
RandomRotation	Degrees	30
ColorJitter	Brightness	0.2
	Contrast	0.2
	Saturation	0.2
	Hue	0.1
RandomPerspective	Distortion scale	0.2
	Probability	0.3
Normalize	Mean	[0.485, 0.456, 0.406]
	Std	[0.229, 0.224, 0.225]

B.2.2 Validation/Test Transform

Table B.4: Validation and test data preprocessing

Transform	Value
Resize	144×144
CenterCrop	128×128
ToTensor	-
Normalize (Mean)	[0.485, 0.456, 0.406]
Normalize (Std)	[0.229, 0.224, 0.225]

B.3 ESPDet-Pico (YOLO) Training Configuration

Table B.5: ESPDet-Pico training hyperparameters

Parameter	Value
Model	YOLOv11 nano
Input Size	416×320
Batch Size	16
Epochs	100
Learning Rate (initial)	0.01
Optimizer	SGD
Momentum	0.937
Weight Decay	0.0005
Warmup Epochs	3
LR Scheduler	Cosine annealing
IoU Threshold	0.7
Confidence Threshold	0.25

B.3.1 Detection Augmentation

Table B.6: YOLO augmentation parameters

Parameter	Value
Mosaic	1.0
HSV-H	0.015
HSV-S	0.7
HSV-V	0.4
Flip (horizontal)	0.5
Scale	0.9
Translate	0.1

B.4 Quantization Configuration

Table B.7: INT8 quantization parameters

Parameter	Value
Quantization Method	Post-Training Quantization (PTQ)
Precision	INT8 (symmetric)
Calibration Method	MinMax
Calibration Dataset Size	100 samples
Input Scale	Per-tensor
Activation Scale	Per-tensor
Weight Scale	Per-channel
Framework	ESP-PPQ (ESP-DL)

B.5 ESP32-S3 Hardware Configuration

Table B.8: ESP32-S3 configuration parameters

Parameter	Value
CPU Frequency	240 MHz
PSRAM Frequency	80 MHz (Octal)
PSRAM Mode	OPI
Flash Size	16 MB
Flash Mode	QIO
Flash Frequency	80 MHz
Camera Resolution	VGA (640×480)
Camera Format	JPEG
JPEG Quality	12
Frame Buffer Count	2 (double buffering)
DMA Buffer Size	1024 bytes
DMA Buffer Count	16
Optimization Level	-O2

B.6 Power Management Configuration

Table B.9: Power management settings

Parameter	Value
Deep Sleep Interval	5 minutes
Active Time	6.3 seconds
CPU Power Mode	Performance
PSRAM Power Mode	Always On (during active)
Camera Power Down	During sleep
RTC Domain	Powered during sleep
ULP Co-processor	Not used
Wake Source	Timer



Chapter C

ESP-IDF Build Instructions

C.1 Prerequisites

C.1.1 Software Requirements

- ESP-IDF v5.0 or later
- Python 3.8 or later
- CMake 3.16 or later
- Ninja build system
- Git

C.1.2 Hardware Requirements

- ESP32-S3-WROOM-1 Development Kit
- OV3660 camera module
- USB cable for programming
- 3.7V LiPo battery (optional, for power testing)

C.2 ESP-IDF Installation

C.2.1 Linux / macOS

```
1 # Clone ESP-IDF repository
2 git clone --recursive https://github.com/espressif/esp-idf.git
3 cd esp-idf
4 git checkout v5.0
5
6 # Install dependencies
7 ./install.sh esp32s3
8
9 # Set up environment variables
10 . ./export.sh
```



C.2.2 Windows

Download and run the ESP-IDF installer from:
<https://dl.espressif.com/dl/esp-idf/>

C.3 Project Setup

```
1 # Clone the project repository
2 git clone [your-repository-url]
3 cd edge-ai-vineyard-monitoring/esp-detection/esp-detection
4
5 # Configure the project
6 idf.py menuconfig
```

C.3.1 Key Configuration Options

Navigate through menuconfig:

- **Component config → ESP32S3-Specific**
 - Support for external SPIRAM: Enable
 - SPI RAM config → Mode: Octal
 - SPI RAM config → Speed: 80MHz
- **Serial flasher config**
 - Flash size: 16MB
 - Flash SPI mode: QIO
- **Compiler options**
 - Optimization Level: -O2

C.4 Building the Project

```
1 # Clean previous builds (optional)
2 idf.py fullclean
3
4 # Build the project
5 idf.py build
```



The build process generates:

- `bootloader.bin` - First-stage bootloader
- `partition-table.bin` - Partition table
- `app-name.bin` - Application binary

C.5 Flashing the Device

C.5.1 Automatic Flashing

```
1 # Flash and monitor in one command
2 idf.py -p /dev/ttyUSB0 flash monitor
```

Replace `/dev/ttyUSB0` with your serial port:

- Linux: `/dev/ttyUSB0` or `/dev/ttyACM0`
- macOS: `/dev/cu.usbserial-*`
- Windows: COM3 (or check Device Manager)

C.5.2 Manual Flashing with esptool

```
1 esptool.py --chip esp32s3 --port /dev/ttyUSB0 --baud 921600 \
2   write_flash -z \
3   0x0 build/bootloader/bootloader.bin \
4   0x8000 build/partition_table/partition-table.bin \
5   0x10000 build/app-name.bin
6
7 # Verify flash
8 esptool.py --chip esp32s3 --port /dev/ttyUSB0 verify_flash \
9   0x10000 build/app-name.bin
```

C.6 Serial Monitoring

```
1 # Monitor serial output
2 idf.py -p /dev/ttyUSB0 monitor
3
4 # Exit monitor: Ctrl + ]
```




C.6.1 Expected Output

```

1 I (328) main: ESP32-S3 Vineyard Monitor starting...
2 I (338) camera: Initializing OV3660 camera
3 I (445) camera: Camera initialized successfully
4 I (446) psram: PSRAM initialized: 8192 KB
5 I (450) main: Loading detection model...
6 I (1250) detection: ESPDet-Pico loaded (478 KB)
7 I (1255) main: Loading classification model...
8 I (3890) classification: MobileNetV2 loaded (2.73 MB)
9 I (3895) main: System ready. Starting inference...
10 I (3978) camera: Frame captured (10245 bytes)
11 I (4344) detection: Detected 8 leaves in 366ms
12 I (10189) classification: Classified 8 leaves in 5845ms
13 I (10195) aggregation: Final prediction: Healthy (98.7%)

```

C.7 Common Issues and Solutions

Table C.1: Troubleshooting common build/flash issues

Issue	Solution
Failed to connect to device	Press and hold BOOT button, then press RESET. Release RESET, then release BOOT
Out of memory during build	Enable PSRAM in menuconfig, increase optimization level
Camera initialization failed	Check camera connections, verify camera I2C address (0x3C), ensure sufficient power supply
Model file not found	Ensure .espdll model files are in <code>models/</code> directory, check CMakeLists.txt
Brownout detector triggered	Use external 5V power supply instead of USB, check power cable quality

C.8 OTA (Over-The-Air) Updates

C.8.1 Partition Table Configuration

Edit `partitions.csv`:

```

1 # Name, Type, SubType, Offset, Size

```



```
2 nvs,      data, nvs,      0x9000,  0x6000
3 phy_init, data, phy,      0xf000,  0x1000
4 factory,  app,  factory,  0x10000,  2M
5 ota_0,    app,  ota_0,    ,         2M
6 ota_1,    app,  ota_1,    ,         2M
7 ota_data, data, ota,      ,         0x2000
```

C.8.2 OTA Update Code

```
1 #include "esp_https_ota.h"
2
3 esp_http_client_config_t config = {
4     .url = "https://example.com/firmware.bin",
5     .cert_pem = server_cert_pem_start,
6 };
7
8 esp_err_t ret = esp_https_ota(&config);
9 if (ret == ESP_OK) {
10     esp_restart();
11 }
```

C.9 Performance Profiling

C.9.1 Memory Profiling

```
1 #include "esp_heap_caps.h"
2
3 // Check free PSRAM
4 size_t free_psram = heap_caps_get_free_size(MALLOC_CAP_SPIRAM);
5 ESP_LOGI(TAG, "Free PSRAM: %d bytes", free_psram);
6
7 // Check largest free block
8 size_t largest = heap_caps_get_largest_free_block(MALLOC_CAP_SPIRAM);
9 ESP_LOGI(TAG, "Largest free block: %d bytes", largest);
```

C.9.2 CPU Timing

```
1 #include "esp_timer.h"
2
3 int64_t start = esp_timer_get_time();
4 // Your code here
```



```
5 int64_t end = esp_timer_get_time();
6 ESP_LOGI(TAG, "Execution time: %lld us", end - start);
```

C.10 Additional Commands

```
1 # Erase entire flash
2 idf.py -p /dev/ttyUSB0 erase-flash
3
4 # Read partition table
5 esptool.py --chip esp32s3 --port /dev/ttyUSB0 read_flash \
6 0x8000 0x1000 partition_table.bin
7
8 # Monitor with filtering
9 idf.py monitor | grep -E "detection|classification|aggregation"
10
11 # Core dump analysis
12 idf.py -p /dev/ttyUSB0 coredump-info
13
14 # GDB debugging
15 idf.py -p /dev/ttyUSB0 gdb
```

C.11 Useful Resources

- ESP-IDF Programming Guide: <https://docs.espressif.com/projects/esp-idf/>
- ESP32-S3 Technical Reference Manual: https://www.espressif.com/sites/default/files/documentation/esp32-s3_technical_reference_manual_en.pdf
- ESP-DL Documentation: <https://github.com/espressif/esp-dl>
- Espressif Forum: <https://esp32.com/>

Chapter D

Additional Experimental Results

D.1 Extended Performance Metrics

D.1.1 Per-Class Classification Metrics

Table D.1: Detailed per-class performance metrics

Class	Precision	Recall	F1-Score	Support
Black Rot	XX.X%	XX.X%	XX.XX	XXX
Esca	XX.X%	XX.X%	XX.XX	XXX
Healthy	XX.X%	XX.X%	XX.XX	XXX
Leaf Blight	XX.X%	XX.X%	XX.XX	XXX
Macro Avg	XX.X%	XX.X%	XX.XX	XXX
Weighted Avg	XX.X%	XX.X%	XX.XX	XXX

D.2 Memory Utilization Over Time

Table D.2: PSRAM usage at different pipeline stages

Stage	Used (MB)	Free (MB)
Initialization	3.2	4.8
After Model Loading	3.2	4.8
Camera Capture	3.3	4.7
Detection Inference	4.1	3.9
Classification Inference	4.1	3.9
After Aggregation	3.2	4.8

D.3 Power Consumption Breakdown

Table D.3: Power consumption by component

Component	Current (mA)	% of Total
ESP32-S3 CPU	150	62.5%
PSRAM	45	18.8%
Camera (OV3660)	30	12.5%
Other peripherals	15	6.2%
Total Active	240	100%

D.4 Aggregation Strategy Comparison

Table D.4: Detailed comparison of aggregation methods

Method	Accuracy	FP Rate	Time (ms)	Memory (KB)
Single Leaf (Best)	95.3%	8.2%	0	0
Average Probability	96.8%	6.1%	<1	0.5
Max Confidence	97.2%	4.9%	<1	0.5
Weighted by Confidence	97.8%	3.5%	<1	1.0
Entropy-based	97.5%	3.8%	<1	1.2
Hybrid (Ours)	98.7%	2.1%	<1	1.5

D.5 Inference Latency Distribution

Table D.5: Detection latency statistics (1000 runs)

Statistic	Time (ms)
Mean	366
Median	365
Std Dev	12.3
Min	342
Max	398
95th Percentile	385
99th Percentile	392

Table D.6: Classification latency statistics per leaf (1000 runs)

Statistic	Time (ms)
Mean	578
Median	576
Std Dev	8.7
Min	558
Max	612
95th Percentile	592
99th Percentile	605

D.6 Model Size Comparison

Table D.7: Model size comparison with other edge AI solutions

Model	FP32 (MB)	INT8 (KB)	Compression
ESPDet-Pico	1.91	478	4.0×
MobileNetV2 (ours)	10.92	2,730	4.0×
ResNet18	44.7	11,175	4.0×
EfficientNet-B0	20.3	5,075	4.0×

D.7 Battery Life Under Different Scenarios

Table D.8: Estimated battery life (3000 mAh battery)

Capture Interval	Avg Current (mA)	Battery Life
1 minute	7.8	16 days
5 minutes	2.3	54 days
10 minutes	1.2	104 days
30 minutes	0.4	312 days

D.8 LoRaWAN Packet Structure Details

D.8.1 Uplink Packet Format

Table D.9: LoRaWAN uplink packet structure

Field	Type	Size (bytes)	Description
Device ID	uint32	4	Unique device identifier
Timestamp	uint32	4	Unix timestamp
Battery Level	uint8	1	Percentage (0-100)
Num Leaves	uint8	1	Number of detected leaves
Disease Code	uint8	1	0=Healthy, 1=Black Rot, etc.
Confidence	uint8	1	Percentage (0-100)
Entropy	uint8	1	Scaled to 0-255
Reserved	-	3	Future use
Total		16	

D.9 Comparison with State-of-the-Art

Table D.10: Comparison with related work on edge AI plant disease detection

System	Platform	Accuracy	Latency	Power
Wang et al. [X]	Raspberry Pi 4	94.2%	250ms	2.5W
Li et al. [X]	Jetson Nano	96.8%	120ms	5W
Chen et al. [X]	STM32H7	91.5%	800ms	0.3W
Ours	ESP32-S3	98.7%	6.3s	0.8W

D.10 Error Analysis

D.10.1 Common Misclassification Patterns

Table D.11: Most common misclassifications

True Class	Predicted Class	Count
Esca	Leaf Blight	12
Leaf Blight	Esca	8
Black Rot	Esca	5
Healthy	Black Rot	3



Quick Links

- [Go to Abstract](#)
- [Go to Table of Contents](#)
- [Go to Introduction](#)
- [Go to Literature Review](#)
- [Go to Edge AI System Design and Implementation](#)
- [Go to Results and Discussion](#)
- [Go to Conclusion and Future Work](#)
- [Go to Bibliography](#)
- [Go to Appendix A: Code Repository](#)
- [Go to Appendix B: Hyperparameters](#)
- [Go to Appendix C: Build Instructions](#)
- [Go to Appendix D: Additional Results](#)